

Lecture 16:

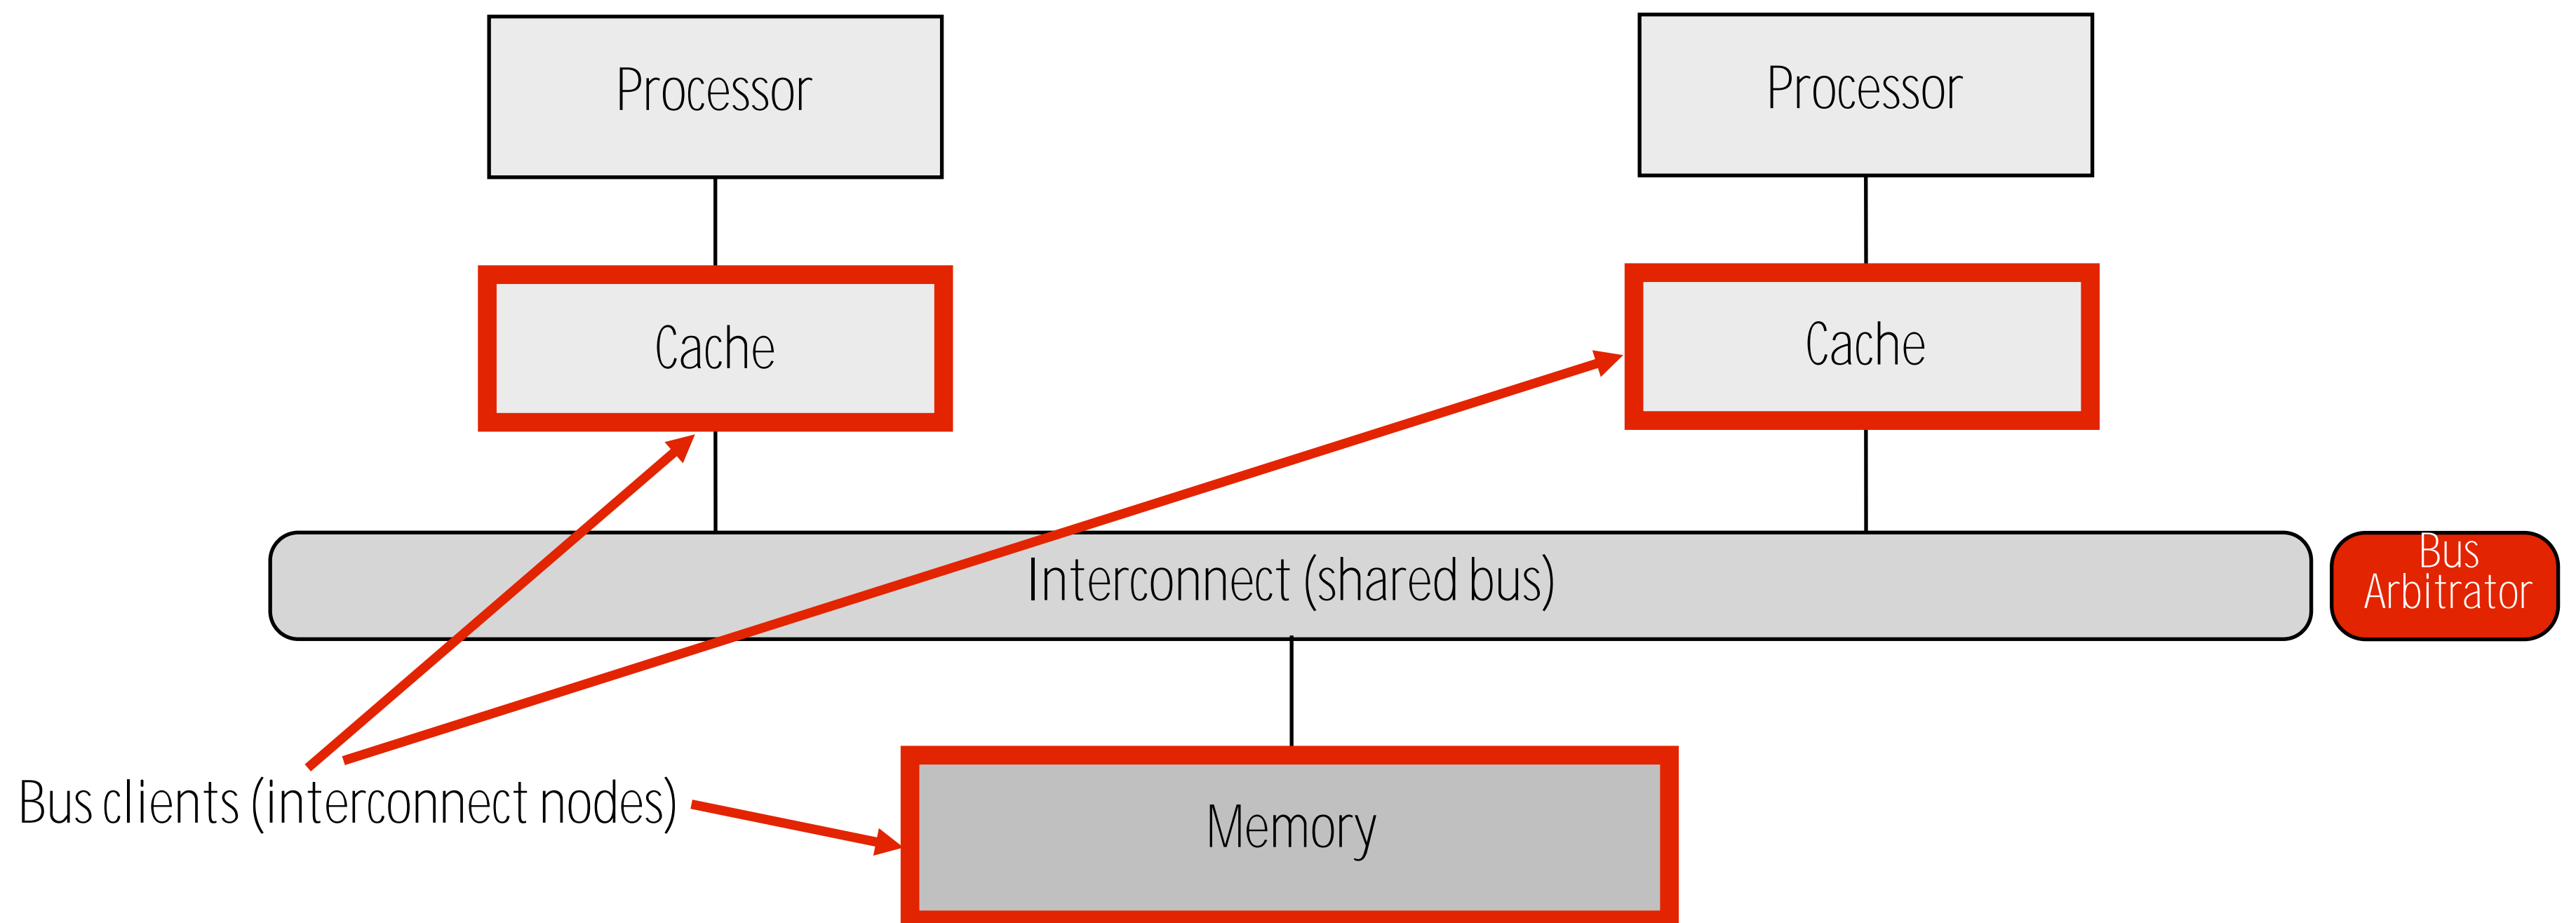
Interconnection Networks

Parallel Computer Architecture and Programming
CMU 15-418/15-618, Spring 2021

Learning Objectives

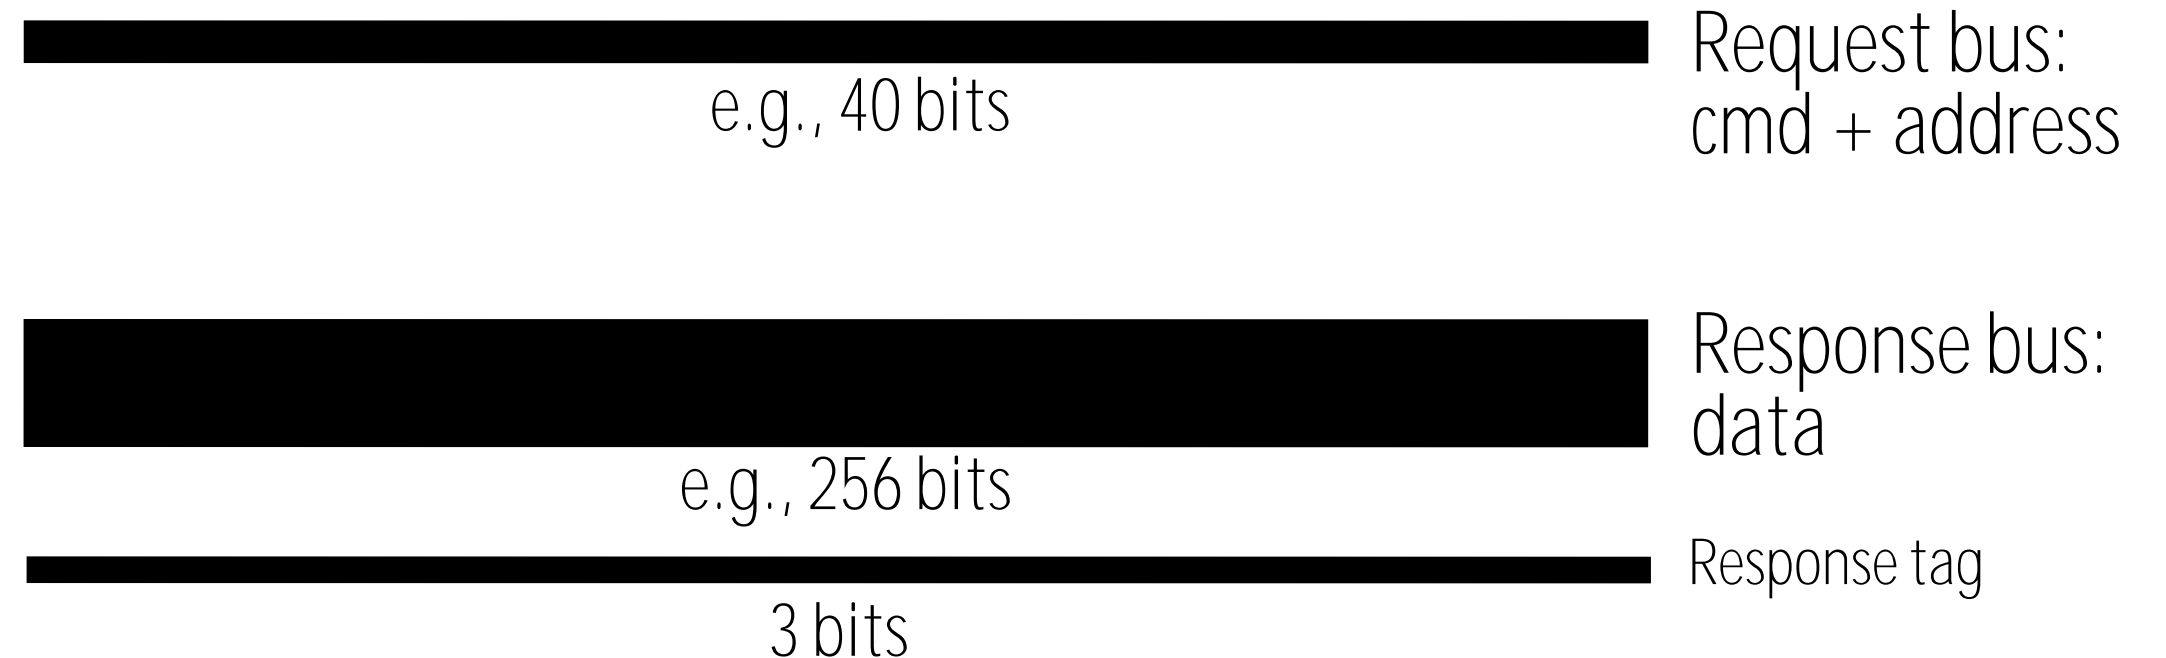
- After this lecture, you should be able to answer:
 - How are hardware components interconnected?
 - What is the basic cost of communication?
 - How does data traverse the network?

Basic system design from previous lectures

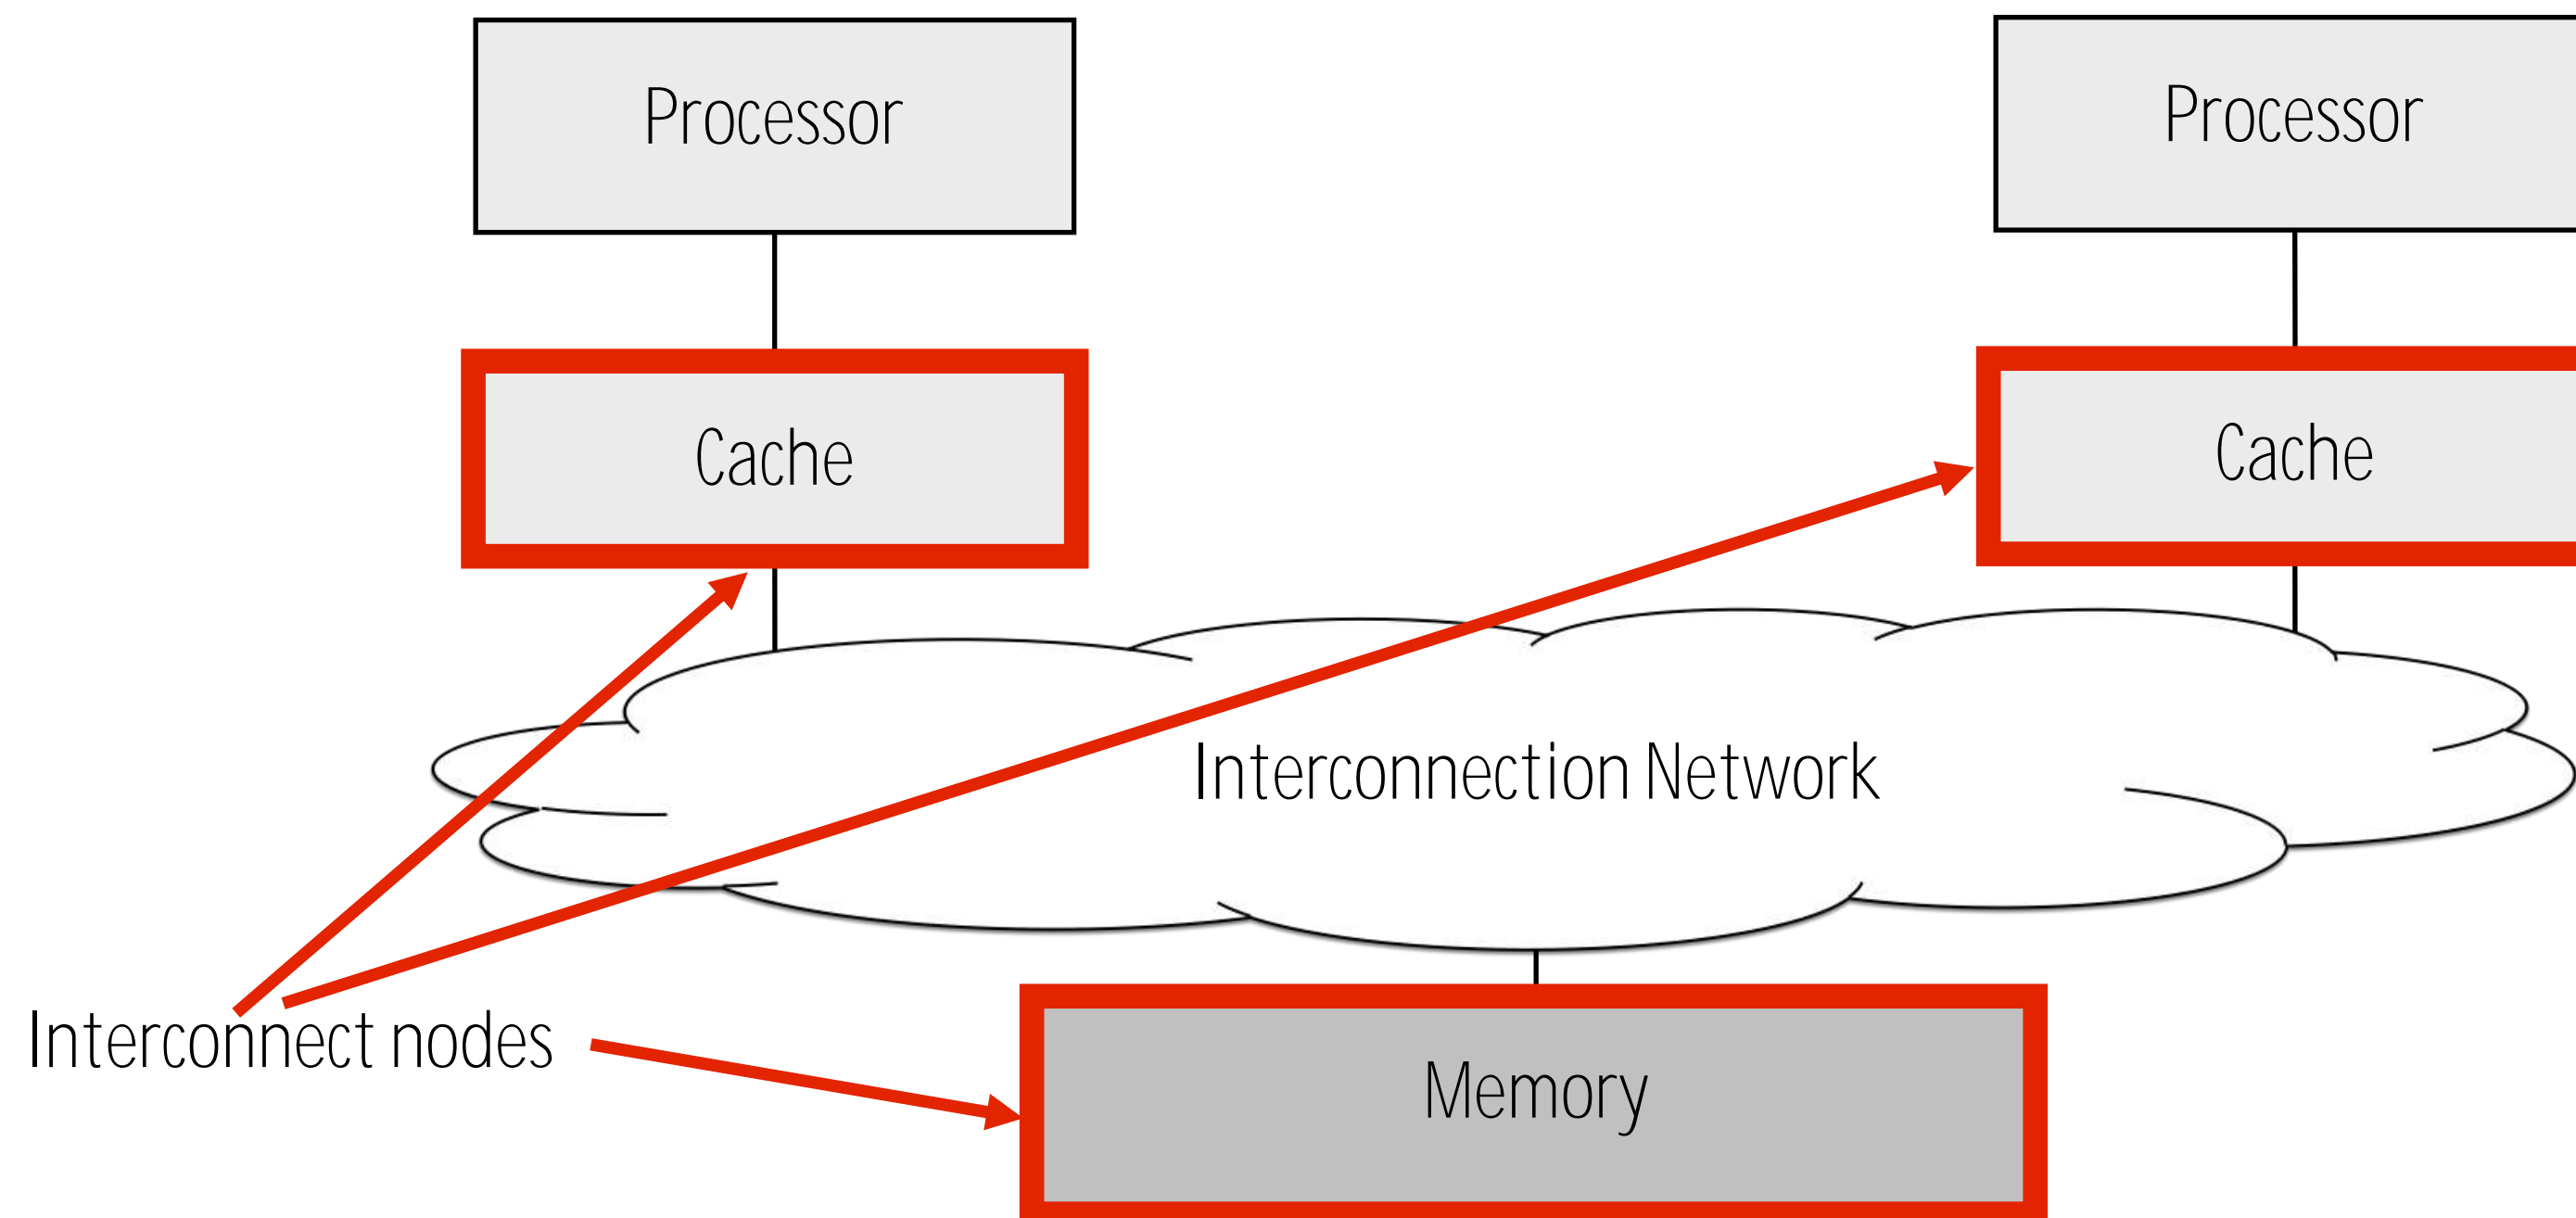


Bus interconnect:

All nodes connected by a shared set of wires



Today: modern interconnect designs



Today's topics: the basic ideas of building a high-performance interconnection network in a parallel processor.
(think: “a network-on-a-chip”)

What are interconnection networks used for?

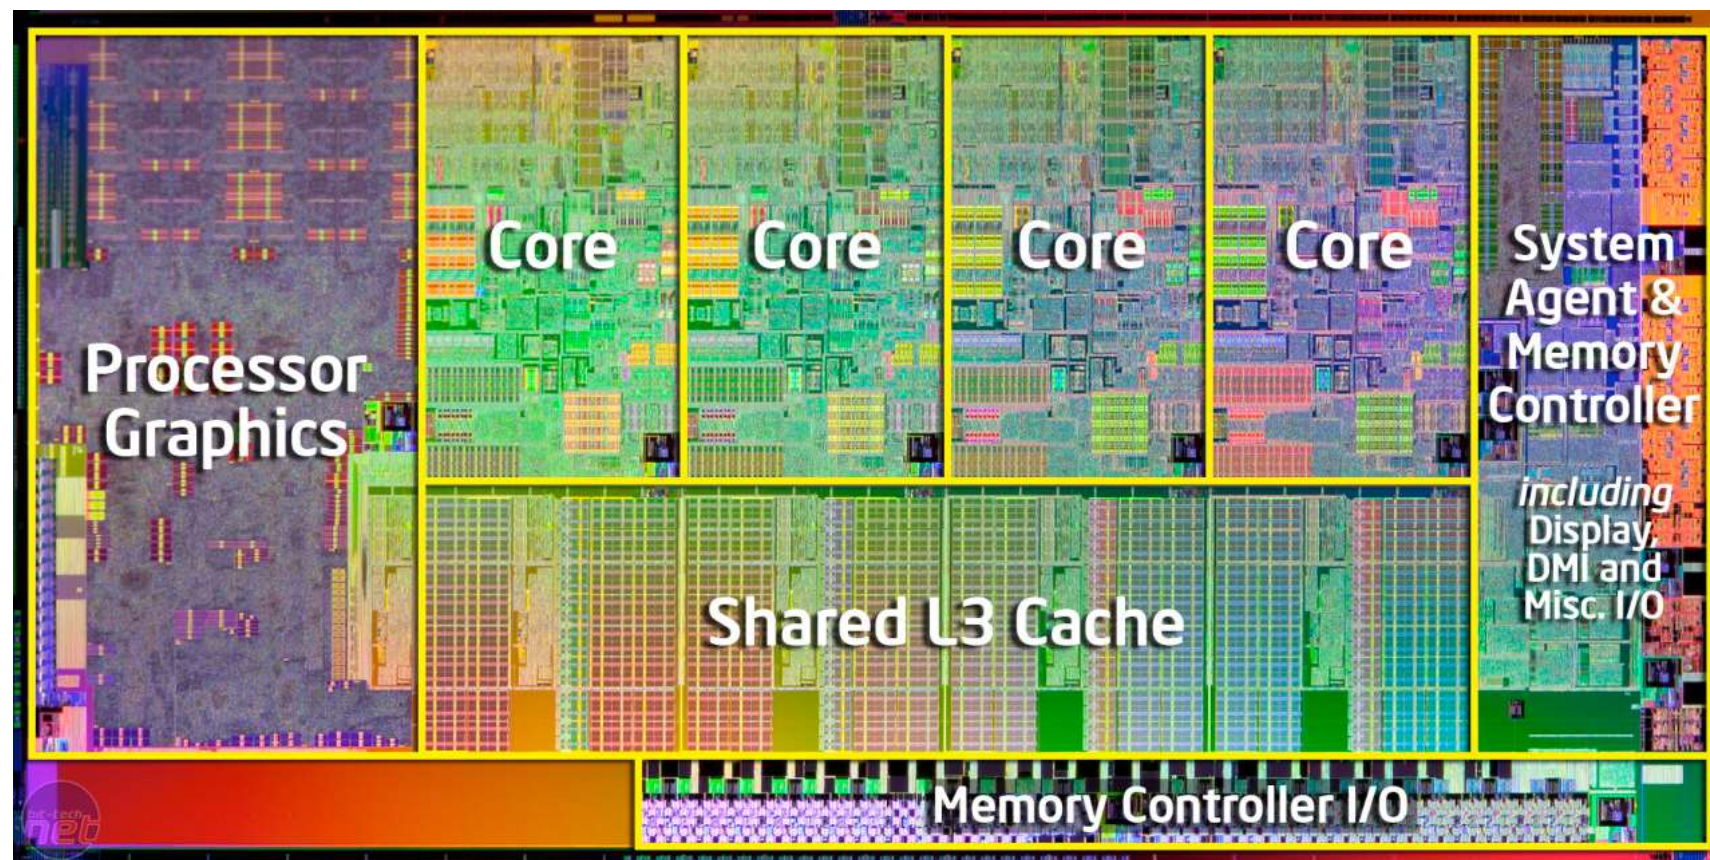
- To connect:
 - Processor cores with other cores
 - Processors and memories
 - Processor cores and caches
 - Caches and caches
 - I/O devices

Why is the design of the interconnection network important?

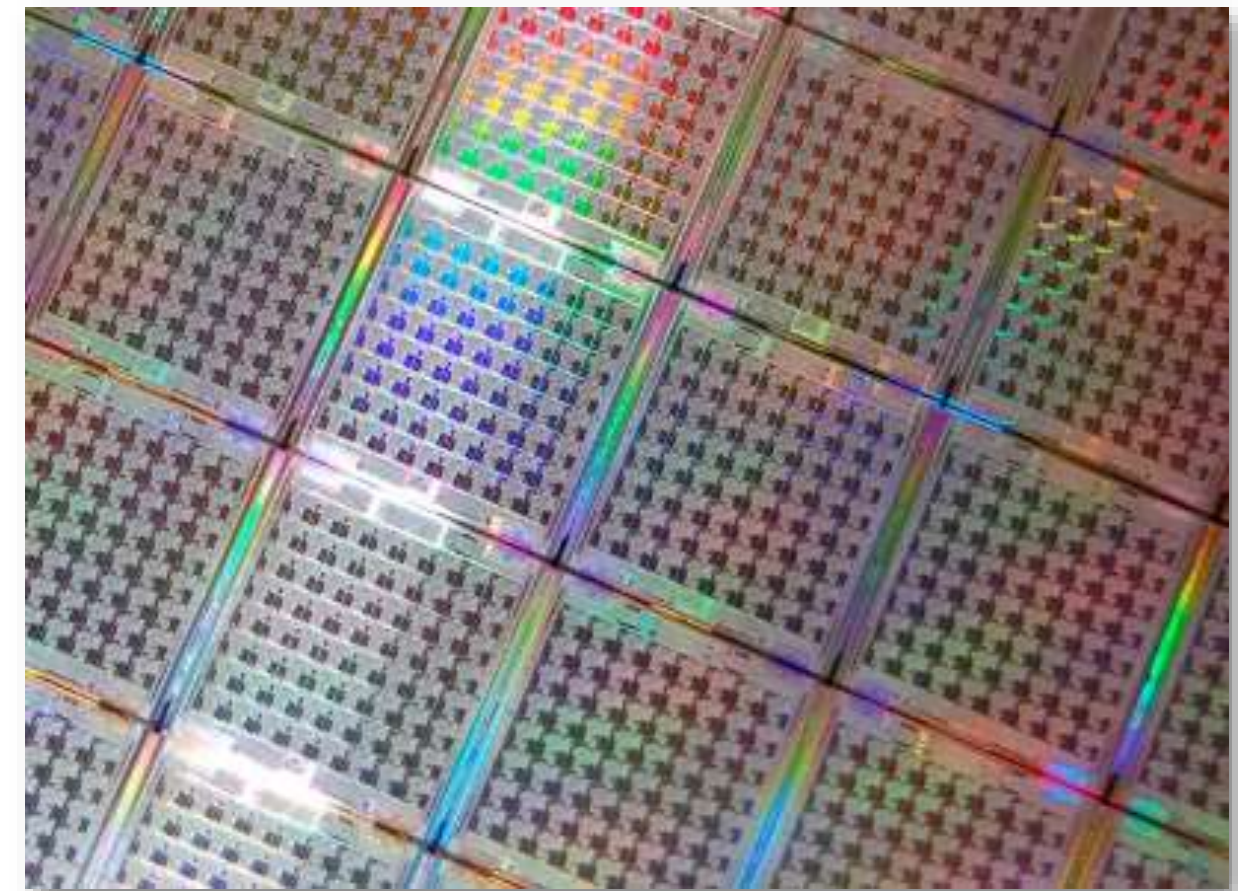
- System scalability
 - How large of a system can be built?
 - How easy is it to add more nodes (e.g., cores)
- System performance and energy efficiency
 - How fast can cores, caches, memory communicate
 - How long is latency to memory?
 - How much energy is spent on communication?

With increasing core counts...

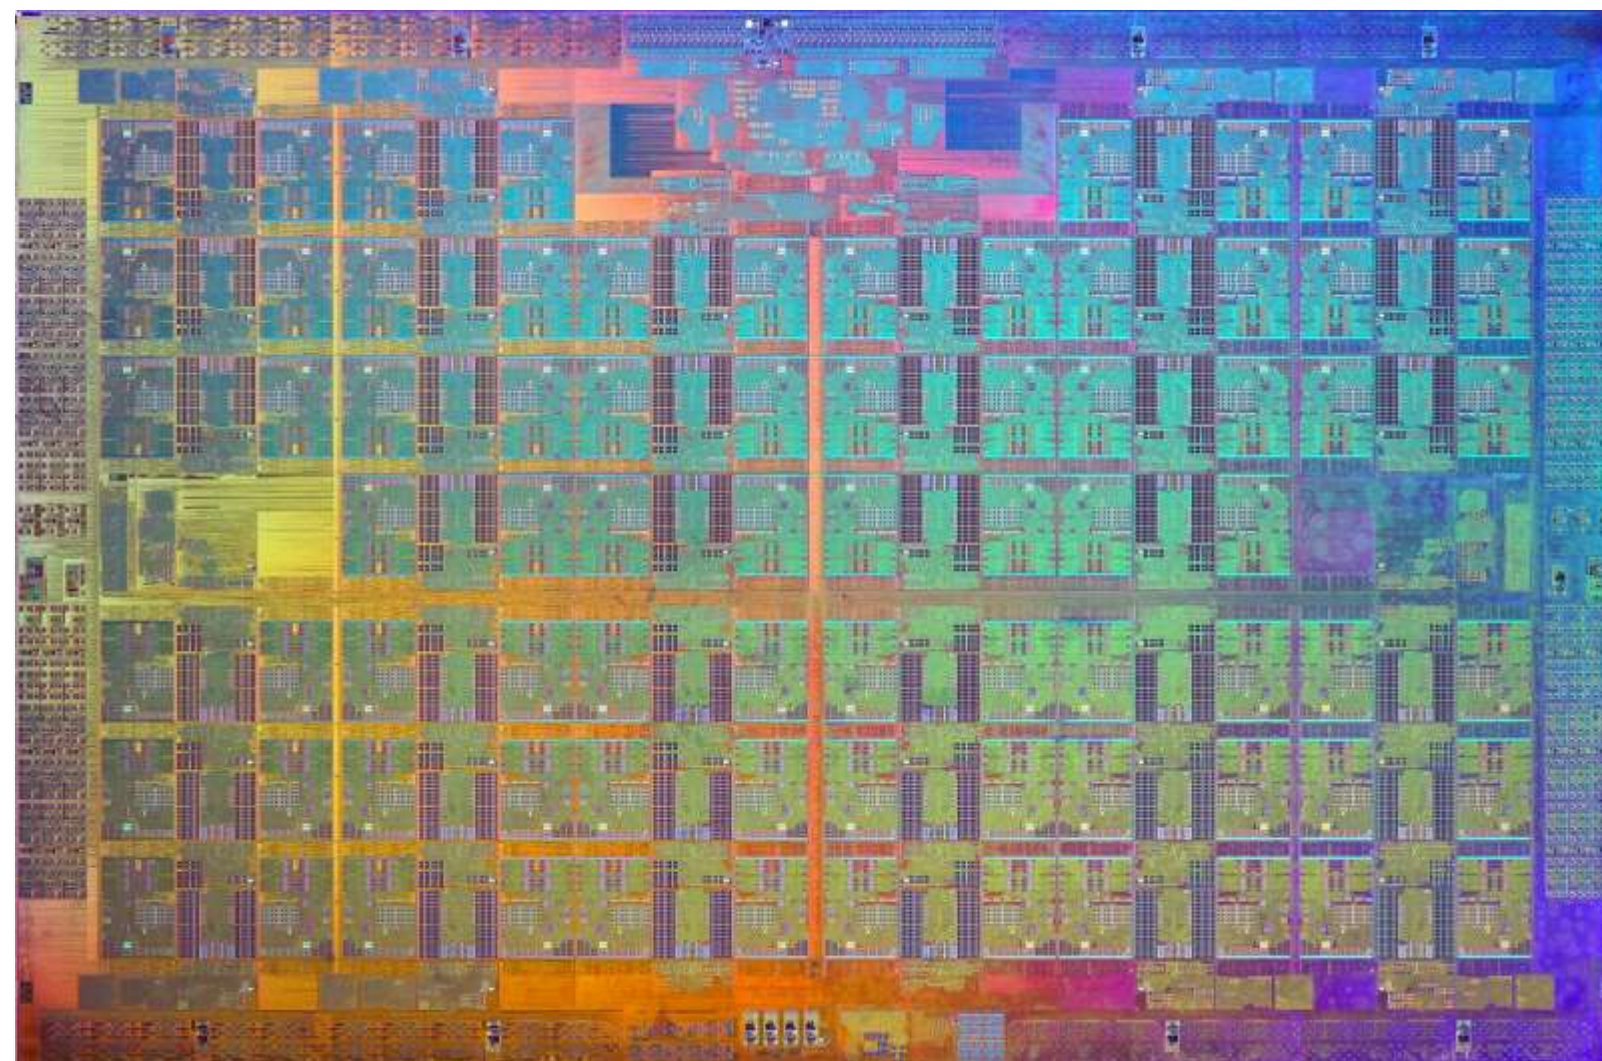
Scalability of on-chip interconnection network becomes increasingly important



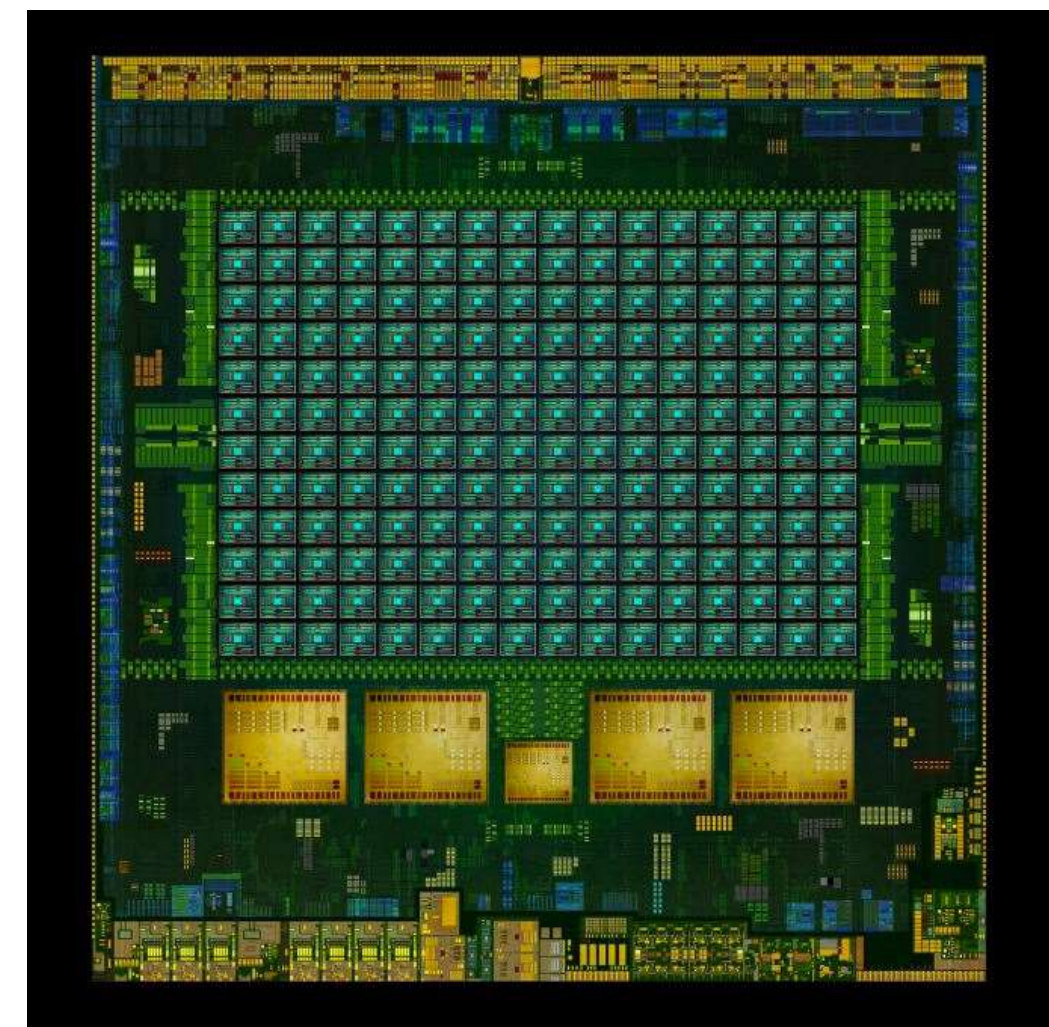
Intel core i7 (4-CPU cores, + GPU)



Tiler GX 64-core chip



Intel Xeon Phi (72-core x86)

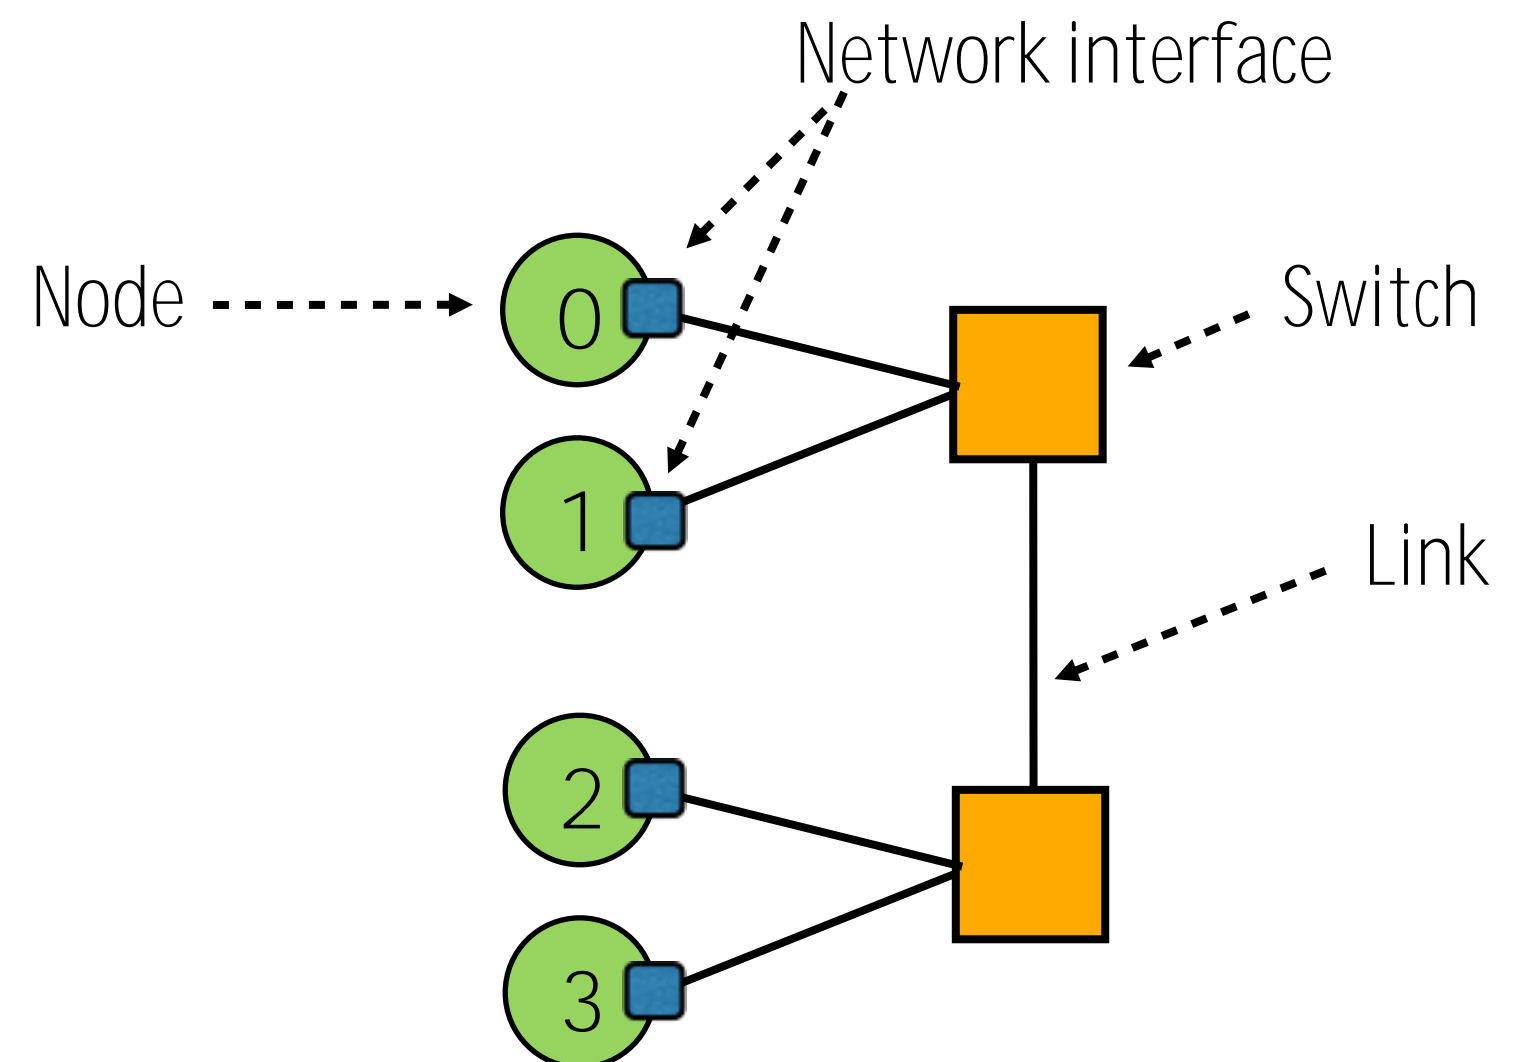


Tegra K1: 4 + 1 ARM cores + GPU cores

Interconnect terminology

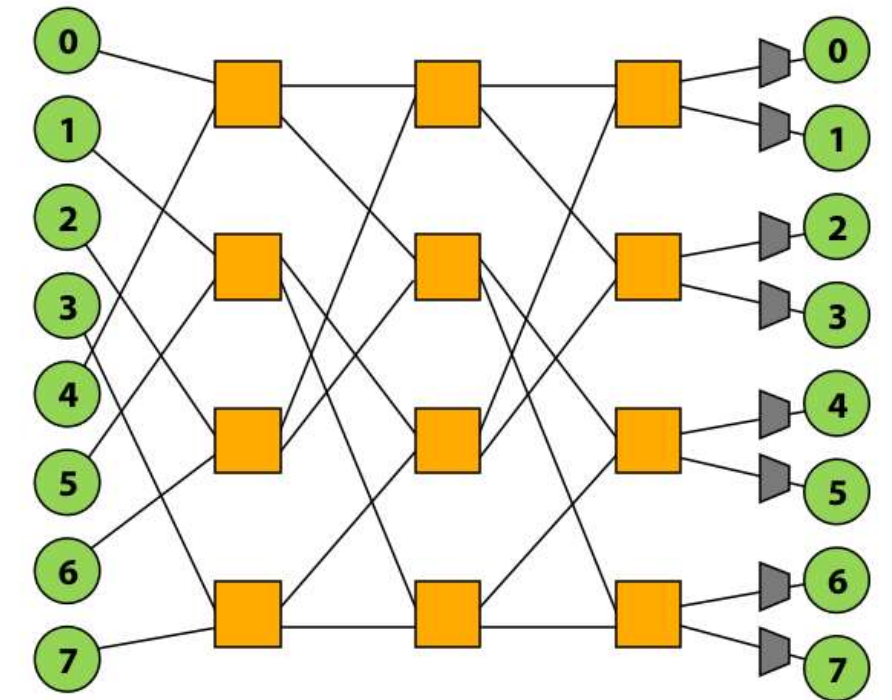
Terminology

- Network node: a network endpoint connected to a router/switch
 - Examples: processor caches, the memory controller
- Network interface:
 - Connects nodes to the network
- Switch/router:
 - Connects a fixed number of input links to a fixed number of output links
- Link:
 - A bundle of wires carrying a signal



Design issues

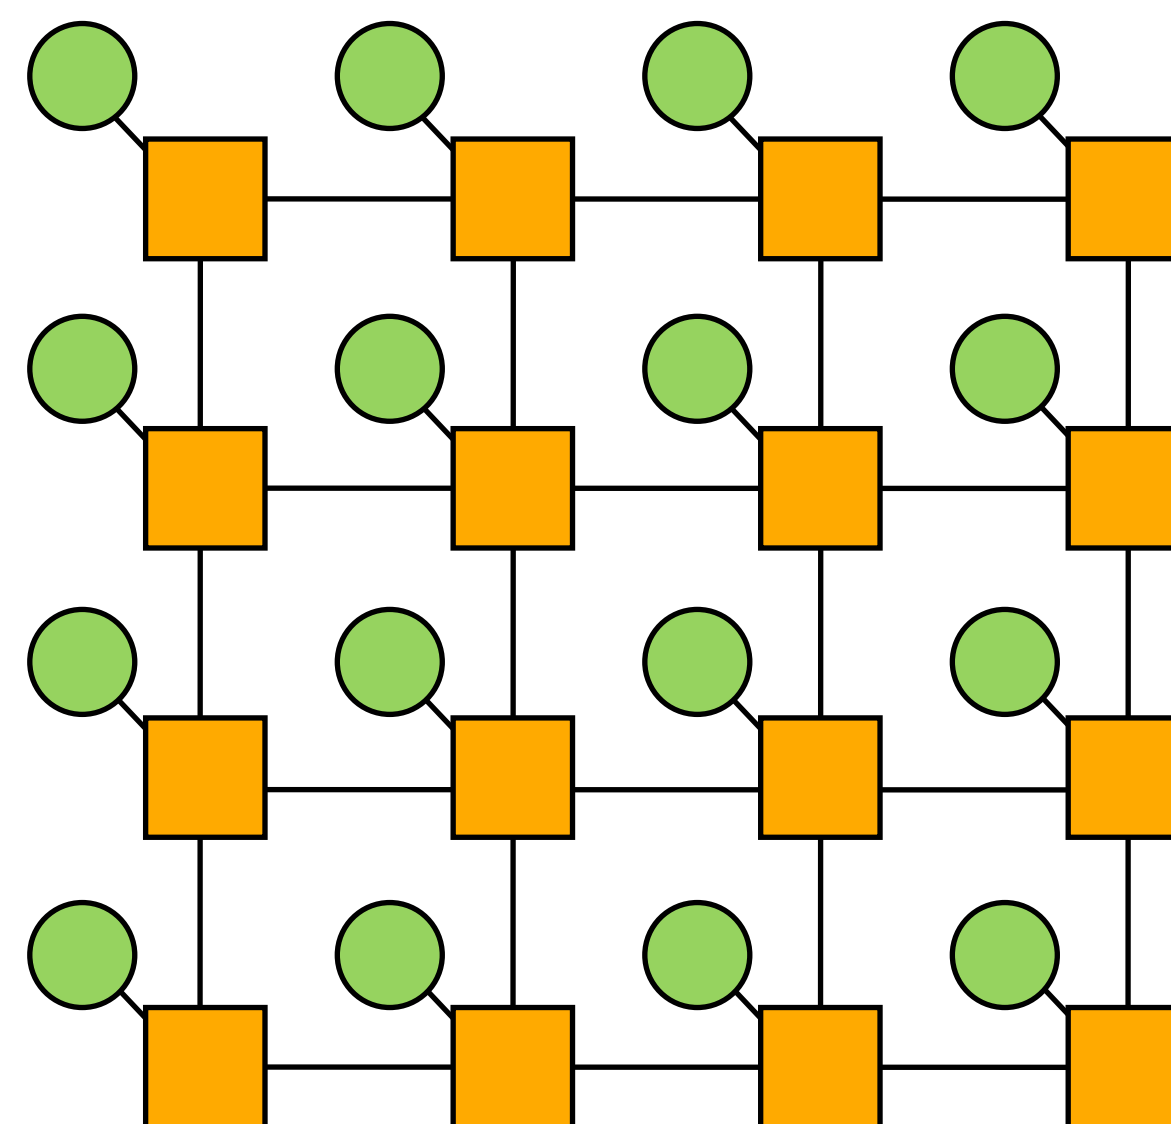
- Topology: how switches are connected via links
 - Affects routing, throughput, latency, complexity/cost of implementation
- Routing: how a message gets from its source to its destination in the network
 - Can be static (messages take a predetermined path) or adaptive based on load
- Buffering and flow control
 - What data is stored in the network? packets, partial packets? etc.
 - How does the network manage buffer space?



Properties of interconnect topology

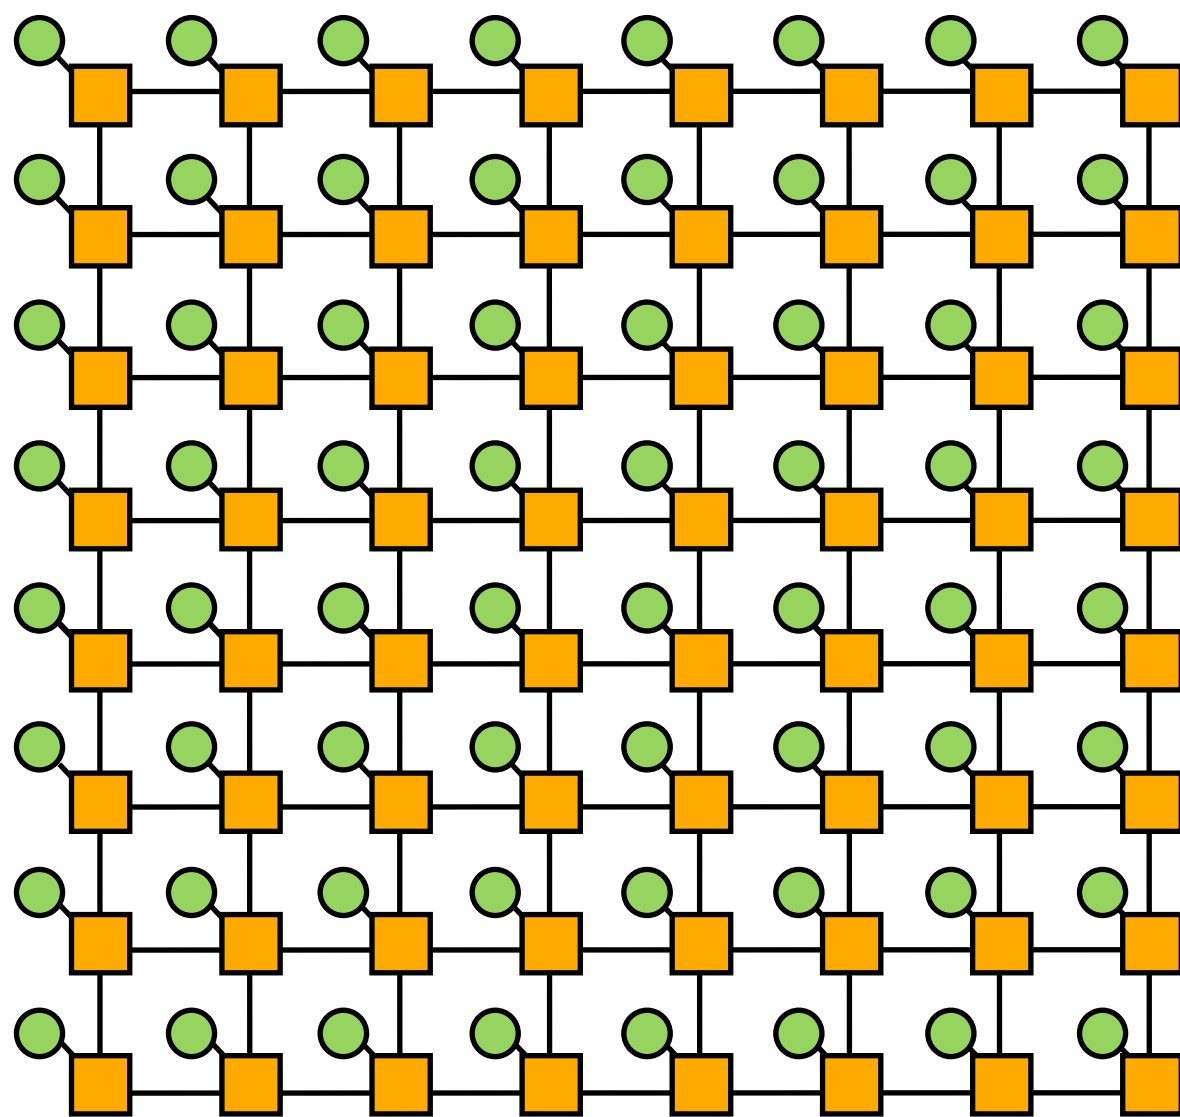
- Routing distance
 - **Number of links (“hops”) along a route between two nodes**
- Diameter: the maximum routing distance
- Average distance: average routing distance over all valid routes

Example:
diameter = 6

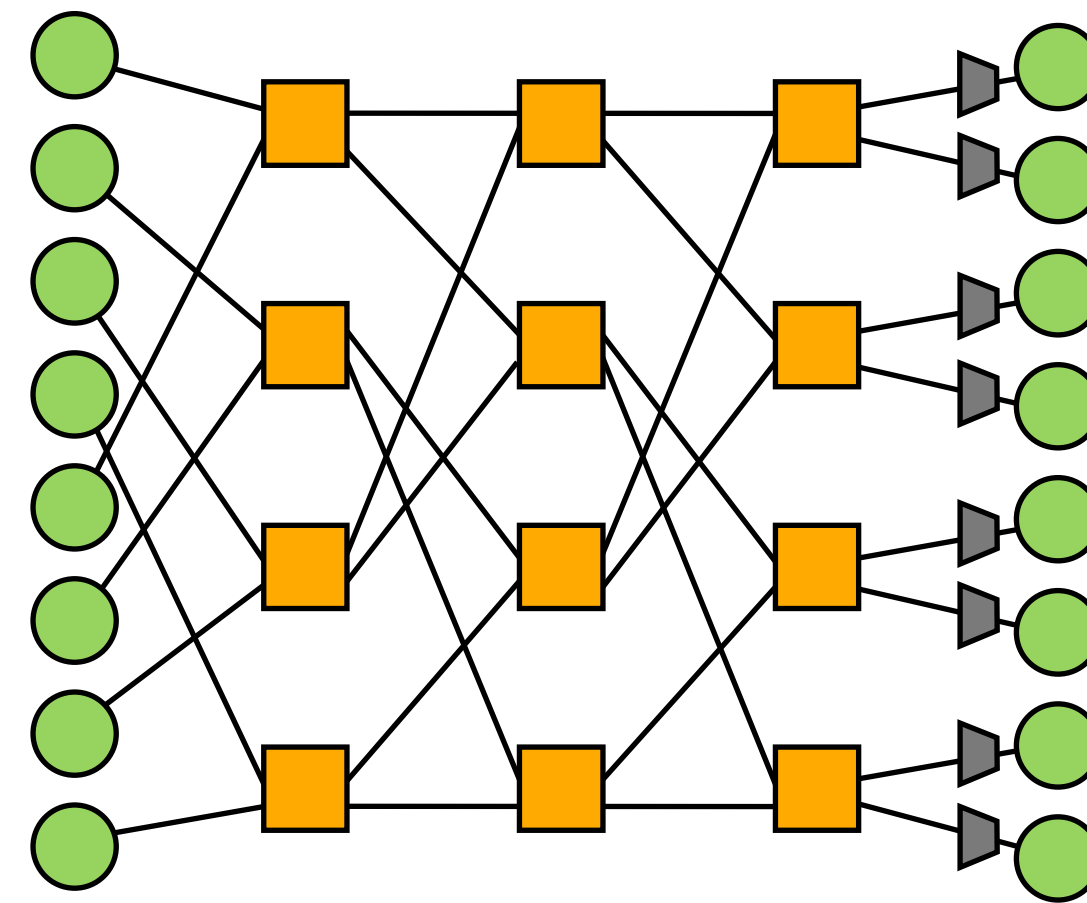


Properties of interconnect topology

- Direct vs. indirect networks
 - **Direct network: endpoints sit “inside” the network**
 - e.g., mesh is direct network: every node is both an endpoint and a switch



Direct network



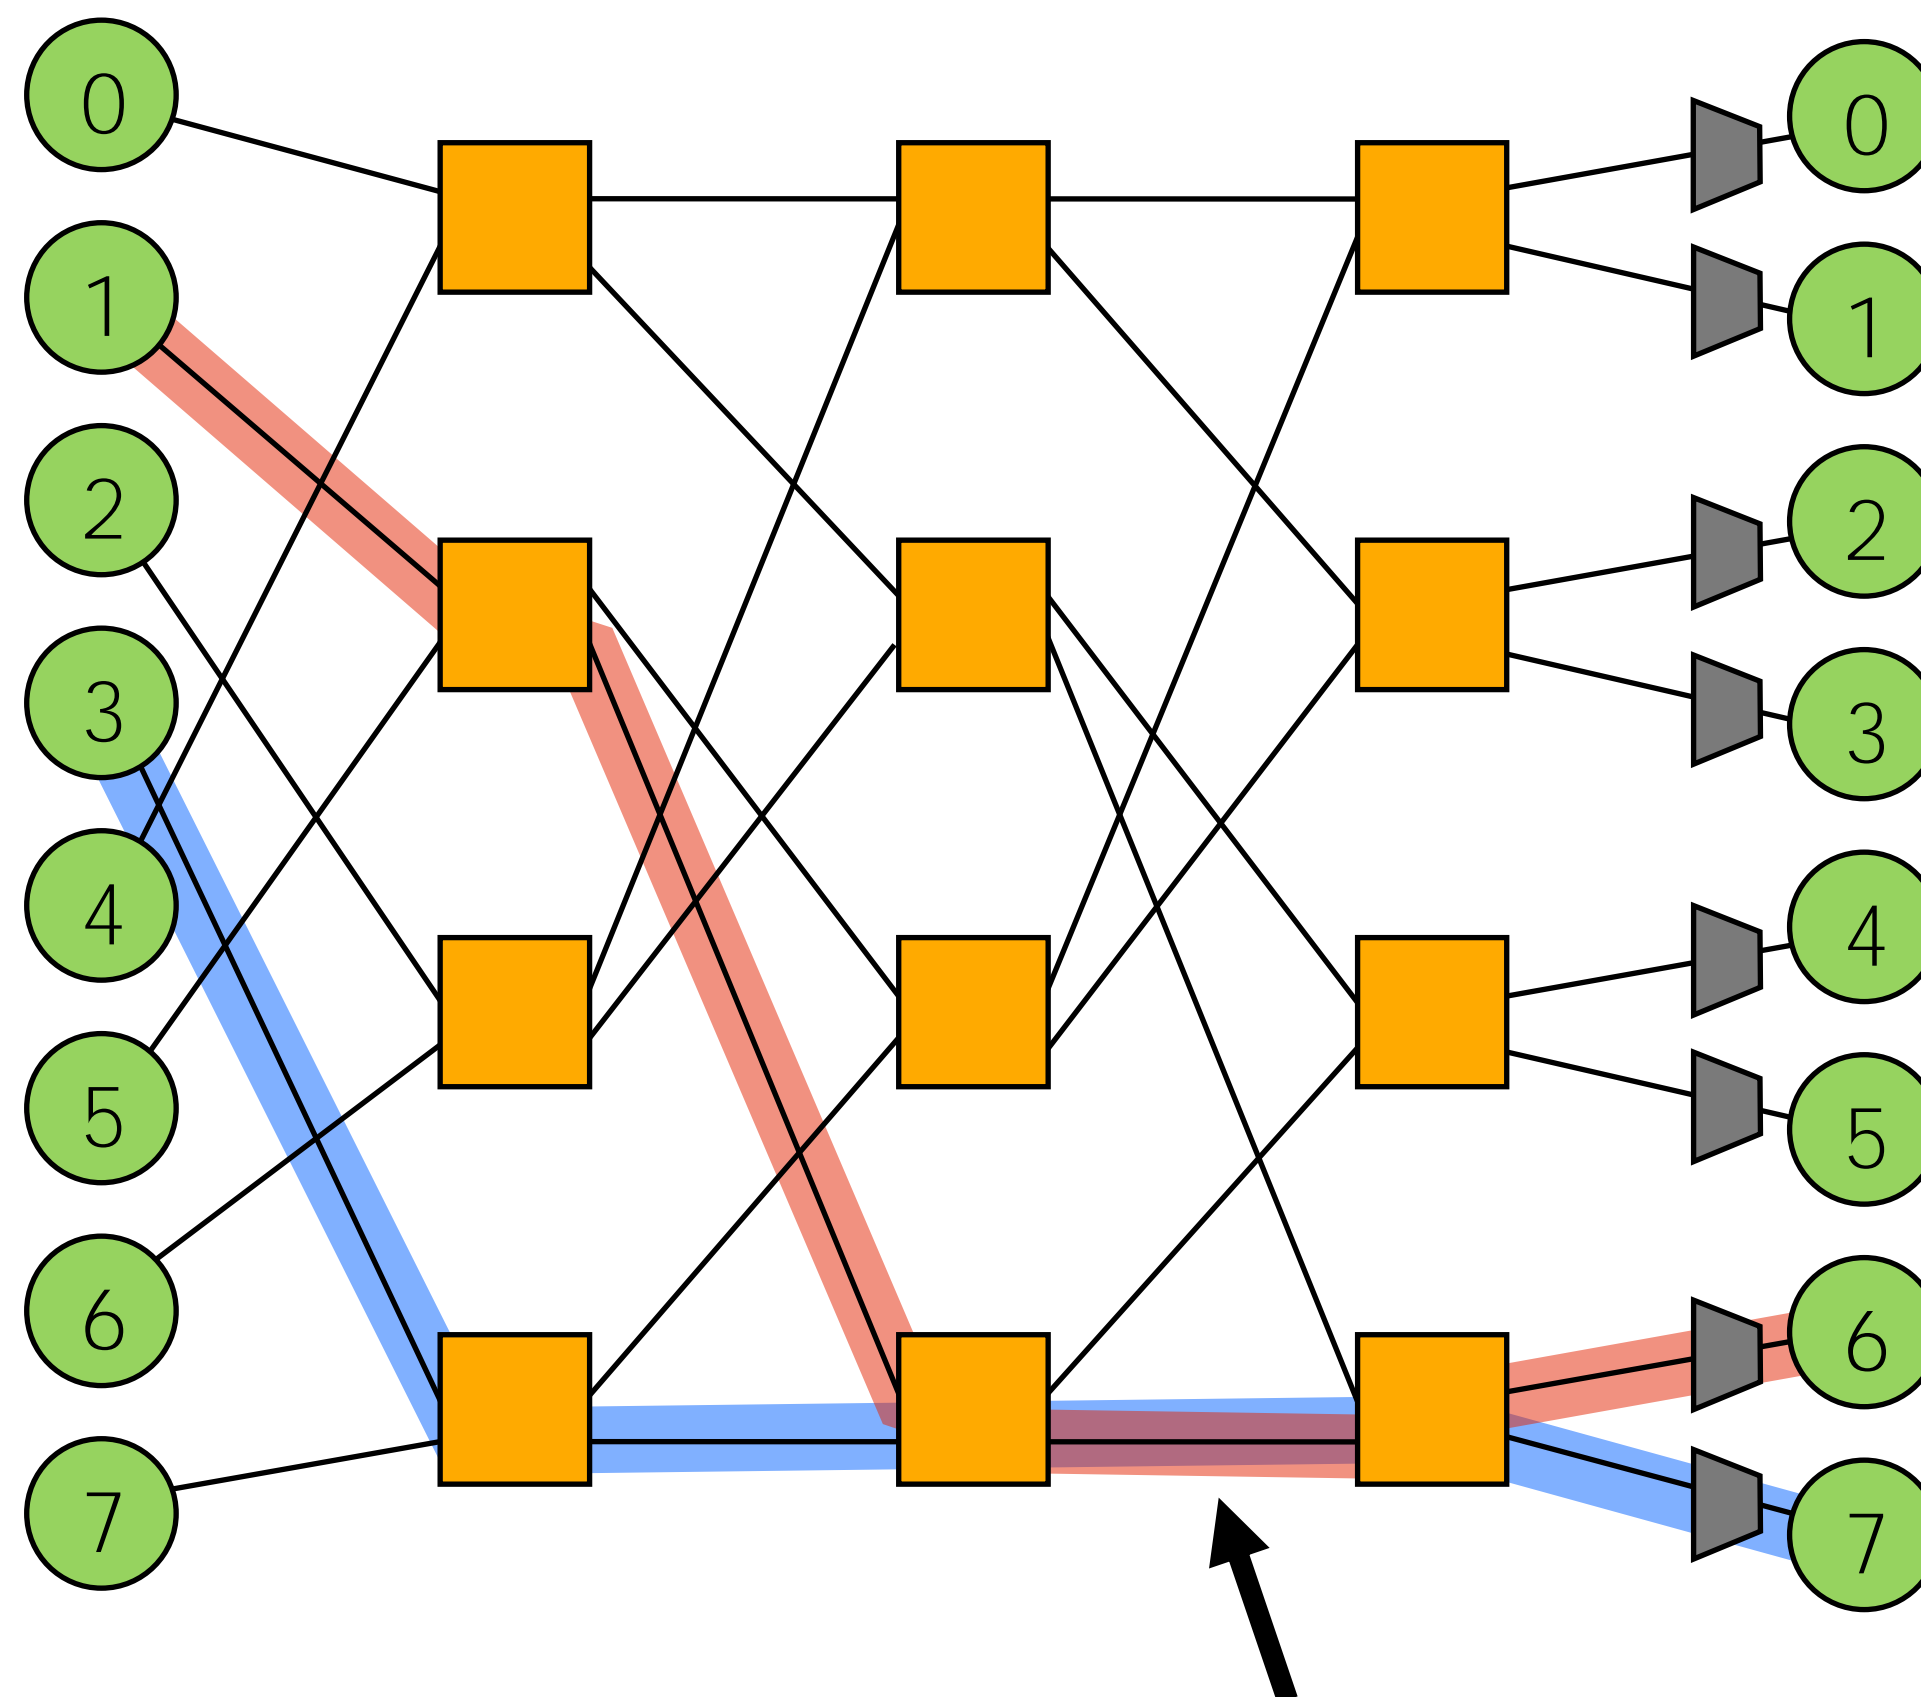
Indirect network

Properties of an interconnect topology

- Bisection bandwidth:
 - Common metric of performance for recursive topologies
 - Cut network in half, sum bandwidth of all severed links
 - Warning: can be misleading as it does not account for switch and routing efficiencies
- Blocking vs. non-blocking:
 - If connecting any pairing of nodes is possible, network is non-blocking (otherwise, **it's blocking**)

Example: blocking vs. non-blocking

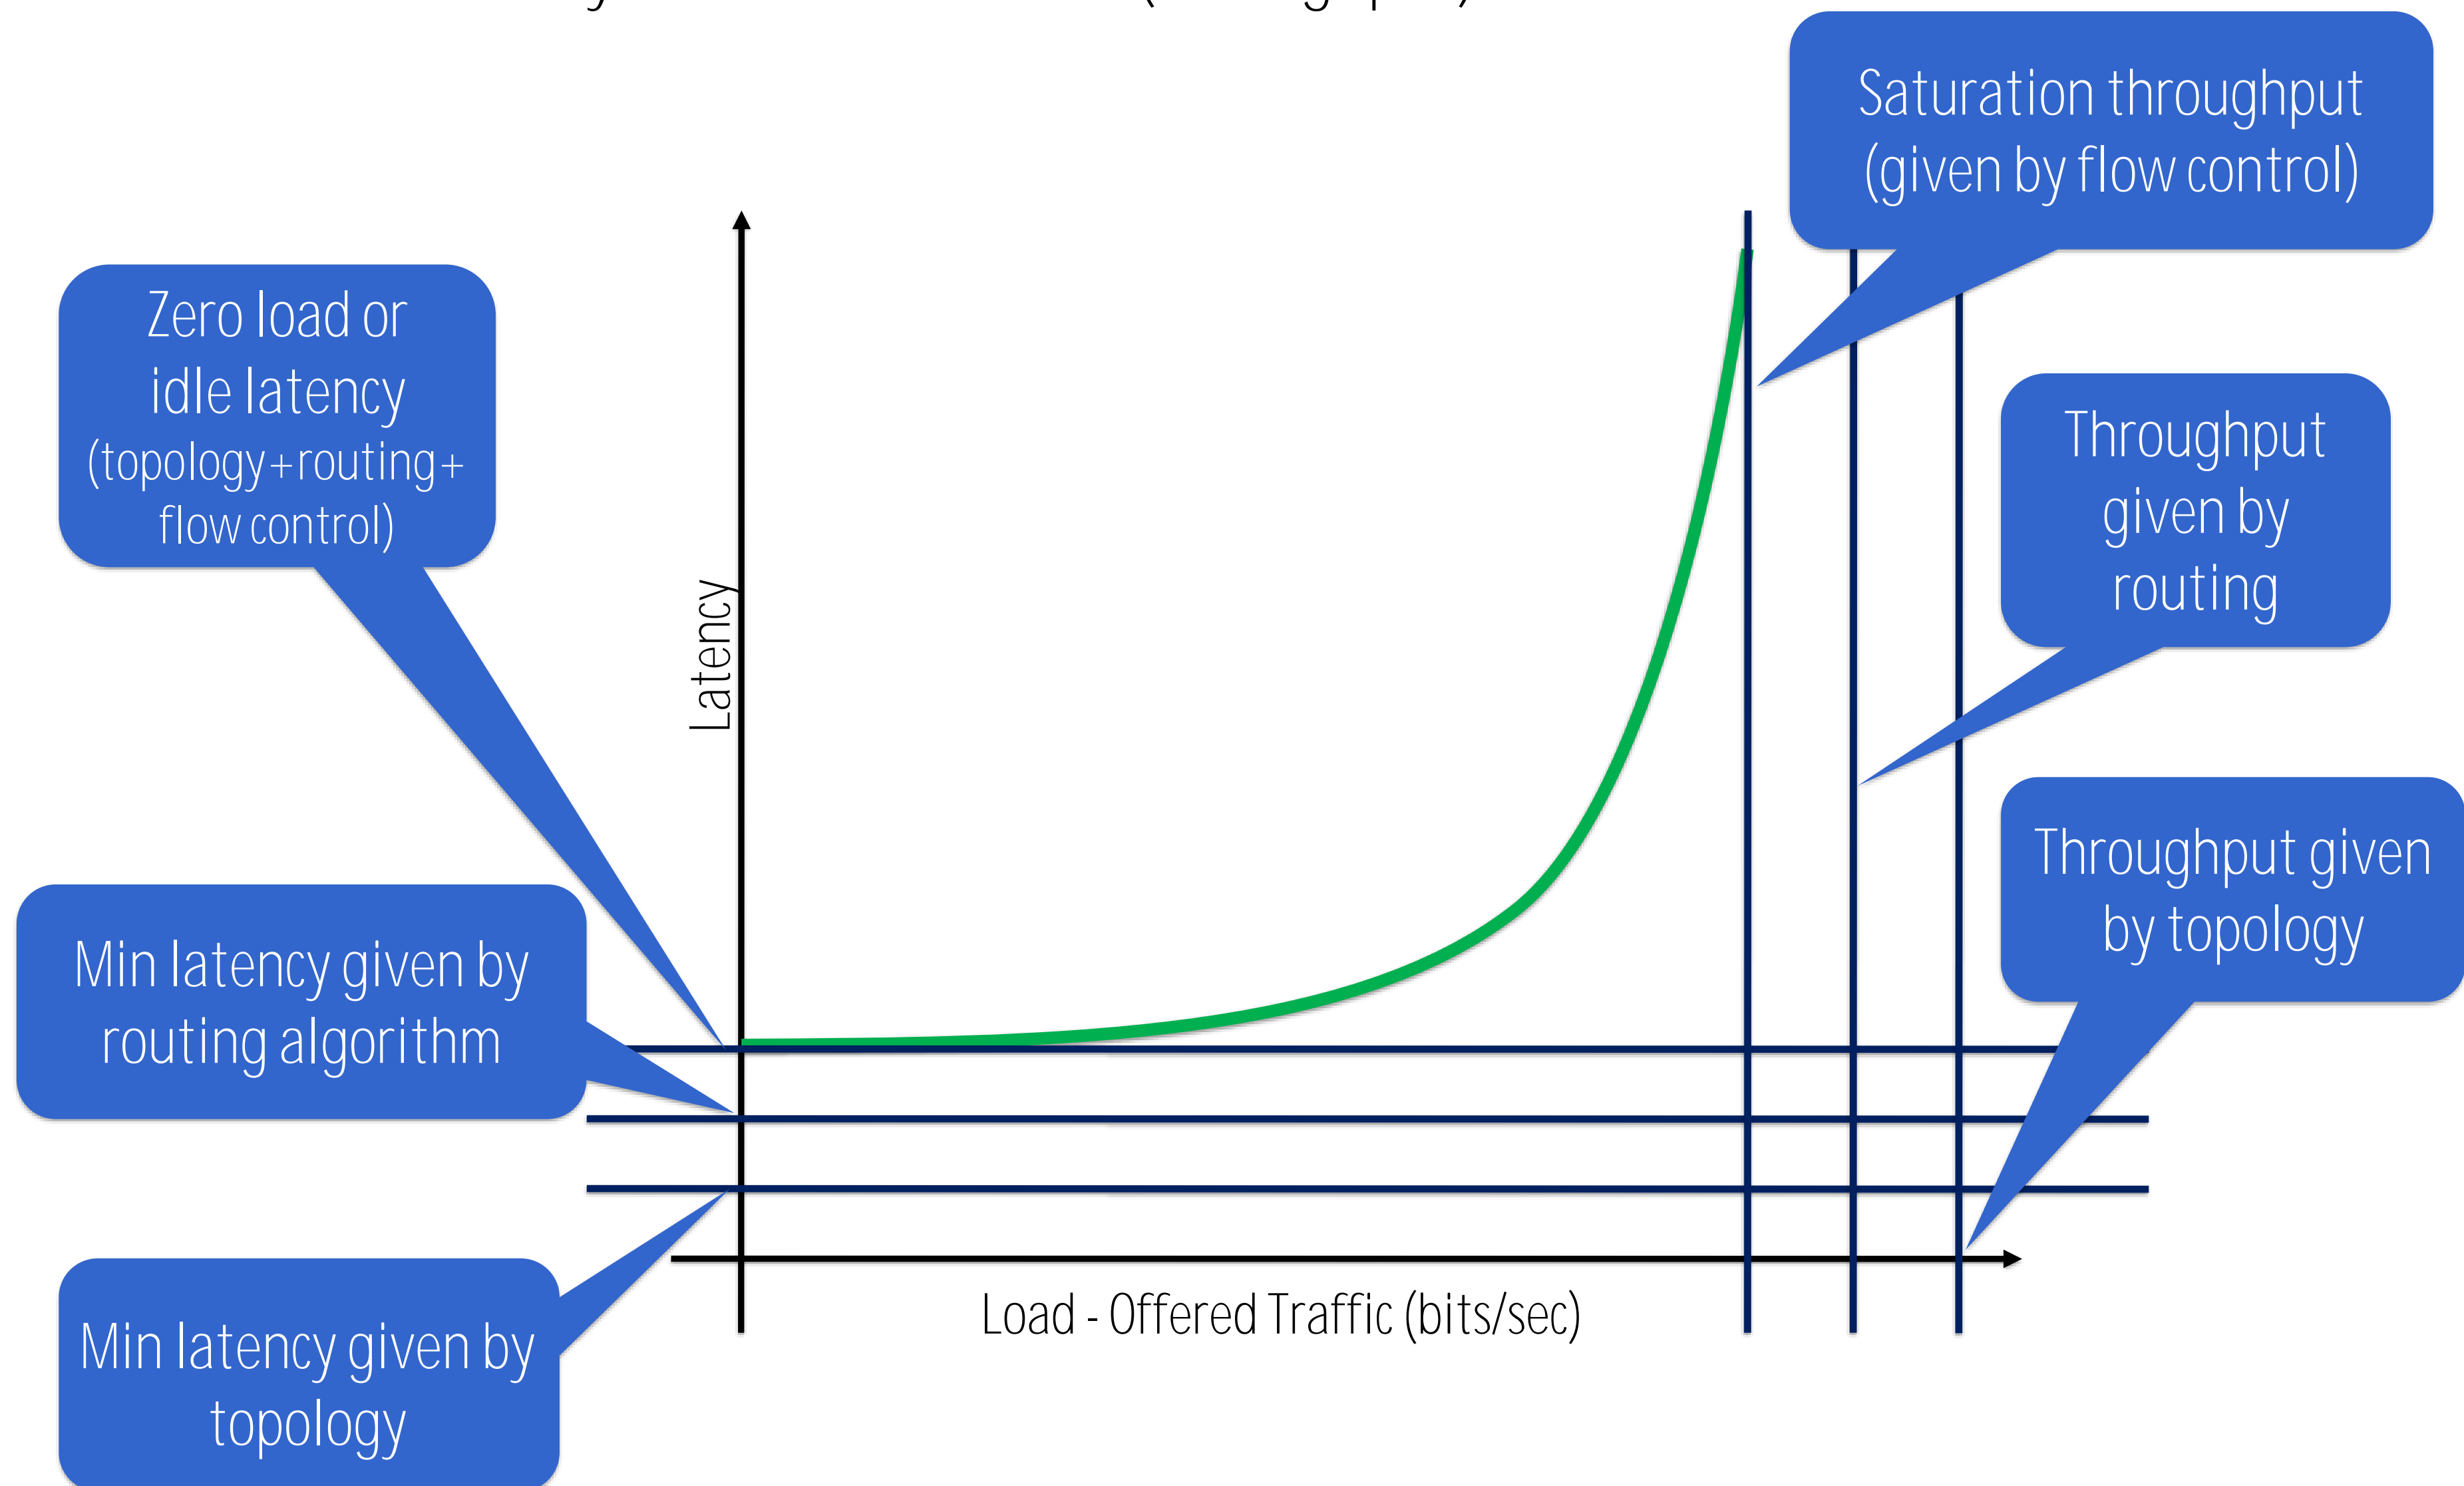
- Is this network blocking or non-blocking?
 - Consider simultaneous messages from 0-to-1 and 3-to-7.
 - Consider simultaneous messages from 1-to-6 and 3-to-7. Blocking!!!



Note: in this network illustration, each node is drawn twice for clarity (at left and at right)

Load-latency behavior of network

General rule: latency increases with load (throughput)



Interconnect topologies

Many possible network topologies

Bus

Crossbar

Ring

Tree

Omega

Hypercube

Mesh

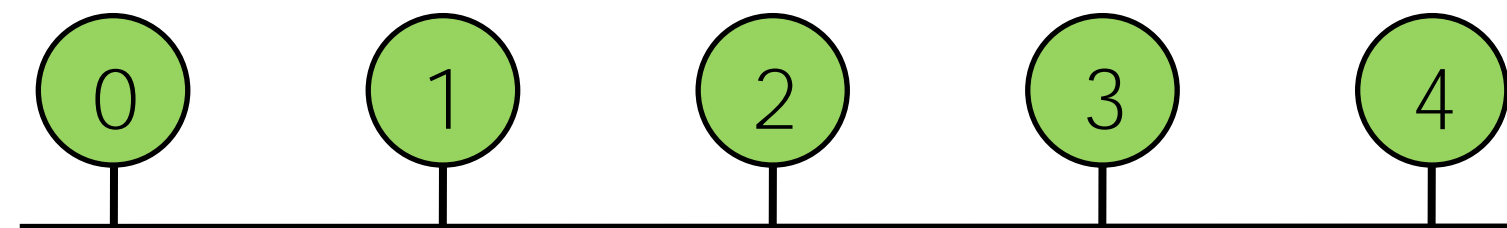
Torus

Butterfly

...

Bus interconnect

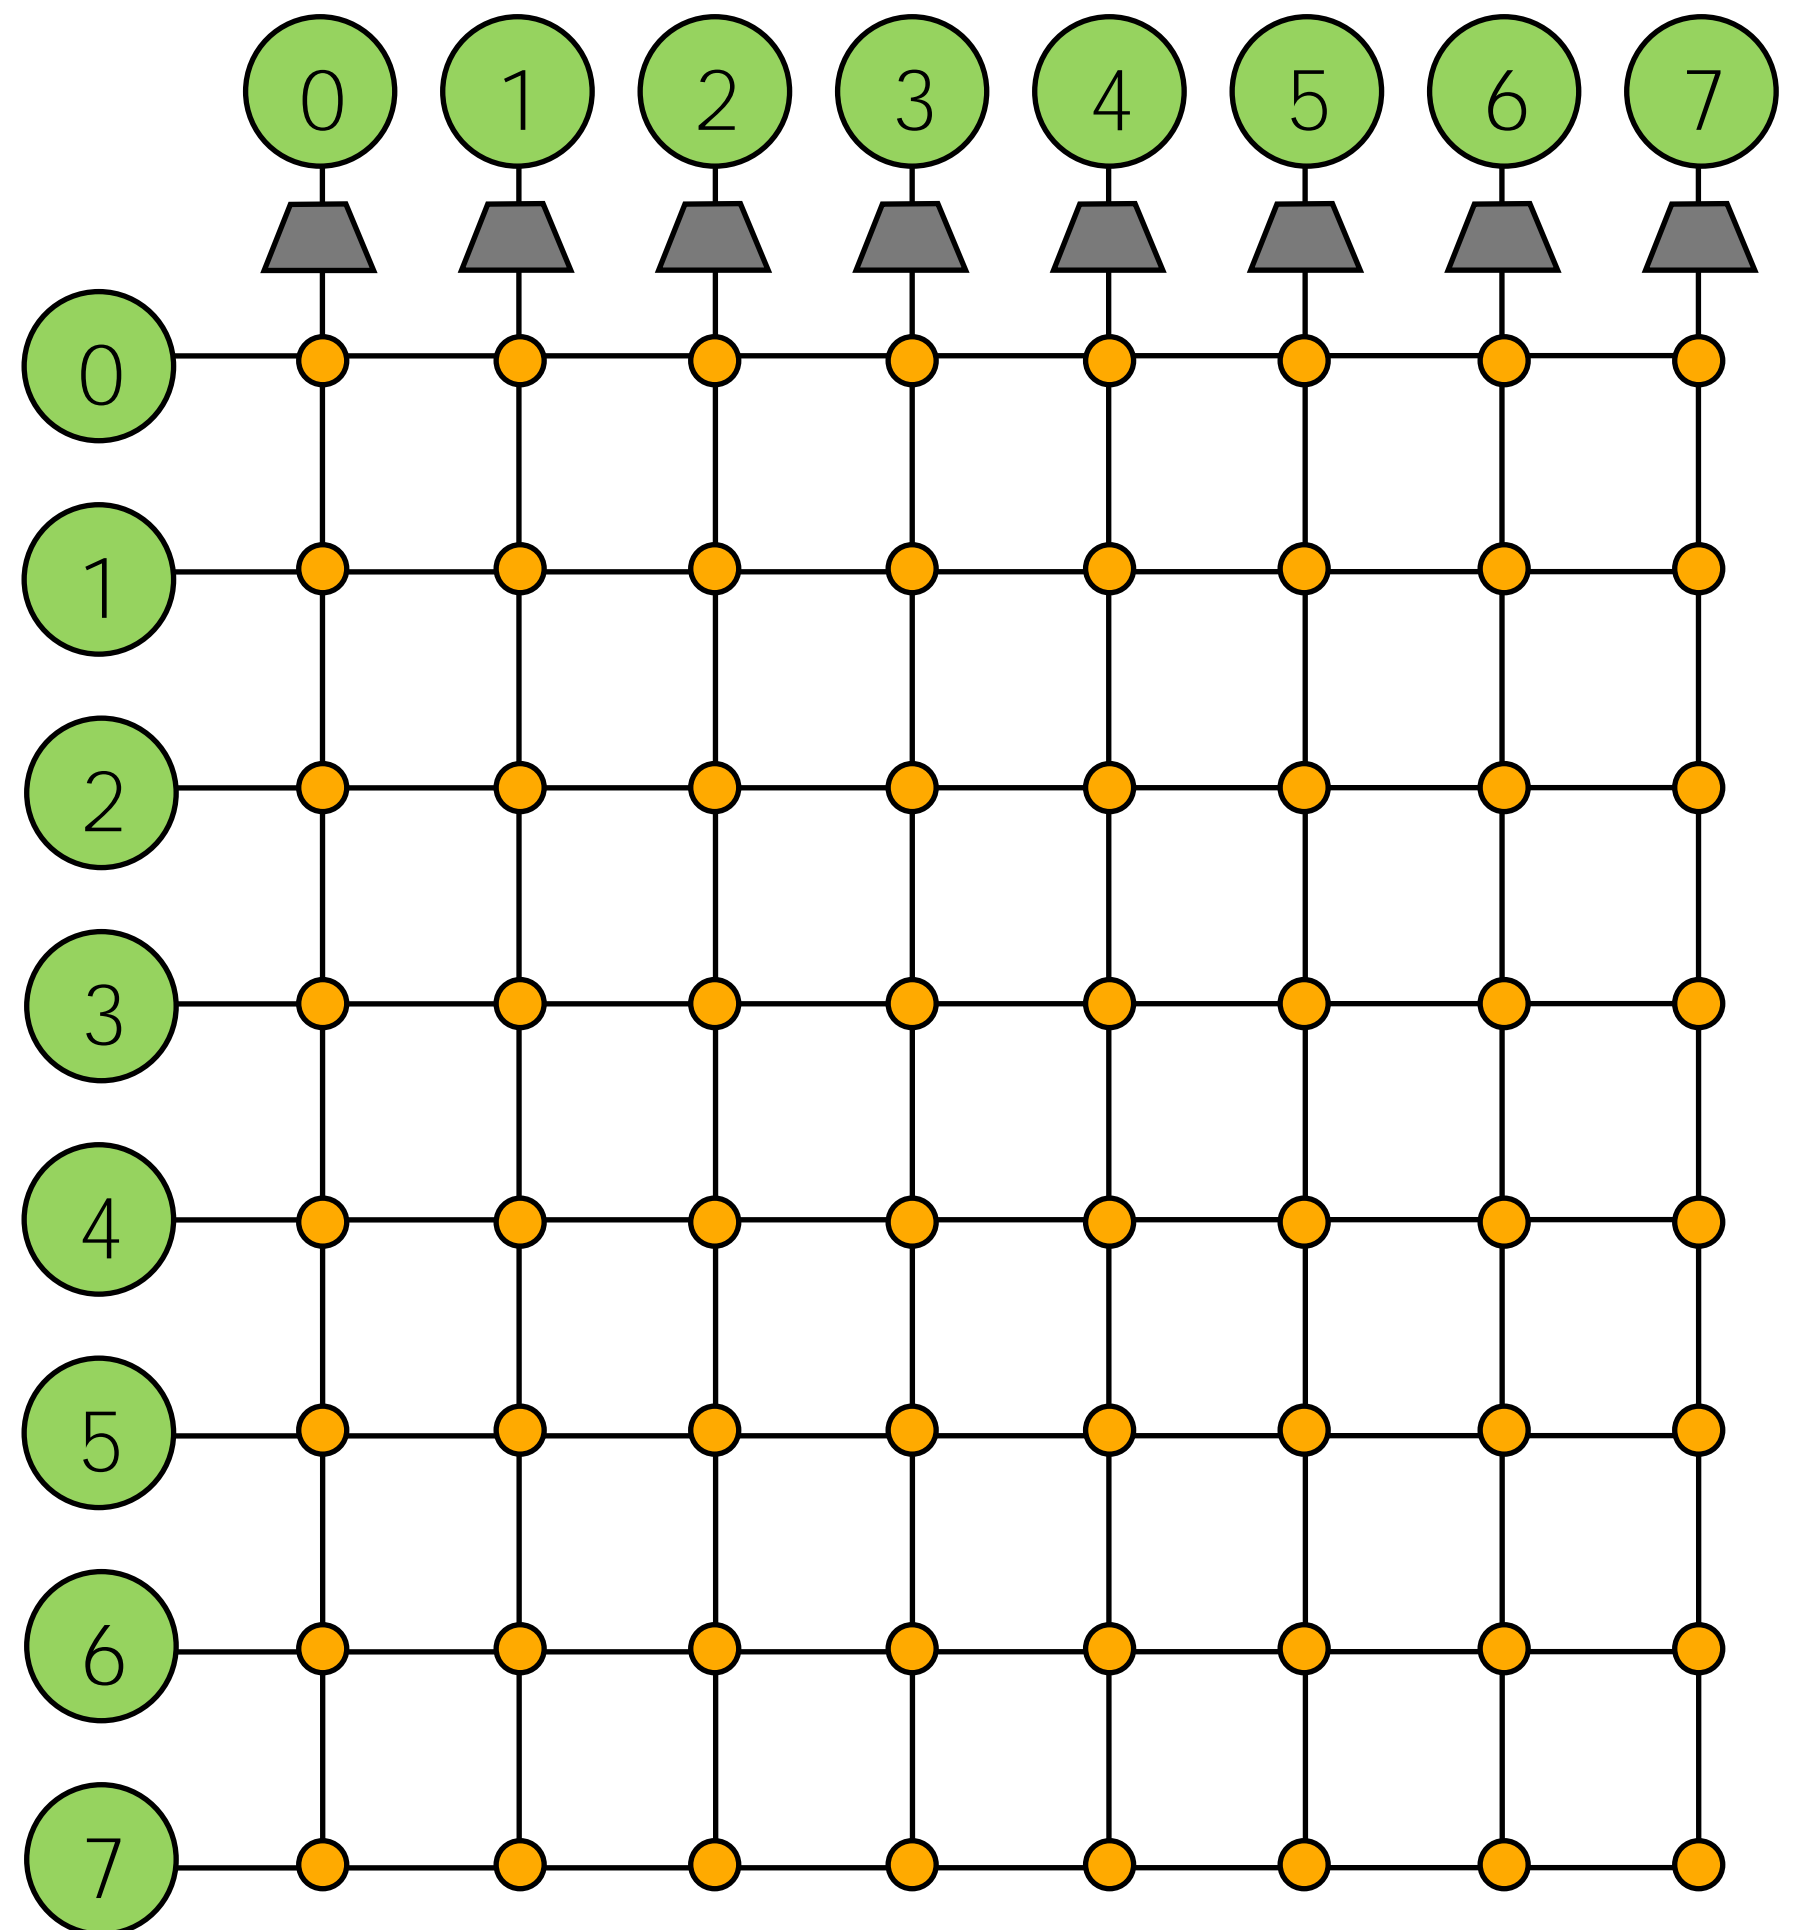
- Good:
 - Simple design
 - Cost effective for a small number of nodes
 - Easy to implement coherence (via snooping)
- Bad:
 - Contention: all nodes contend for shared bus
 - Limited bandwidth: all nodes communicate over same wires (one communication at a time)
 - High electrical load = low frequency, high power



Crossbar interconnect

- Every node is connected to every other node (non-blocking, indirect)
- Good:
 - $O(1)$ latency and high bandwidth
- Bad:
 - Not scalable: $O(N^2)$ switches
 - High cost
 - Difficult to arbitrate at scale

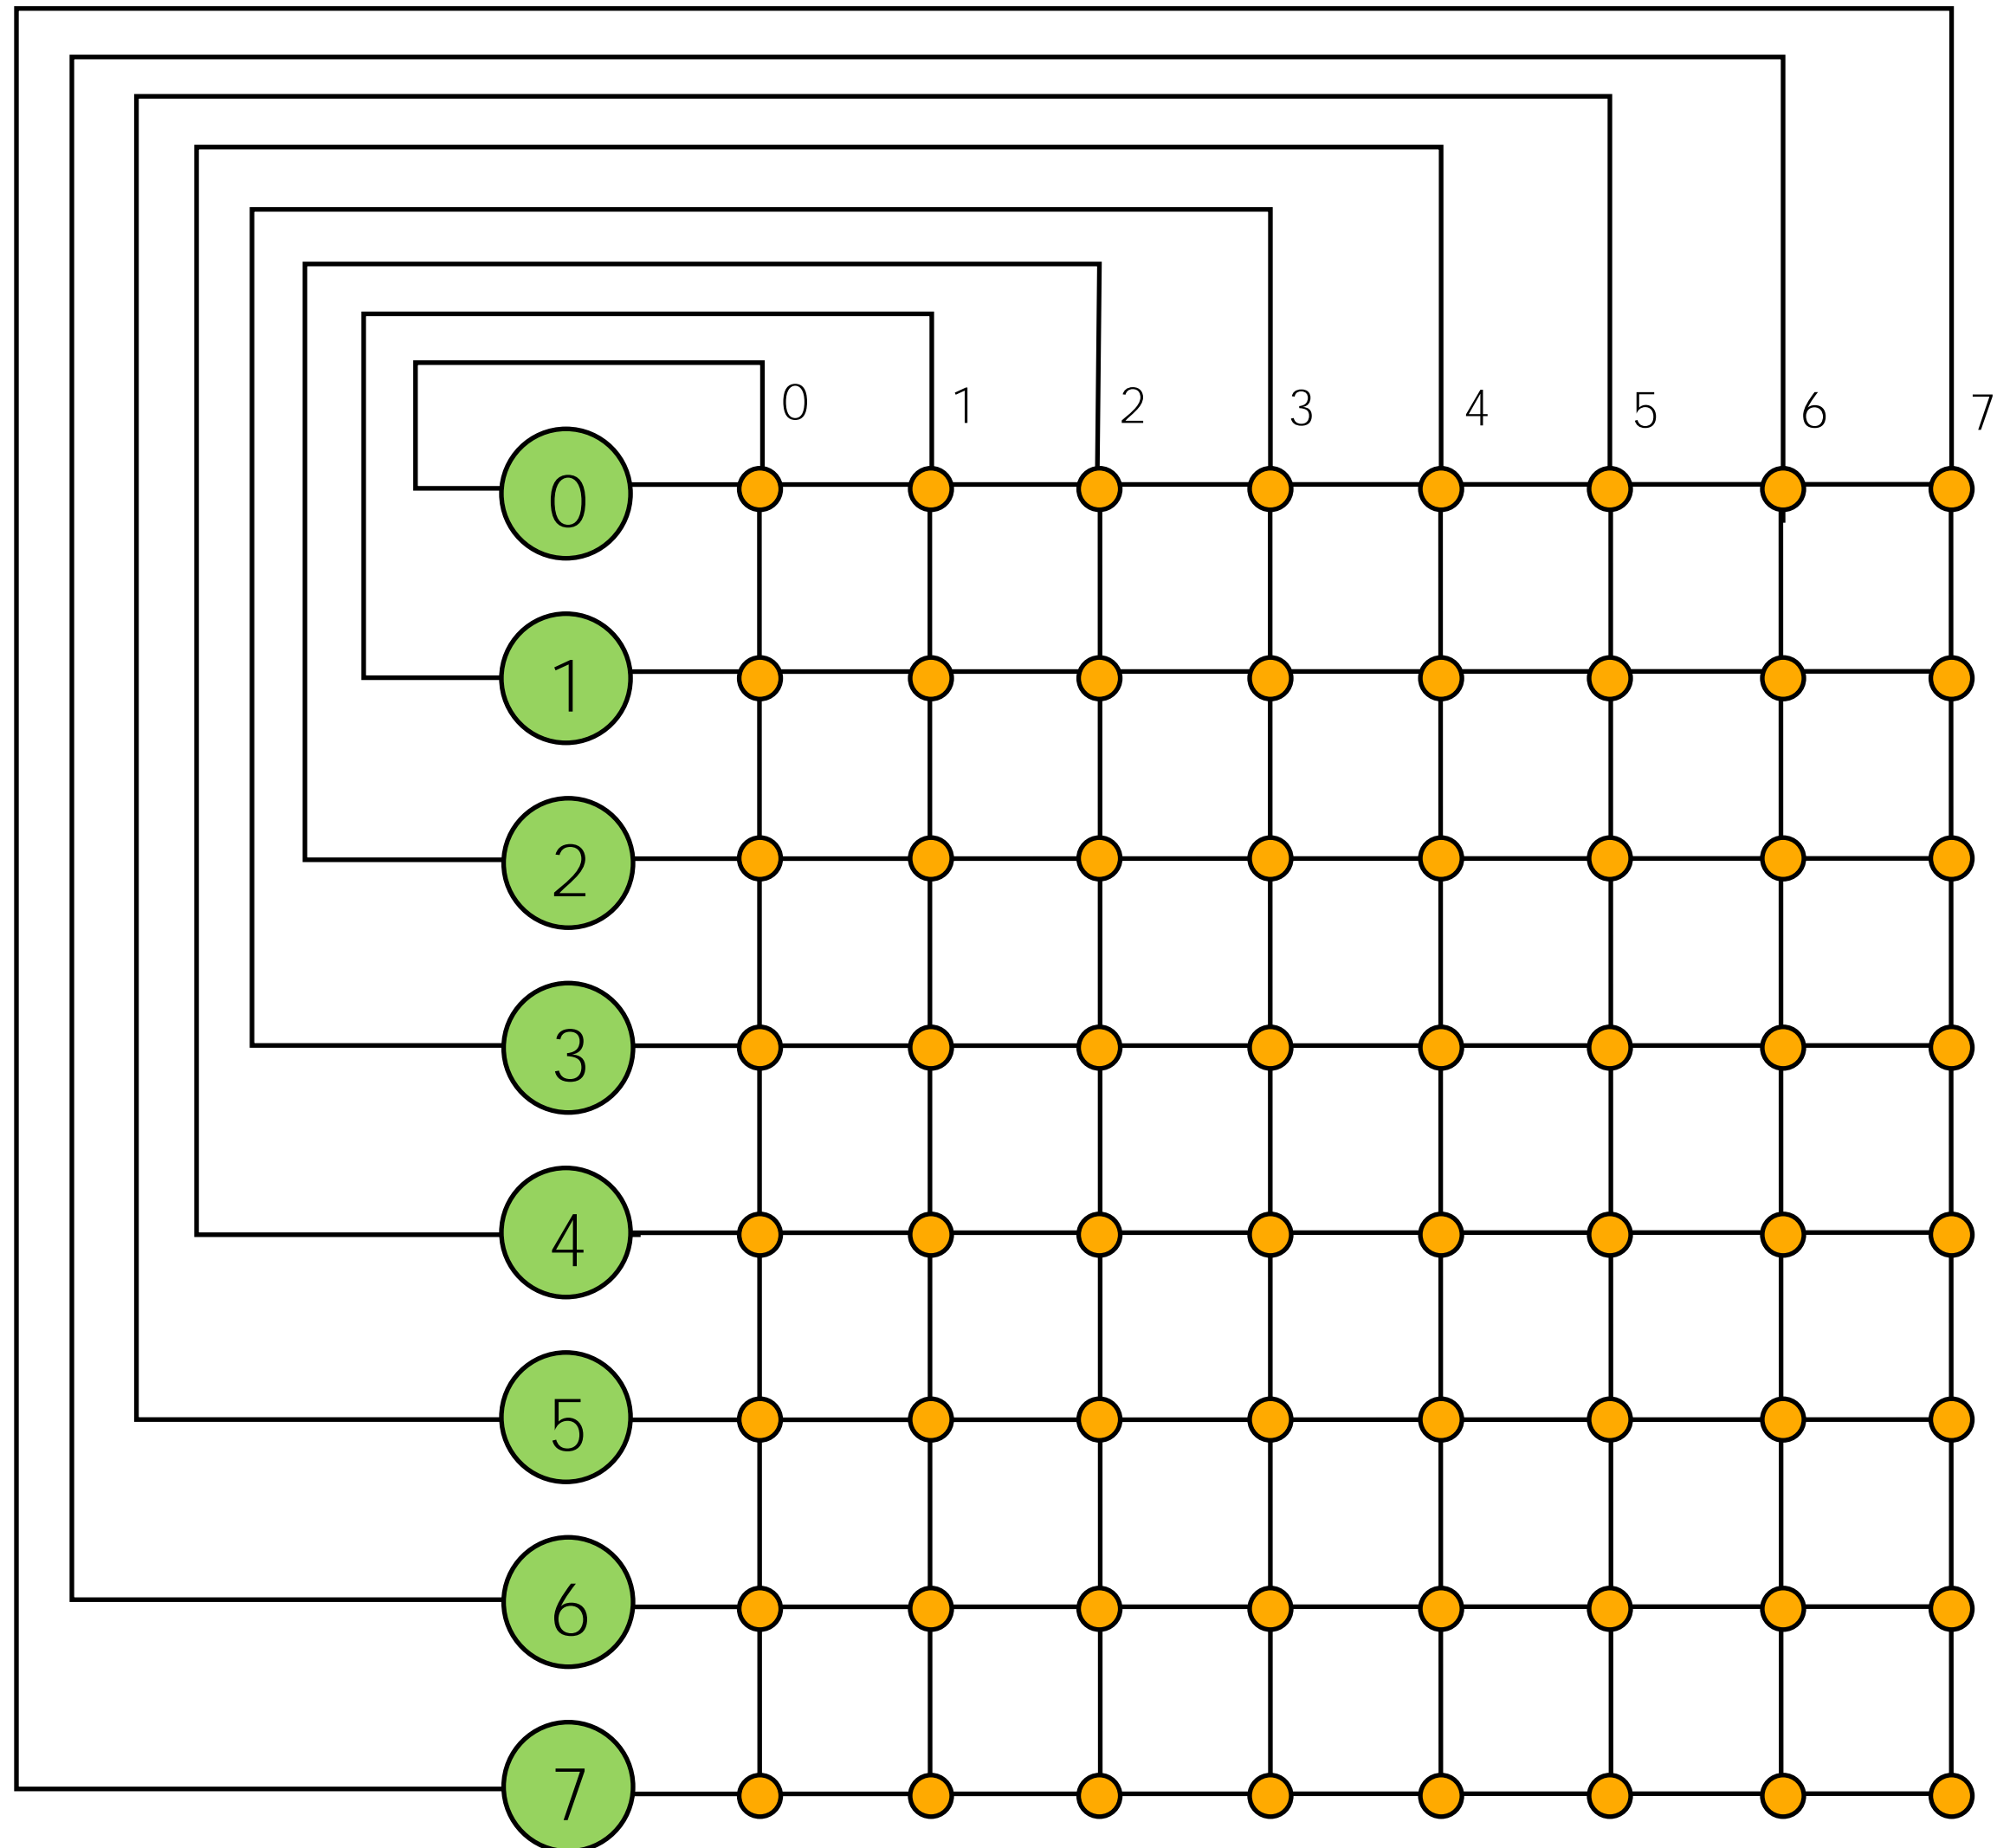
↗
Crossbar scheduling algorithms / efficient hardware implementations are still active research areas.



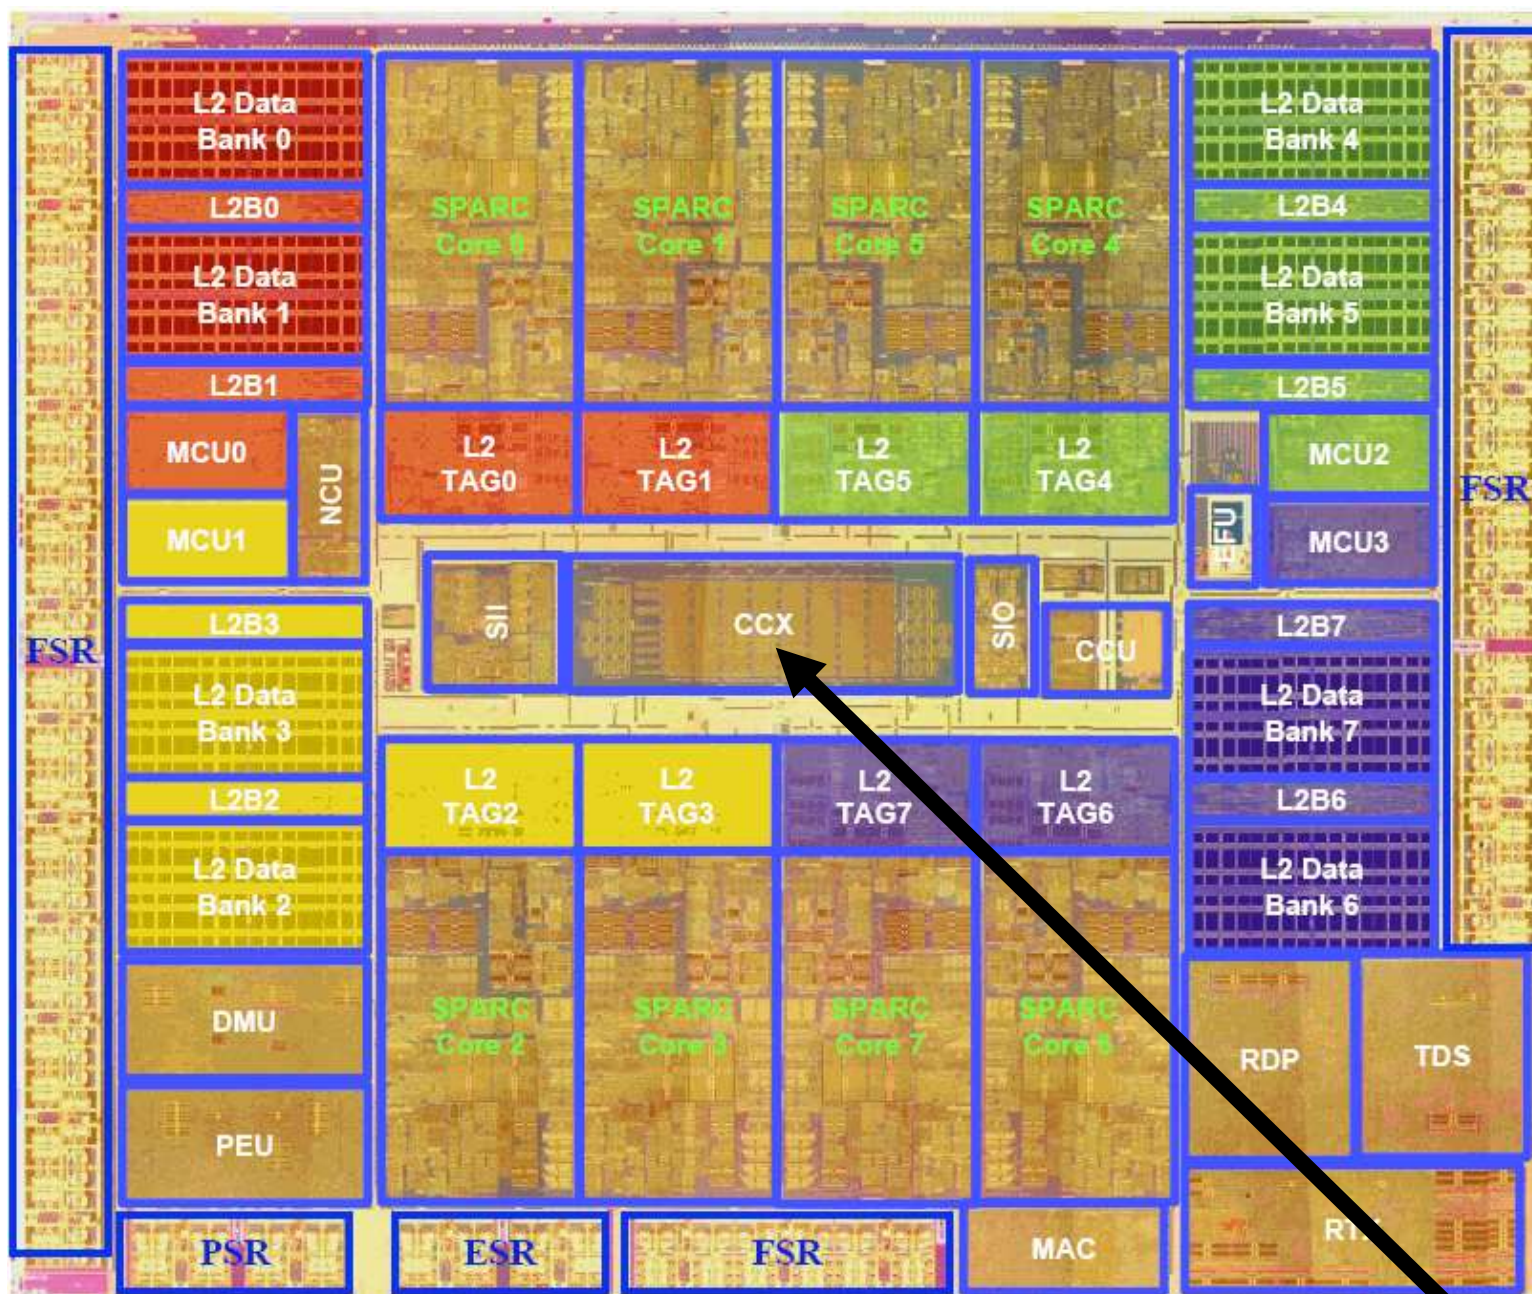
8-node crossbar network (N=8)

Crossbar interconnect

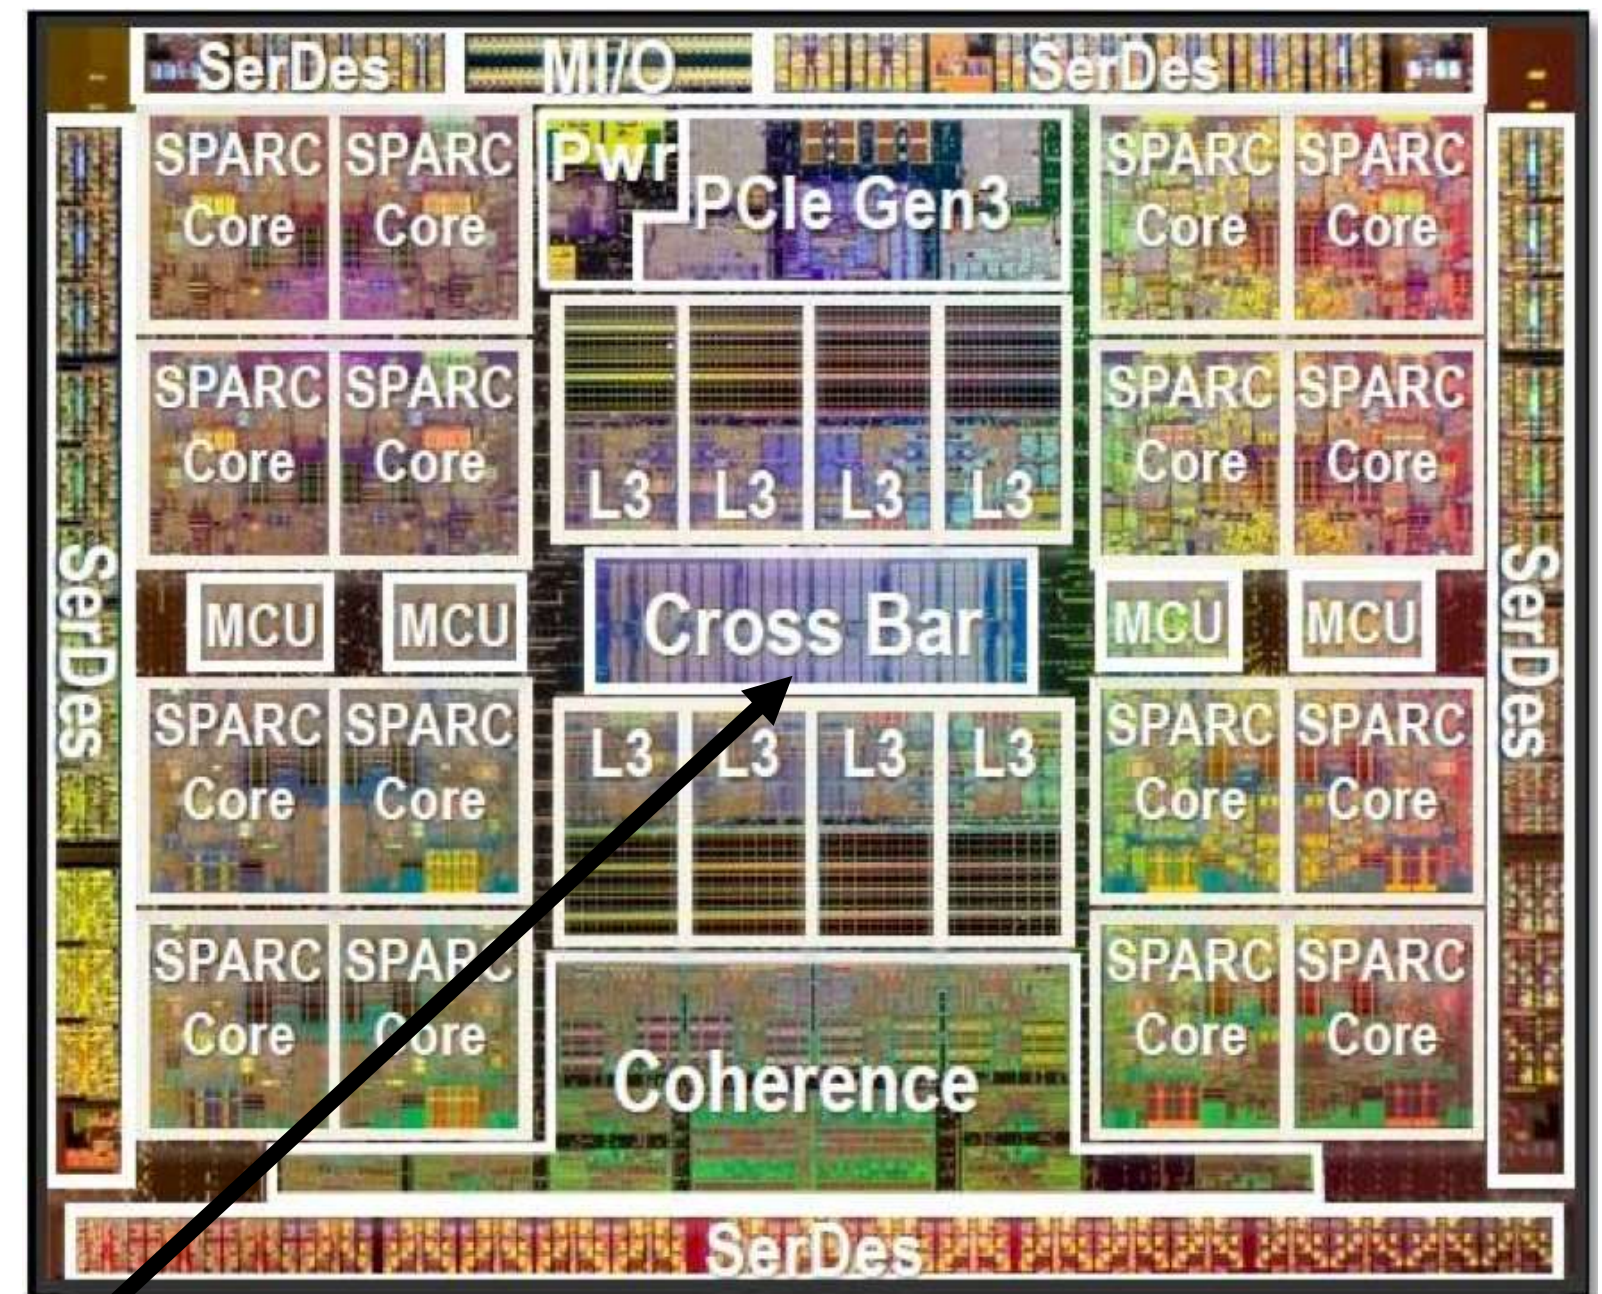
(Here is a more verbose illustration than that on previous slide)



Crossbars were used in recent multi-core processing from Oracle (previously Sun)



Sun SPARC T2 (8 cores, 8 L2 cache banks)

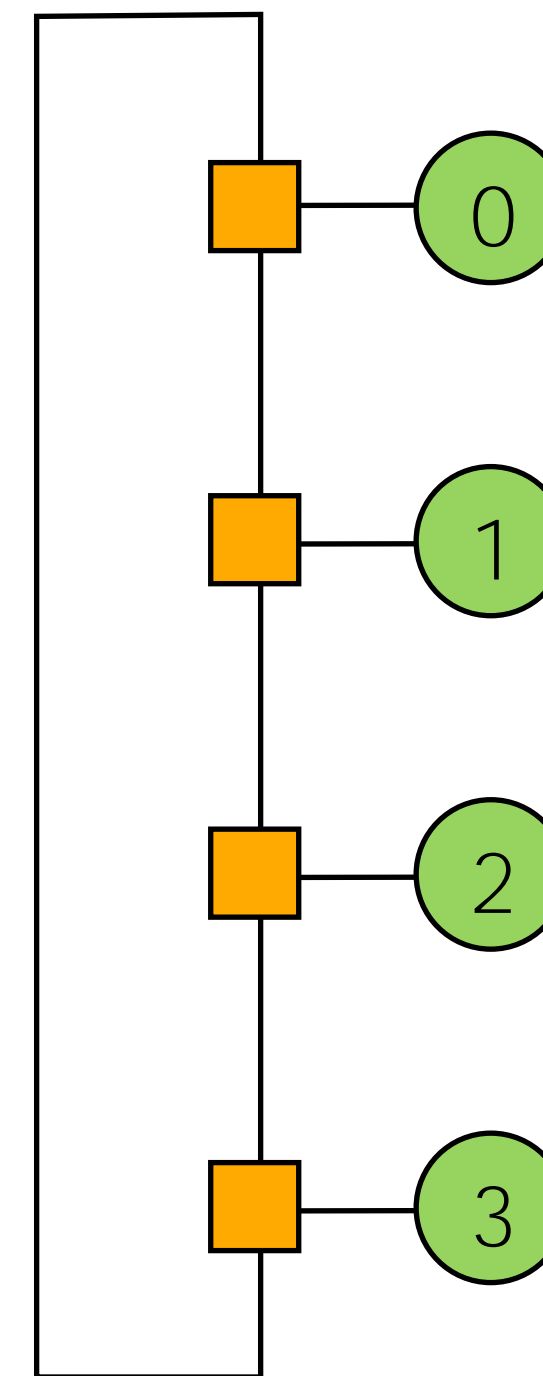


Oracle SPARC T5 (16 cores, 8 L3 cache banks)

Note that crossbar (CCX) occupies about the same chip area as a core

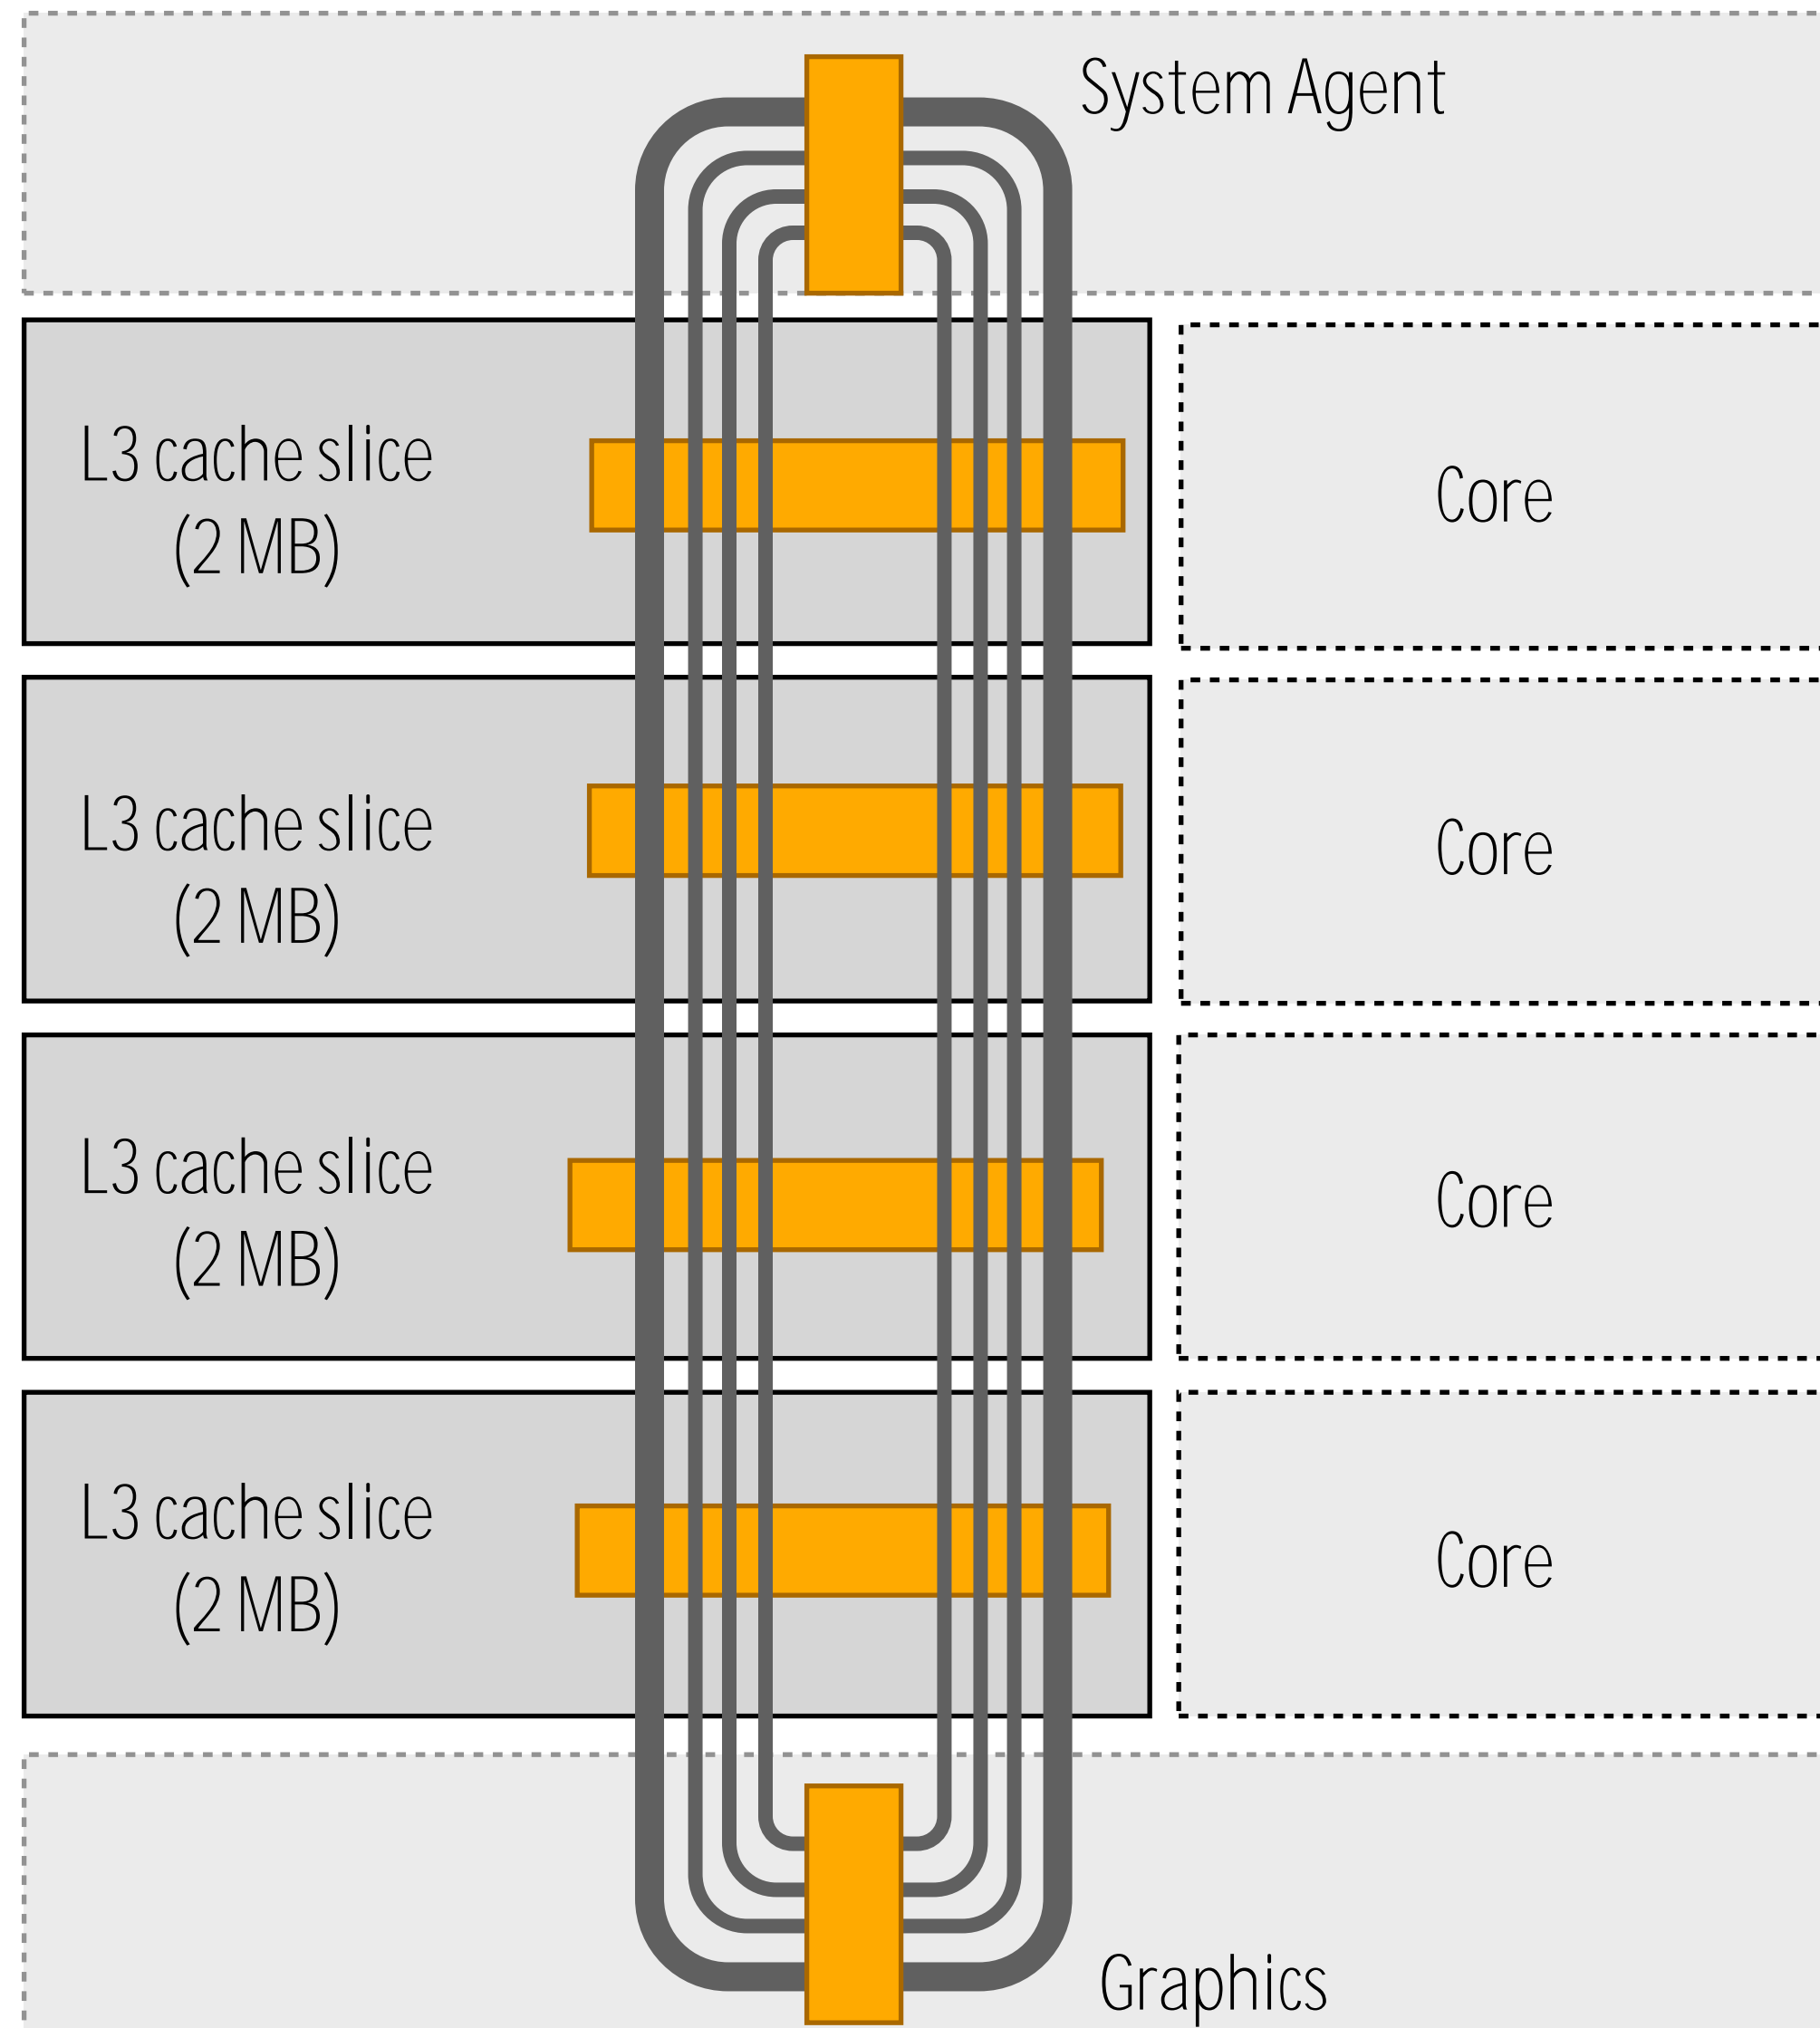
Ring

- Good:
 - Simple
 - Cheap: $O(N)$ cost
- Bad:
 - High latency: $O(N)$
 - Bisection bandwidth remains constant as nodes are added (scalability issue)
- Used in recent Intel architectures
 - Core i7
- Also used in IBM CELL Broadband Engine (9 cores)



Intel's ring interconnect

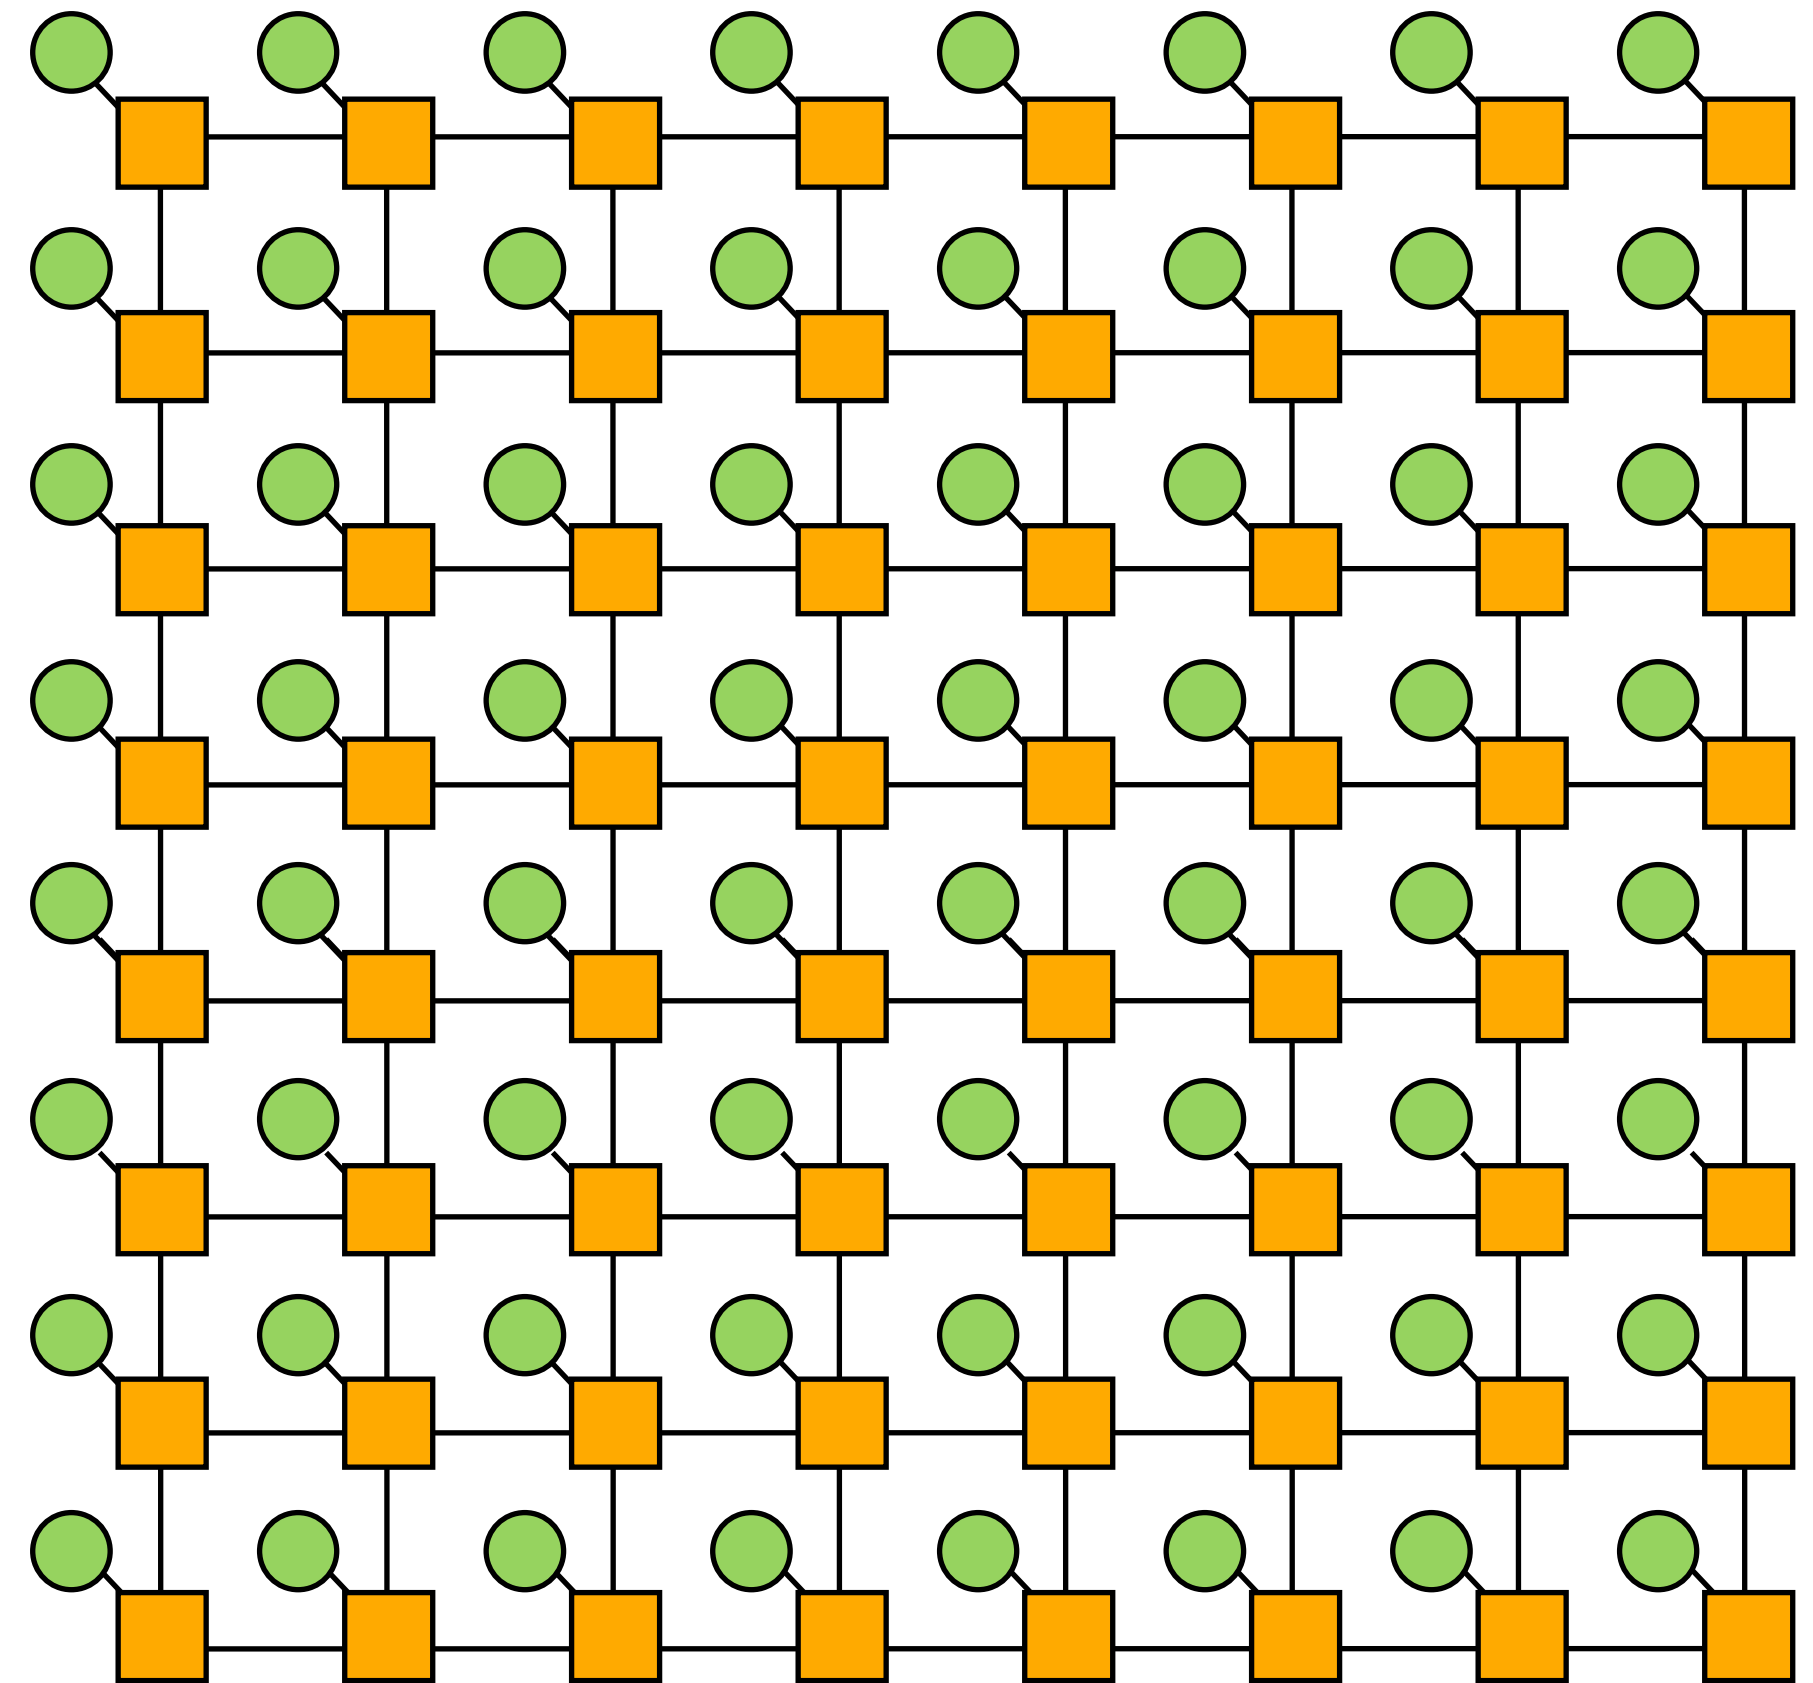
Introduced in Sandy Bridge microarchitecture



- Four rings
 - request
 - snoop
 - Ack
 - data (32 bytes)
- Six interconnect nodes: four **“slices” of L3 cache + system agent + graphics**
- Each bank of L3 connected to ring bus twice
- Theoretical peak BW from cores to L3 at 3.4 GHz is approx. 435 GB/sec
 - When each core is accessing its local slice

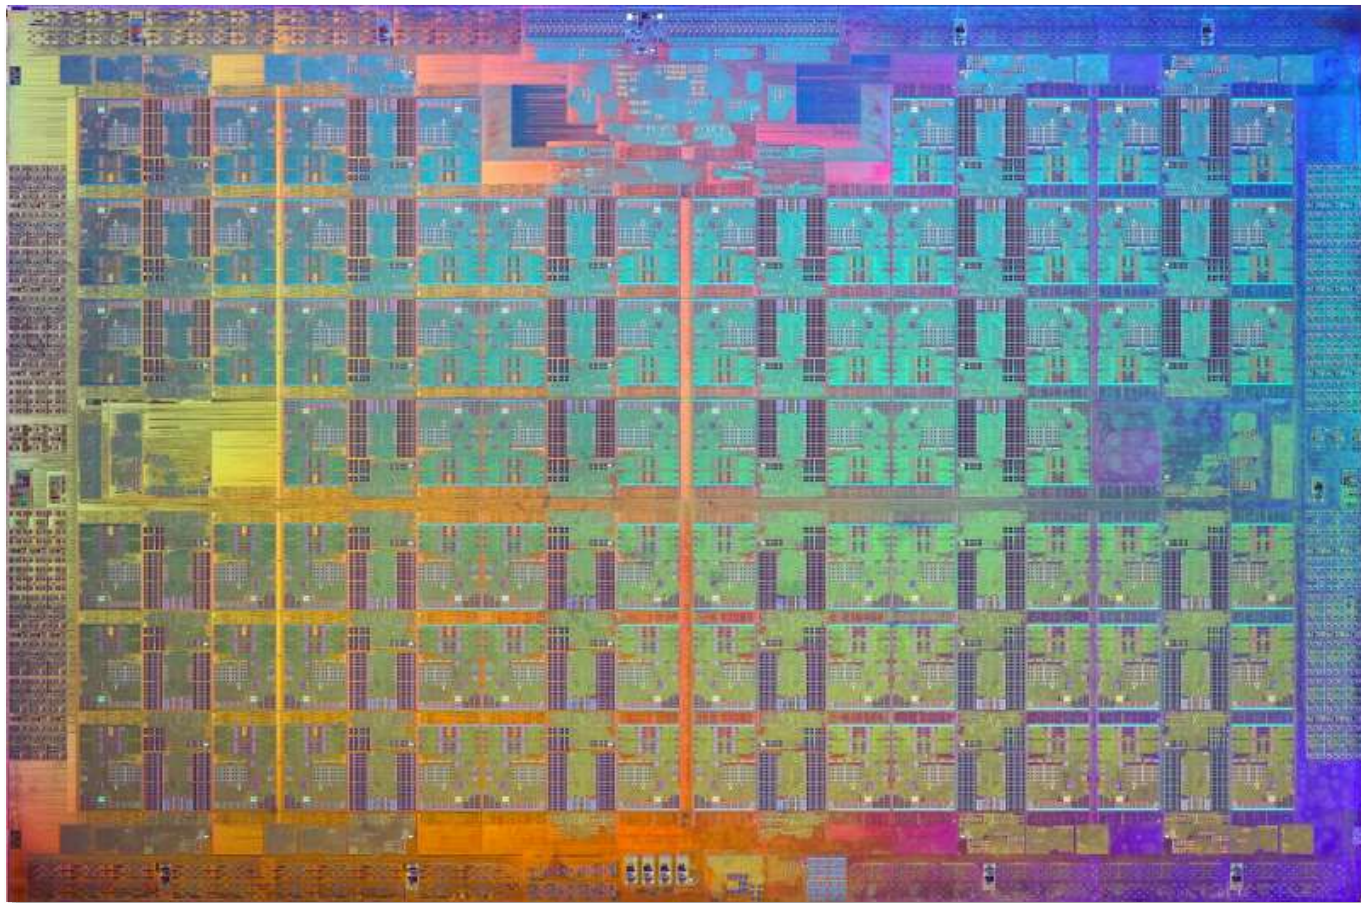
Mesh

- Direct network
- Echoes locality in grid-based applications
- $O(N)$ cost
- Average latency: $O(\sqrt{N})$
- Easy to lay out on chip: fixed-length links
- Path diversity: many ways for message to travel from one node to another
- Used by:
 - Tiler processors
 - Prototype Intel chips

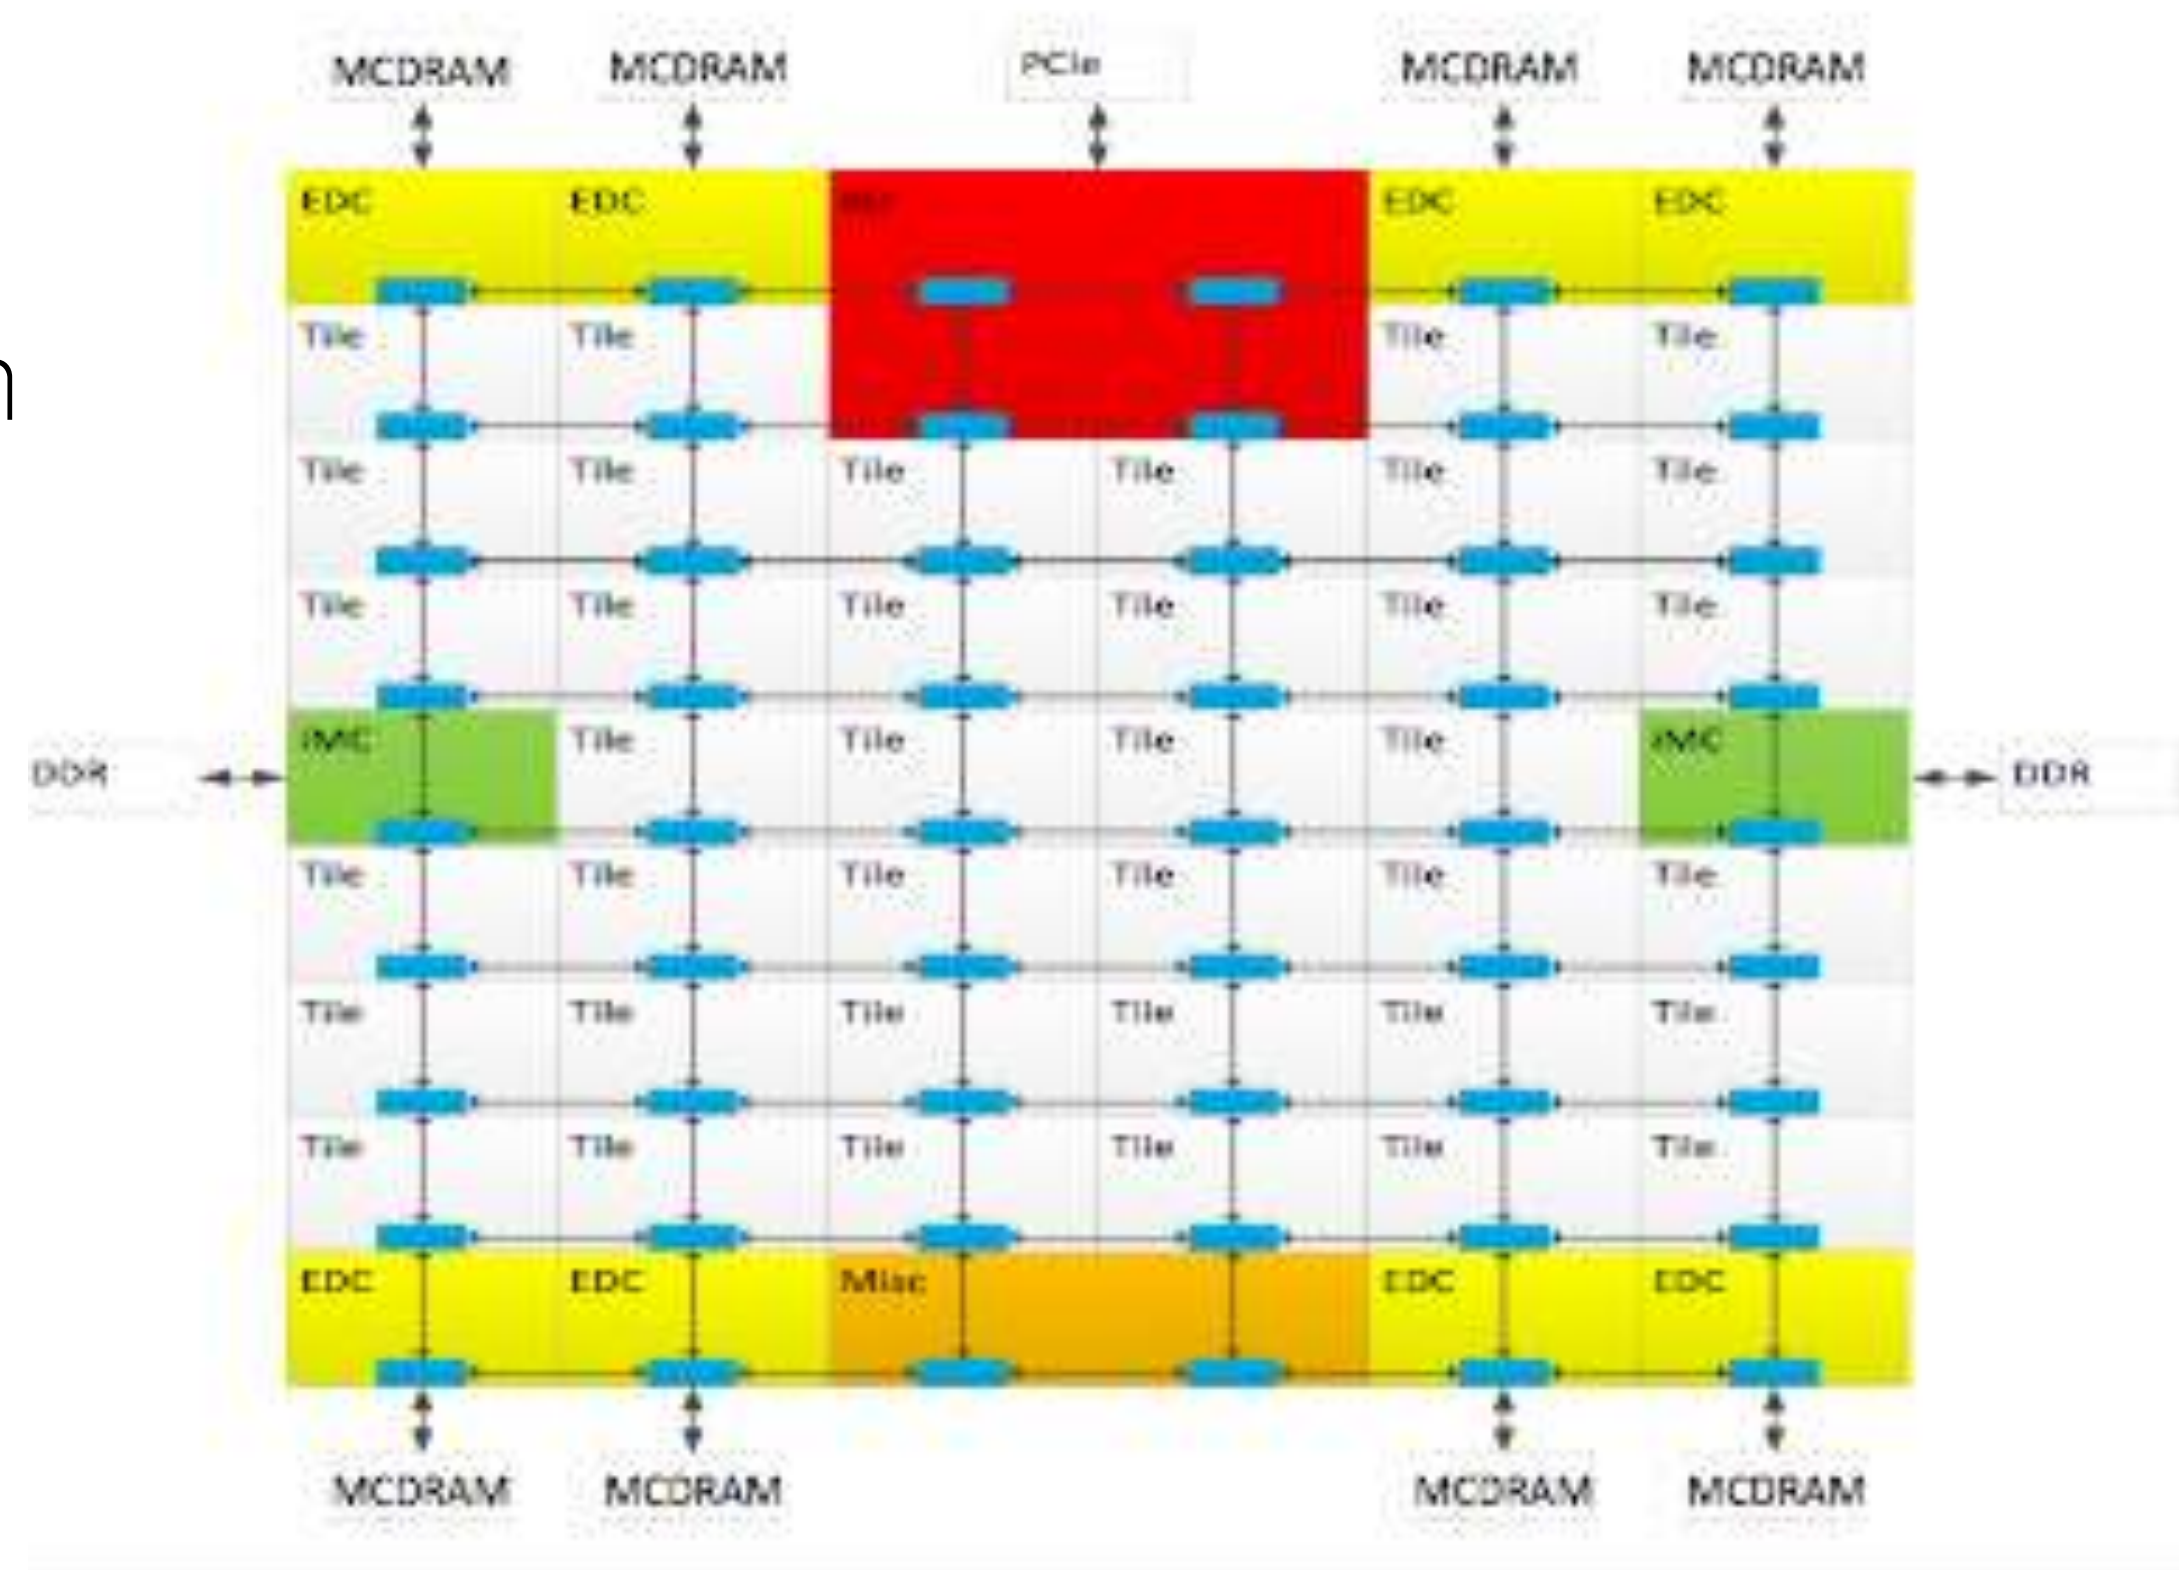


2D Mesh

Xeon Phi (Knights Landing)

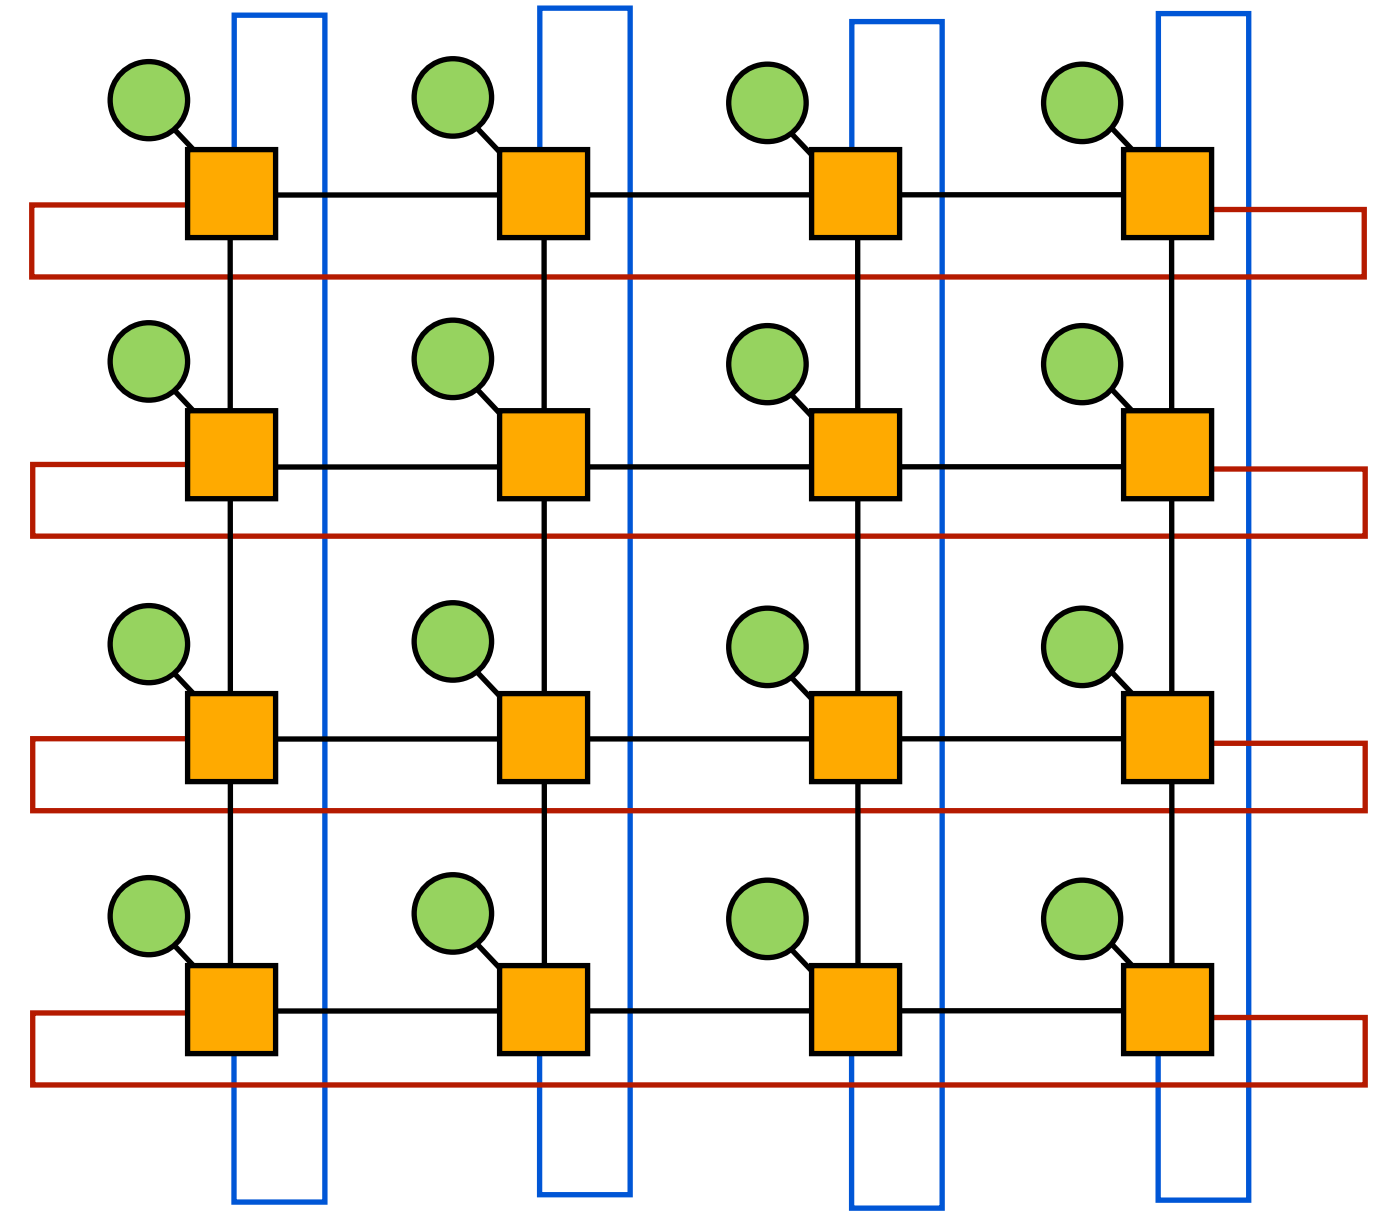


- 72 cores, arranged as 6 x 6 mesh of tiles (2 cores/tile)
- YX routing of messages:
 - Move in Y
 - **“Turn”**
 - Move in X



Torus

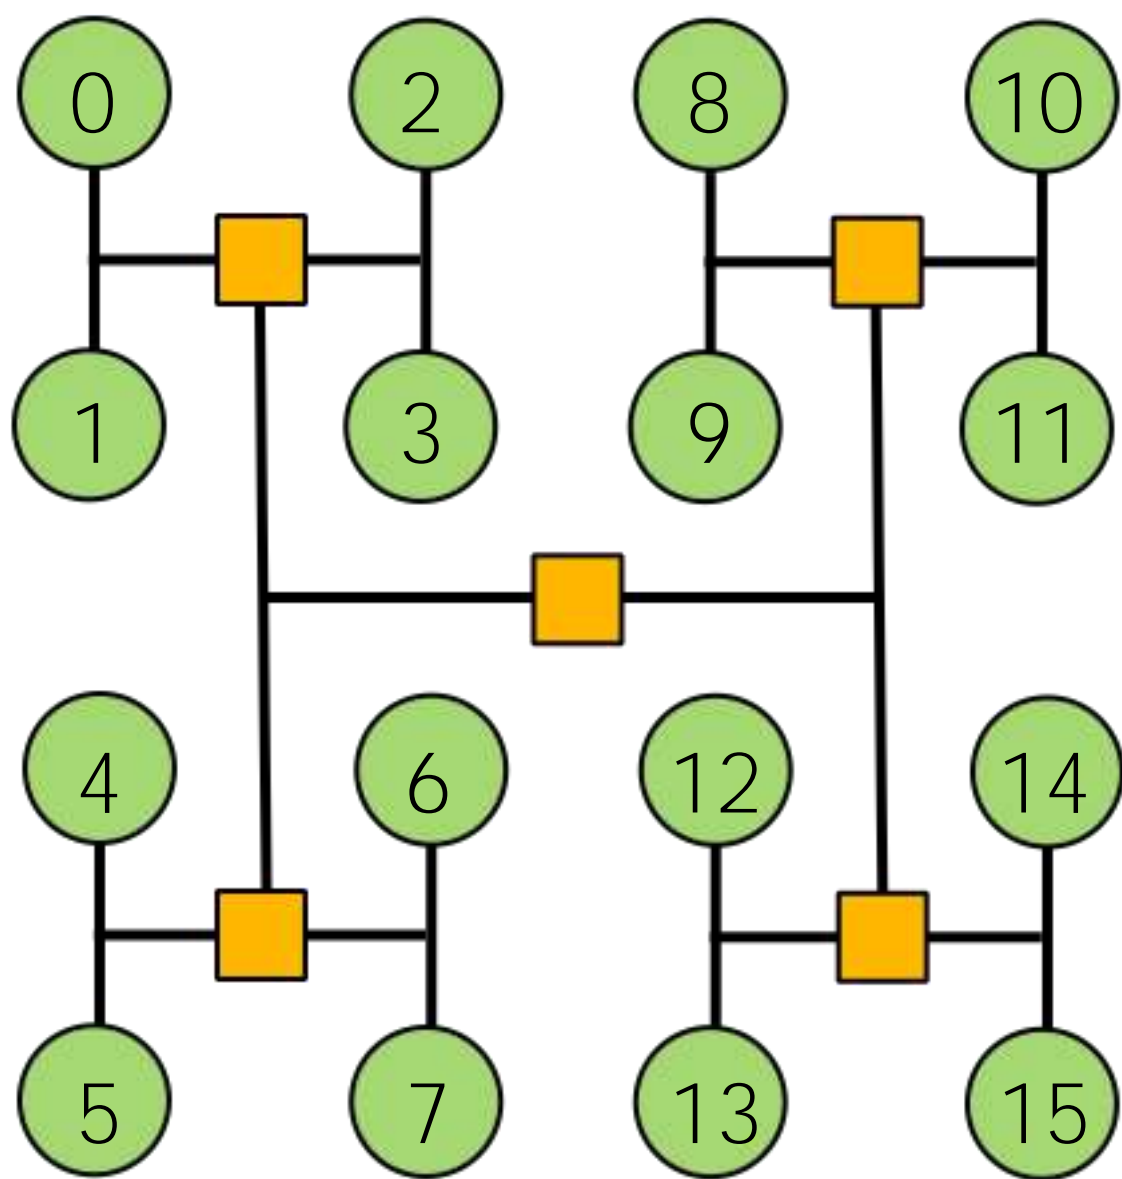
- Characteristics of mesh topology are different based on whether node is near edge or middle of network (torus topology introduces new links to avoid this problem)
- Still $O(N)$ cost, but higher cost than 2D grid
- Higher path diversity and bisection BW than mesh
- Higher complexity
 - Difficult to layout on chip
 - Unequal link lengths



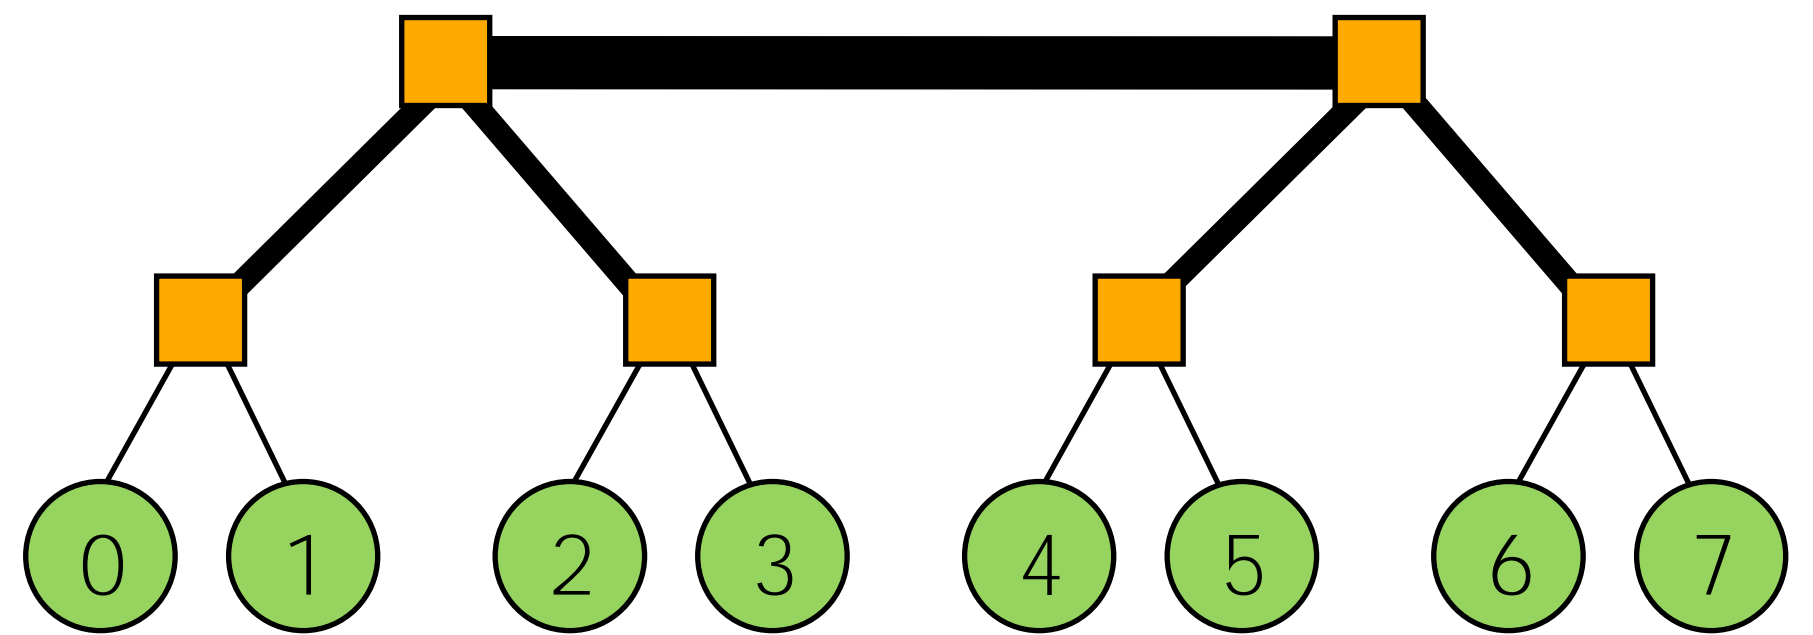
2D Torus

Trees

- Planar, hierarchical topology
- Like mesh/torus, good when traffic has locality
- Latency: $O(\lg N)$
- **Use “fat trees” to alleviate root bandwidth problem (higher bandwidth links near root)**



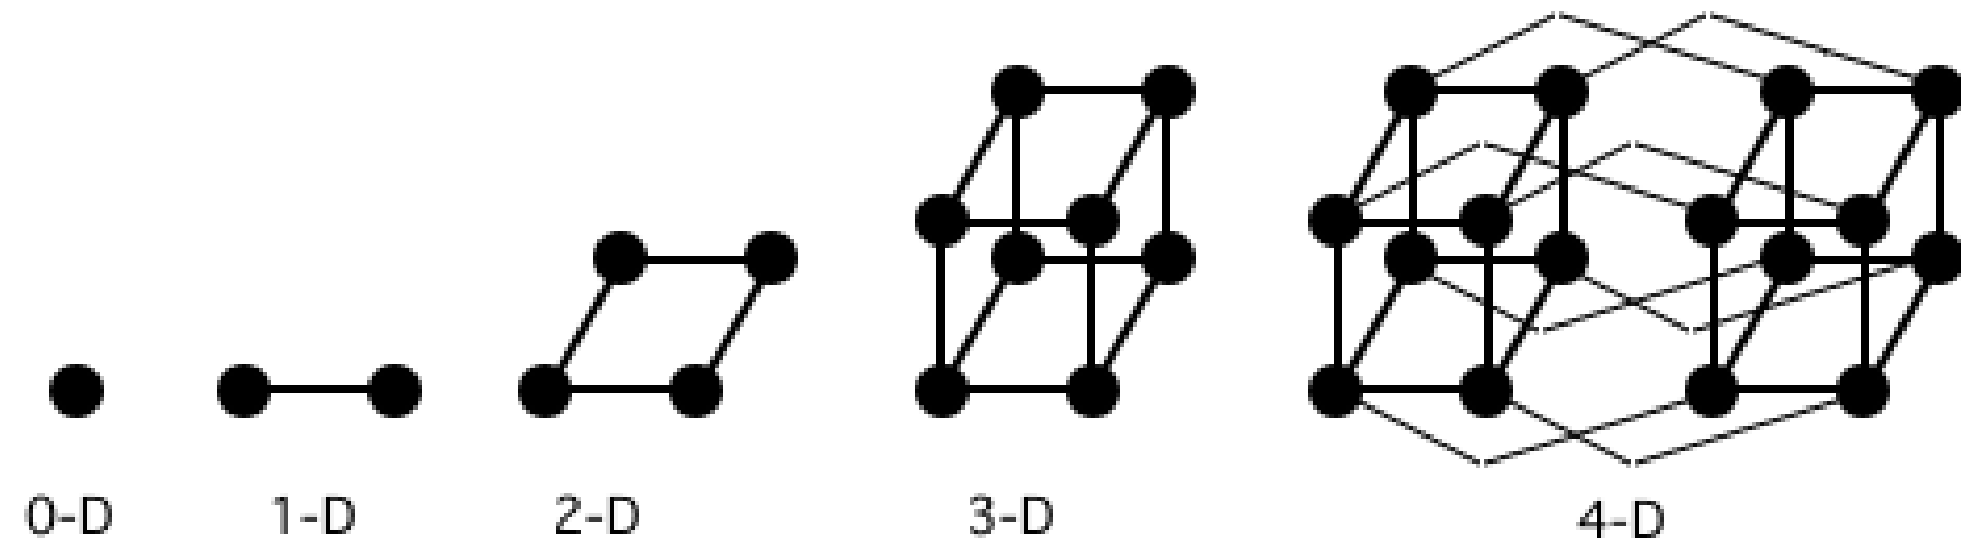
H-Tree



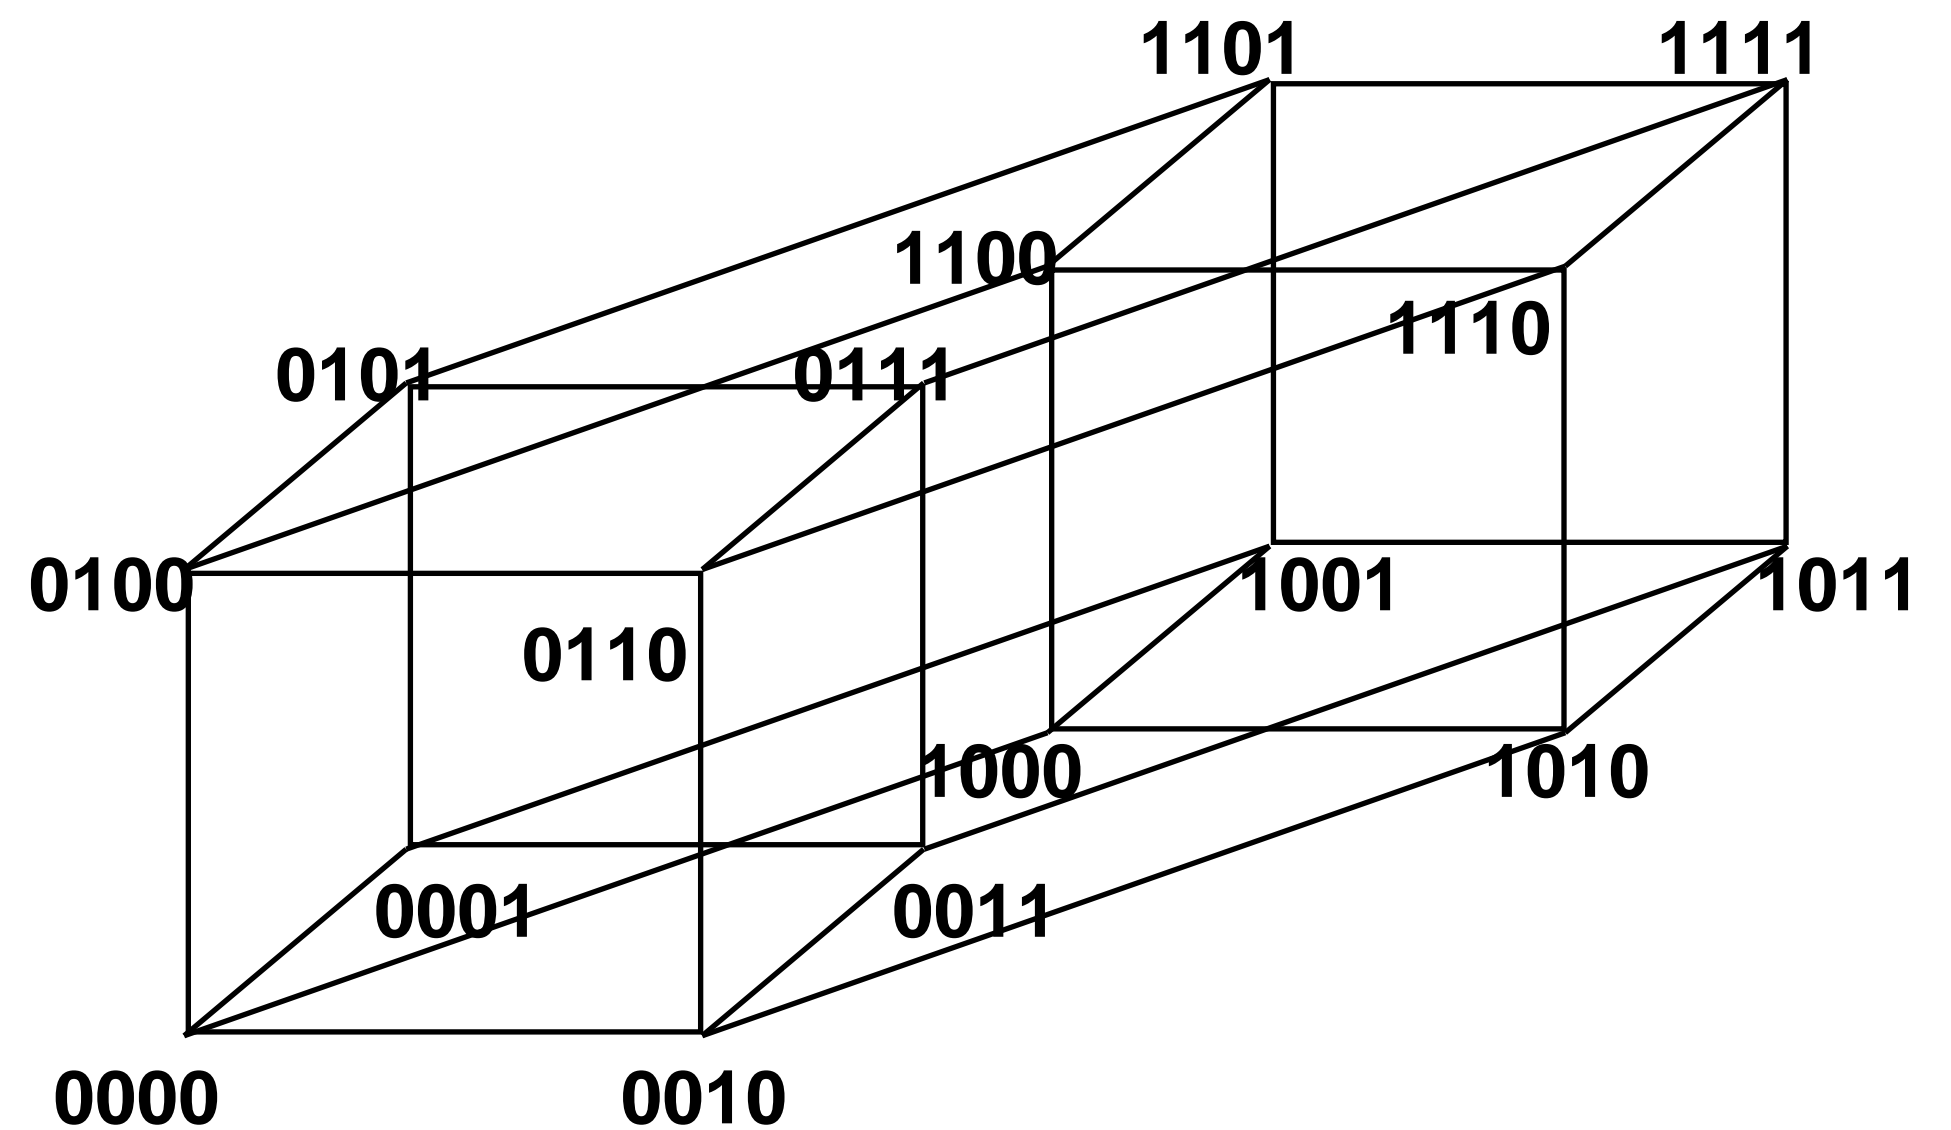
Fat Tree

Hypercube

- Low latency: $O(\lg N)$
- Radix: $O(\lg N)$
- Number of links $O(N \lg N)$

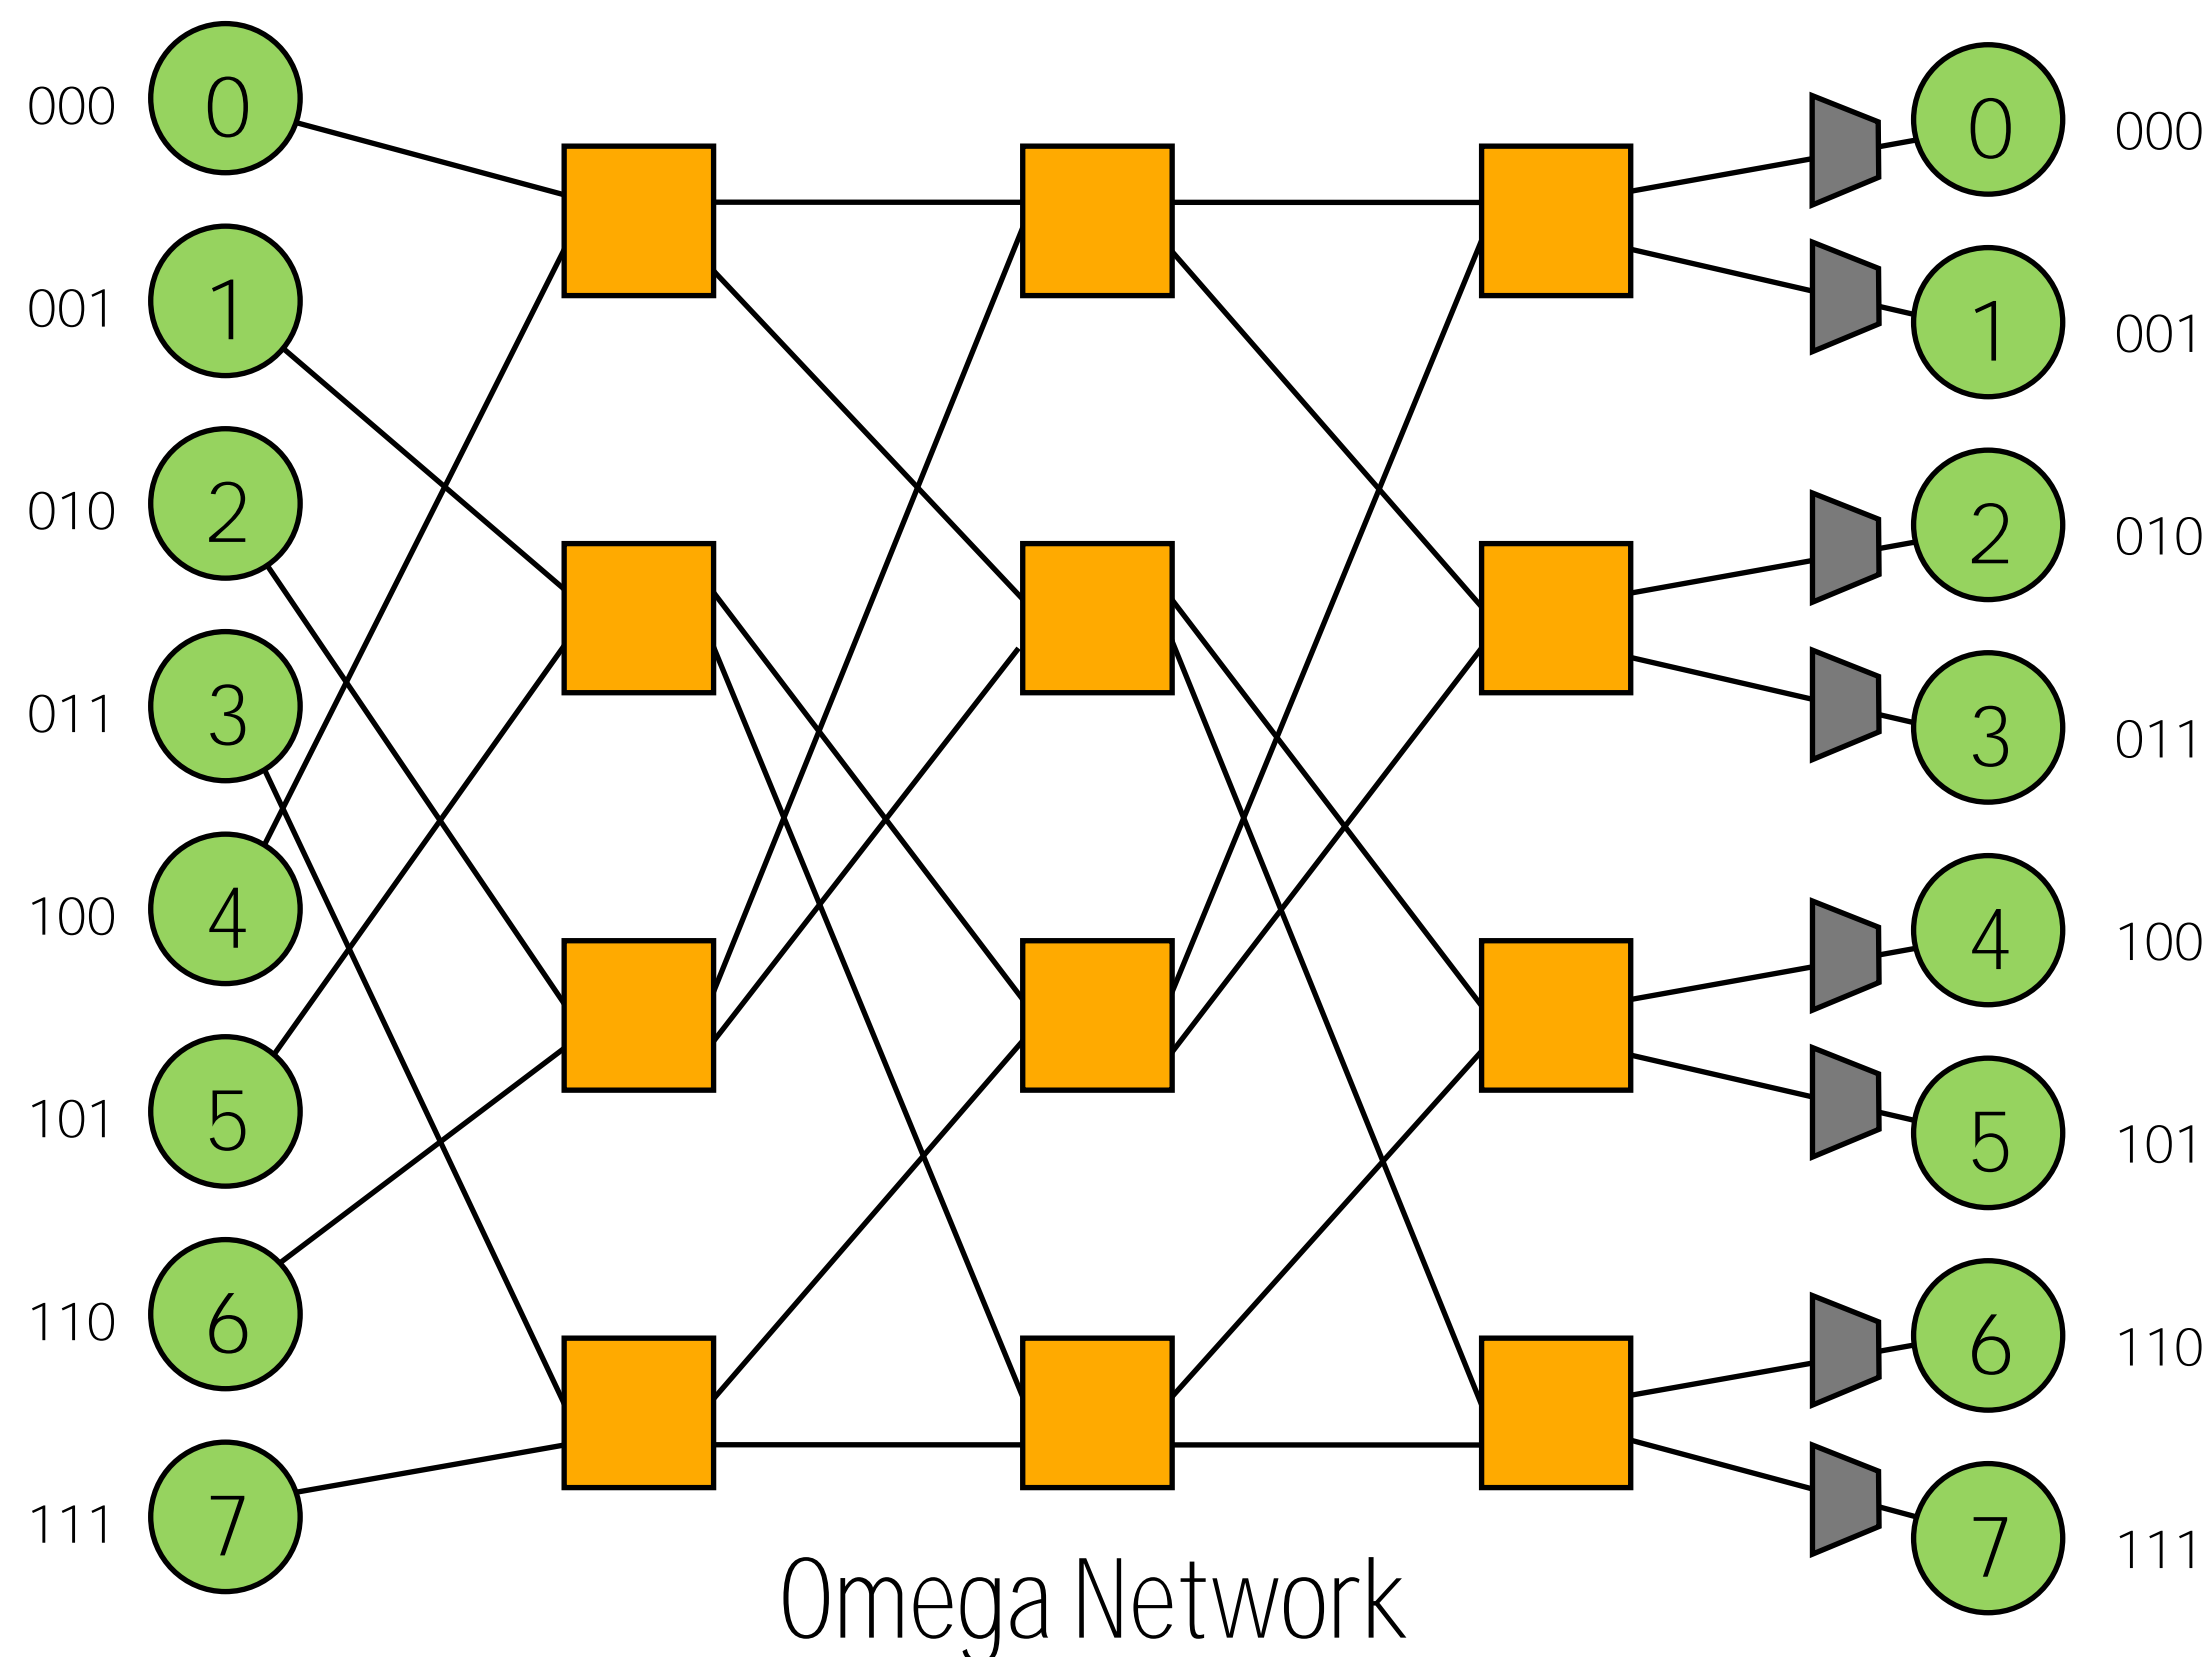


- 6D hypercube used in 64-core Cosmic Cube computer developed at Caltech in the 80s
- SGI Origin used a hypercube

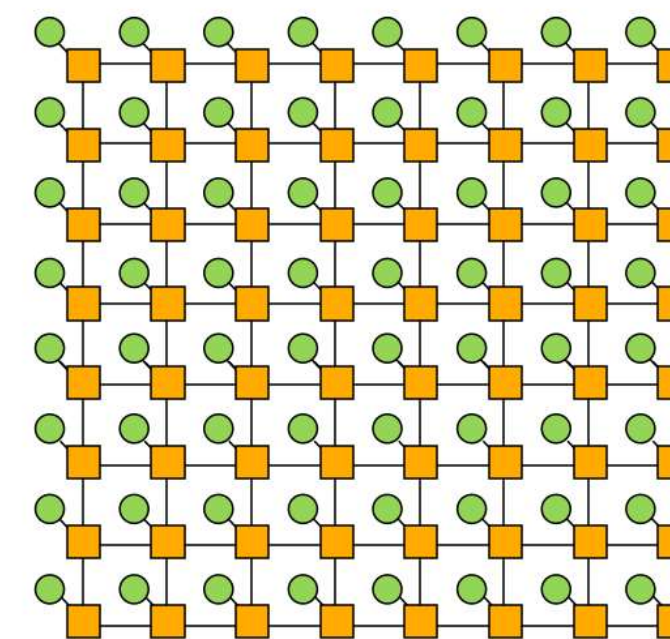
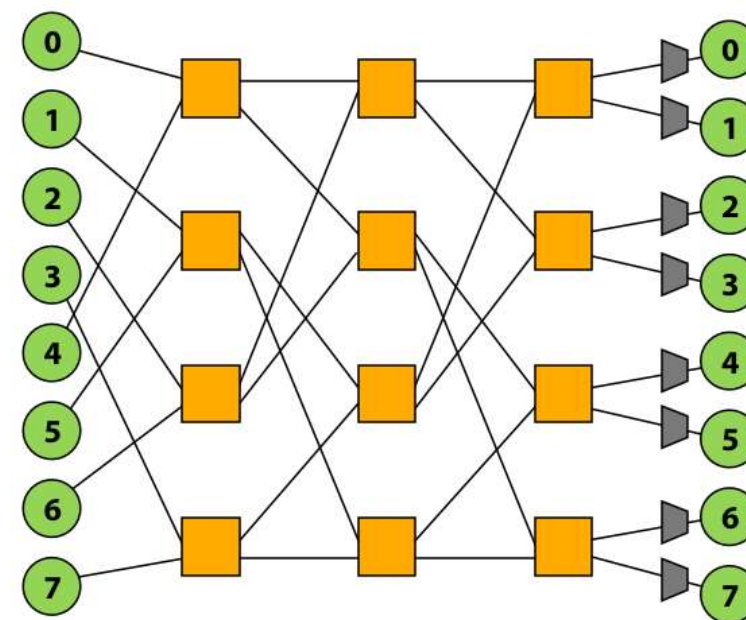
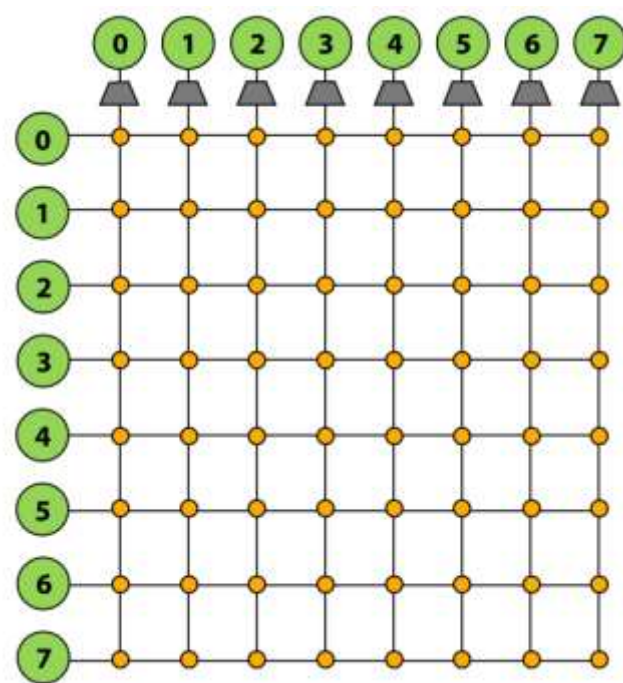


Multi-stage logarithmic

- Indirect network with multiple switches between terminals
- Cost: $O(N \lg N)$
- Latency: $O(\lg N)$
- **Many variations: Omega, butterfly, Clos networks, etc...**



Review: network topologies



Topology

Crossbar

Multi-stage log.

Mesh

Direct/Indirect

Indirect

Indirect

Direct

Blocking/
Non-blocking

Non-blocking

Blocking
(one discussed in class is, others
are not)

Blocking

Cost

$O(N^2)$

$O(N \lg N)$

$O(N)$

Latency

$O(1)$

$O(\lg N)$

$O(\sqrt{N})$
(average)

Buffering and flow control

Circuit switching vs. packet switching

- Circuit switching sets up a full path (acquires all resources) between sender and receiver prior to sending a message

- Establish route (reserve links) then send all data for message
- Higher bandwidth transmission (no per-packet link mgmt overhead)
- Does incur overhead to set up/tear down path
- Reserving links can result in low utilization



- Packet switching makes routing decisions per packet

- Route each packet individually (possibly over different network links)
- Opportunity to use link for a packet whenever link is idle
- Overhead due to dynamic switching logic during transmission
- No setup/tear down overhead

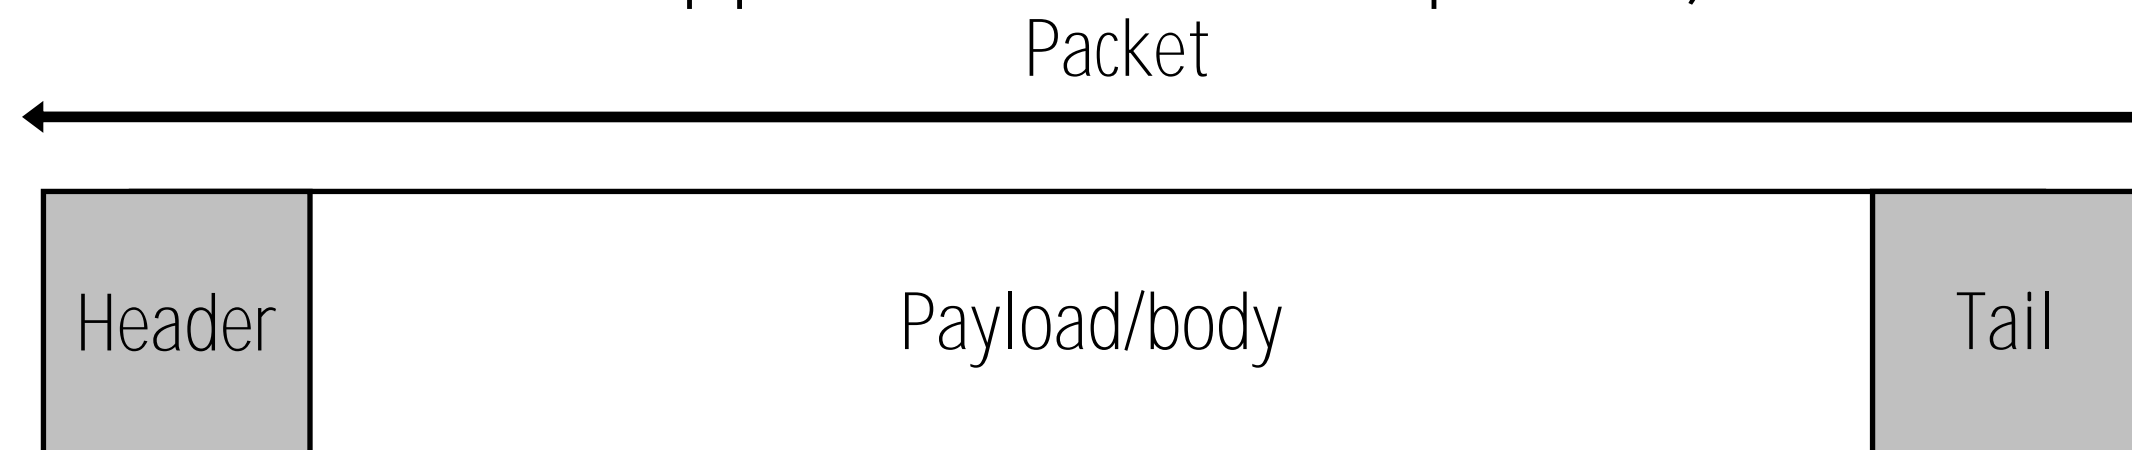


Granularity of communication

- Message
 - Unit of transfer between network clients (e.g., cores, memory)
 - Can be transmitted using many packets
- Packet
 - Unit of transfer for network
 - Can be transmitted using multiple flits (will discuss later)
- Flit (flow control digit)
 - **Packets broken into smaller units called “flits”**
 - **Flit: (“flow control digit”) a unit of flow control in the network**
 - Flits become minimum granularity of routing/buffering

Packet format

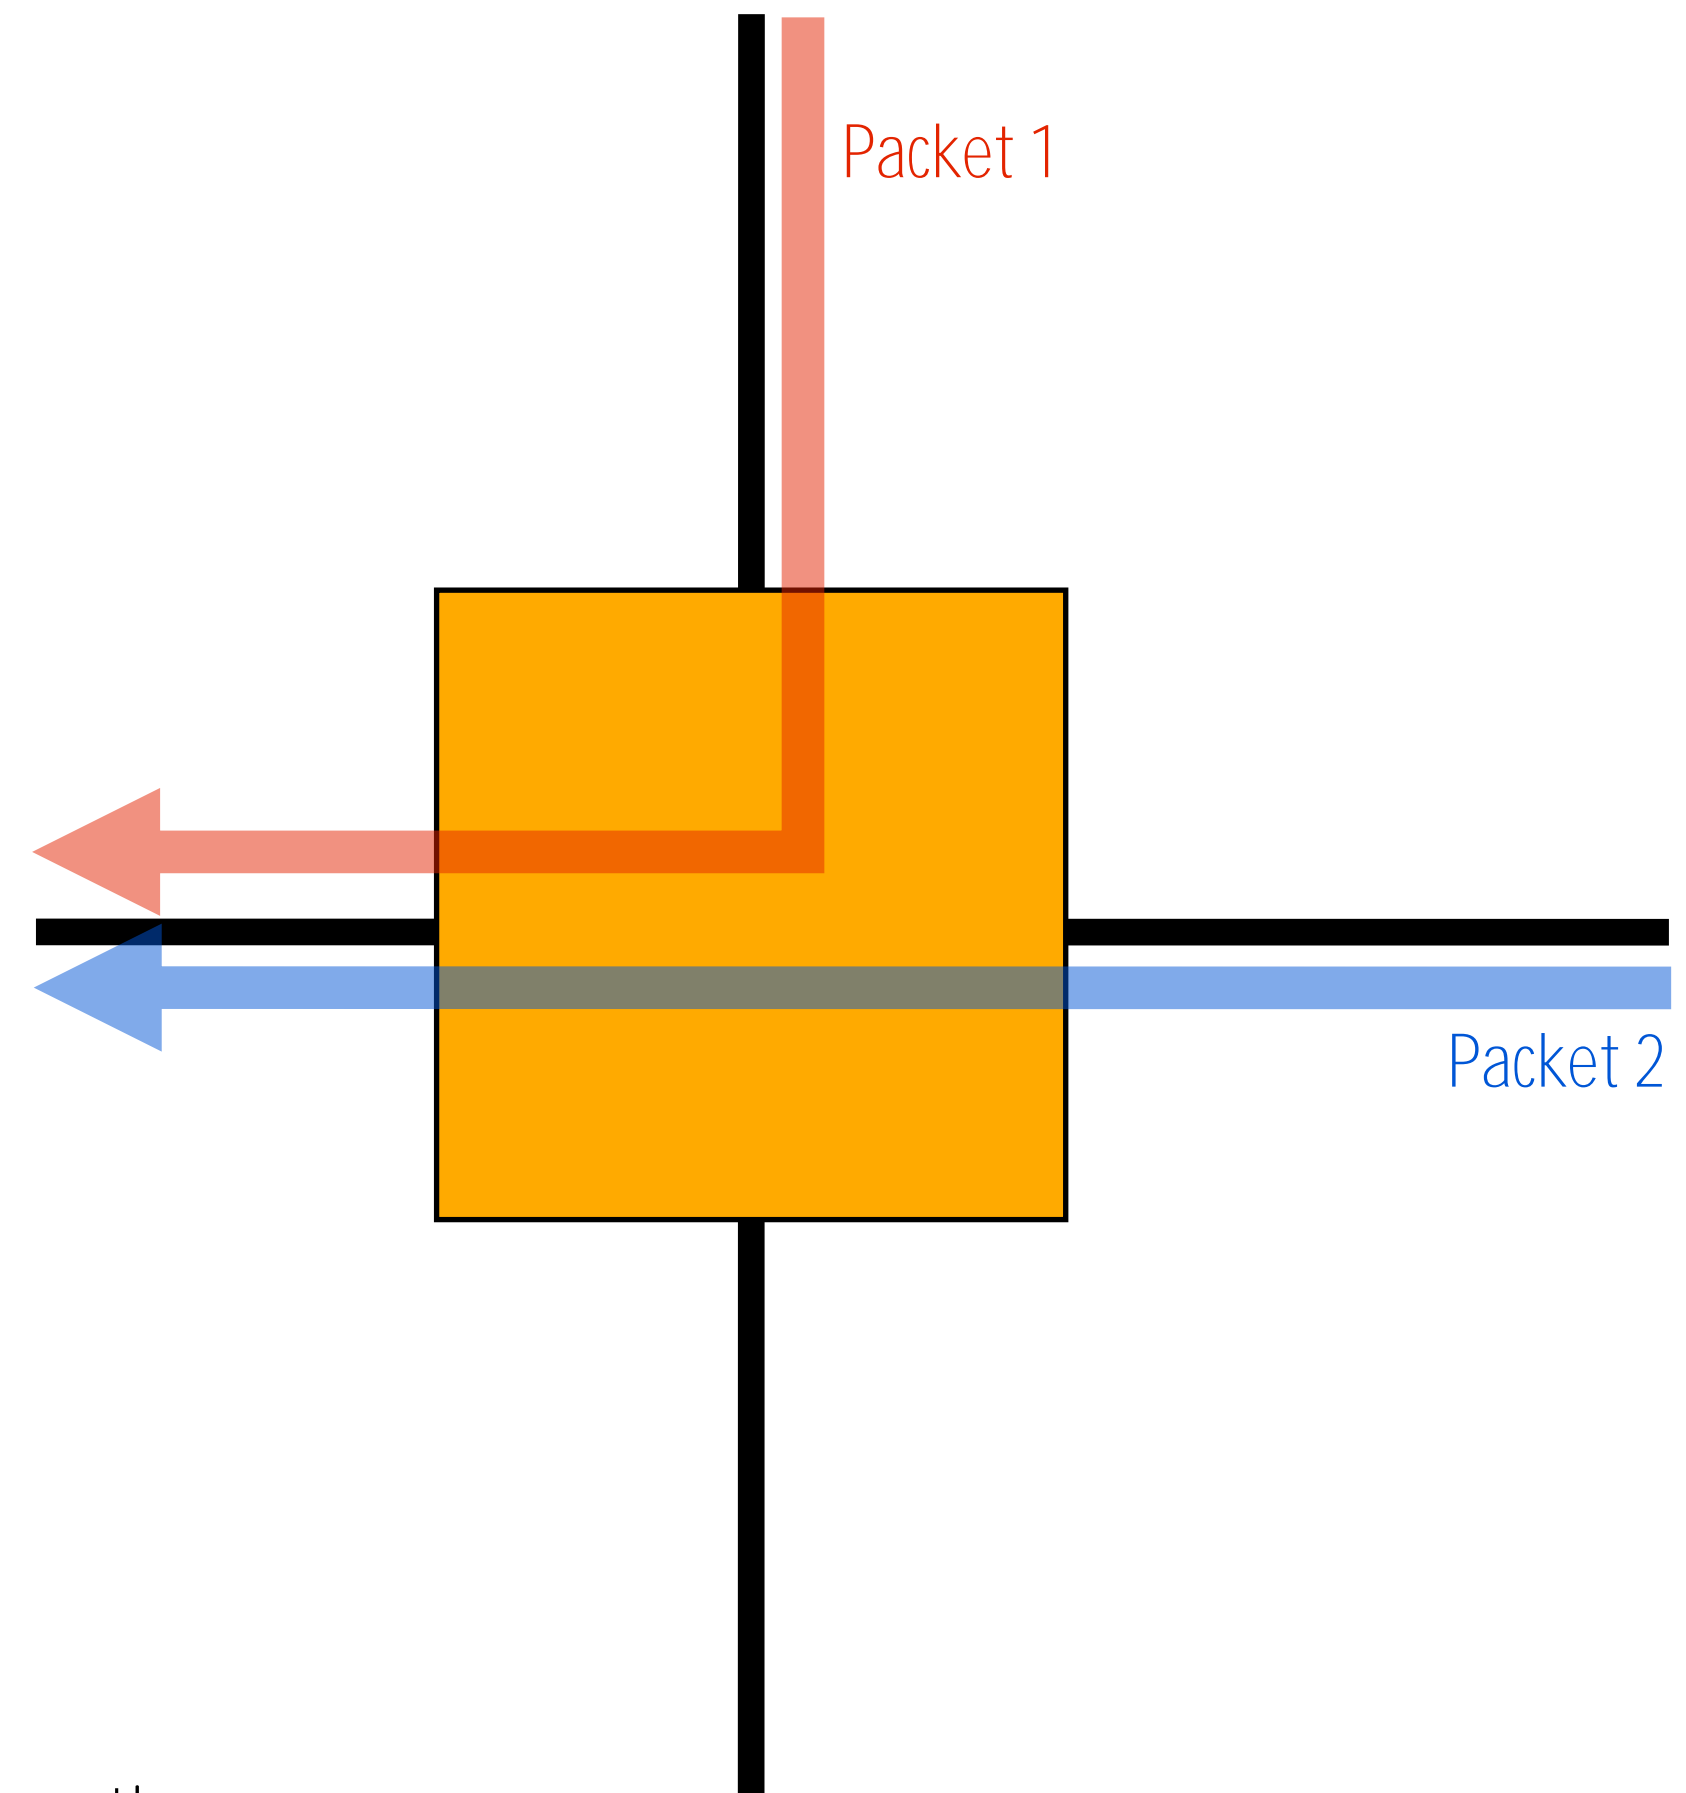
- A packet consists of:
 - Header:
 - Contains routing and control information
 - At start of packet to router can start forwarding early
 - Payload/body: containing the data to be sent
 - Tail
 - Contains control information, e.g., error code
 - **Generally located at end of packet so it can be generated “on the way out”**
- (sender computes checksum, appends it to end of packet)



Handling contention

Scenario: two packets need to be routed onto the same outbound link at the same time

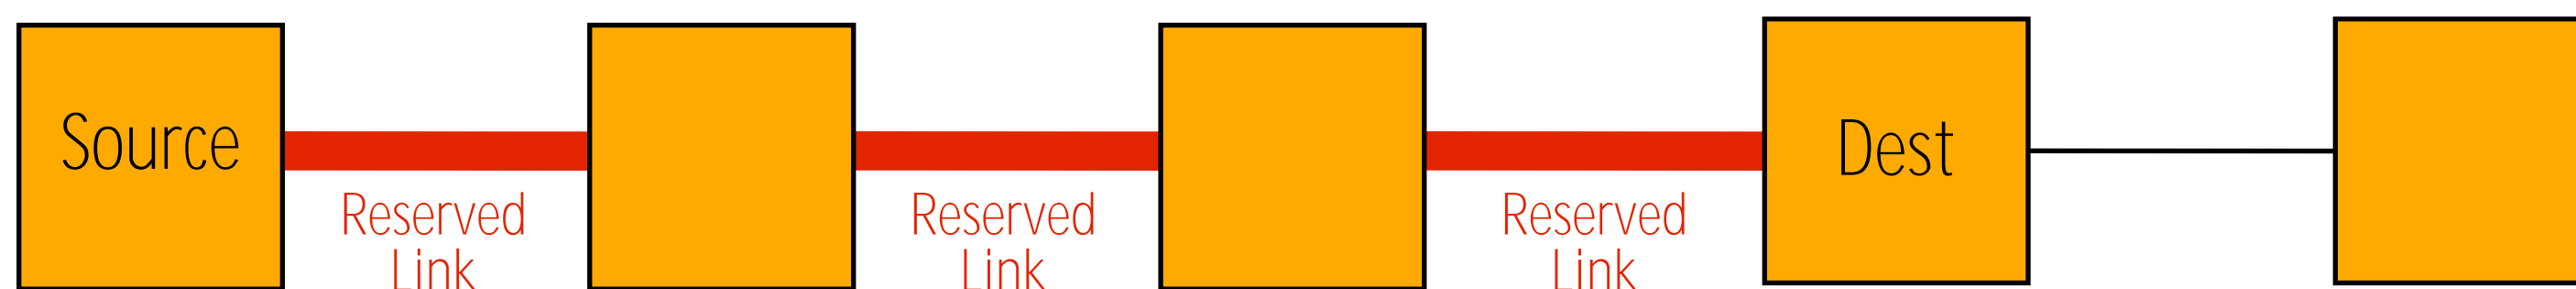
- Options:
 - Buffer one packet, send it over link later
 - Drop one packet
 - Reroute one packet (deflection)
- In this lecture: we only consider buffering *



* But recent research has looked at using bufferless networks with deflection routing as a power-efficient interconnect for chip multiprocessors.

Circuit-switched routing

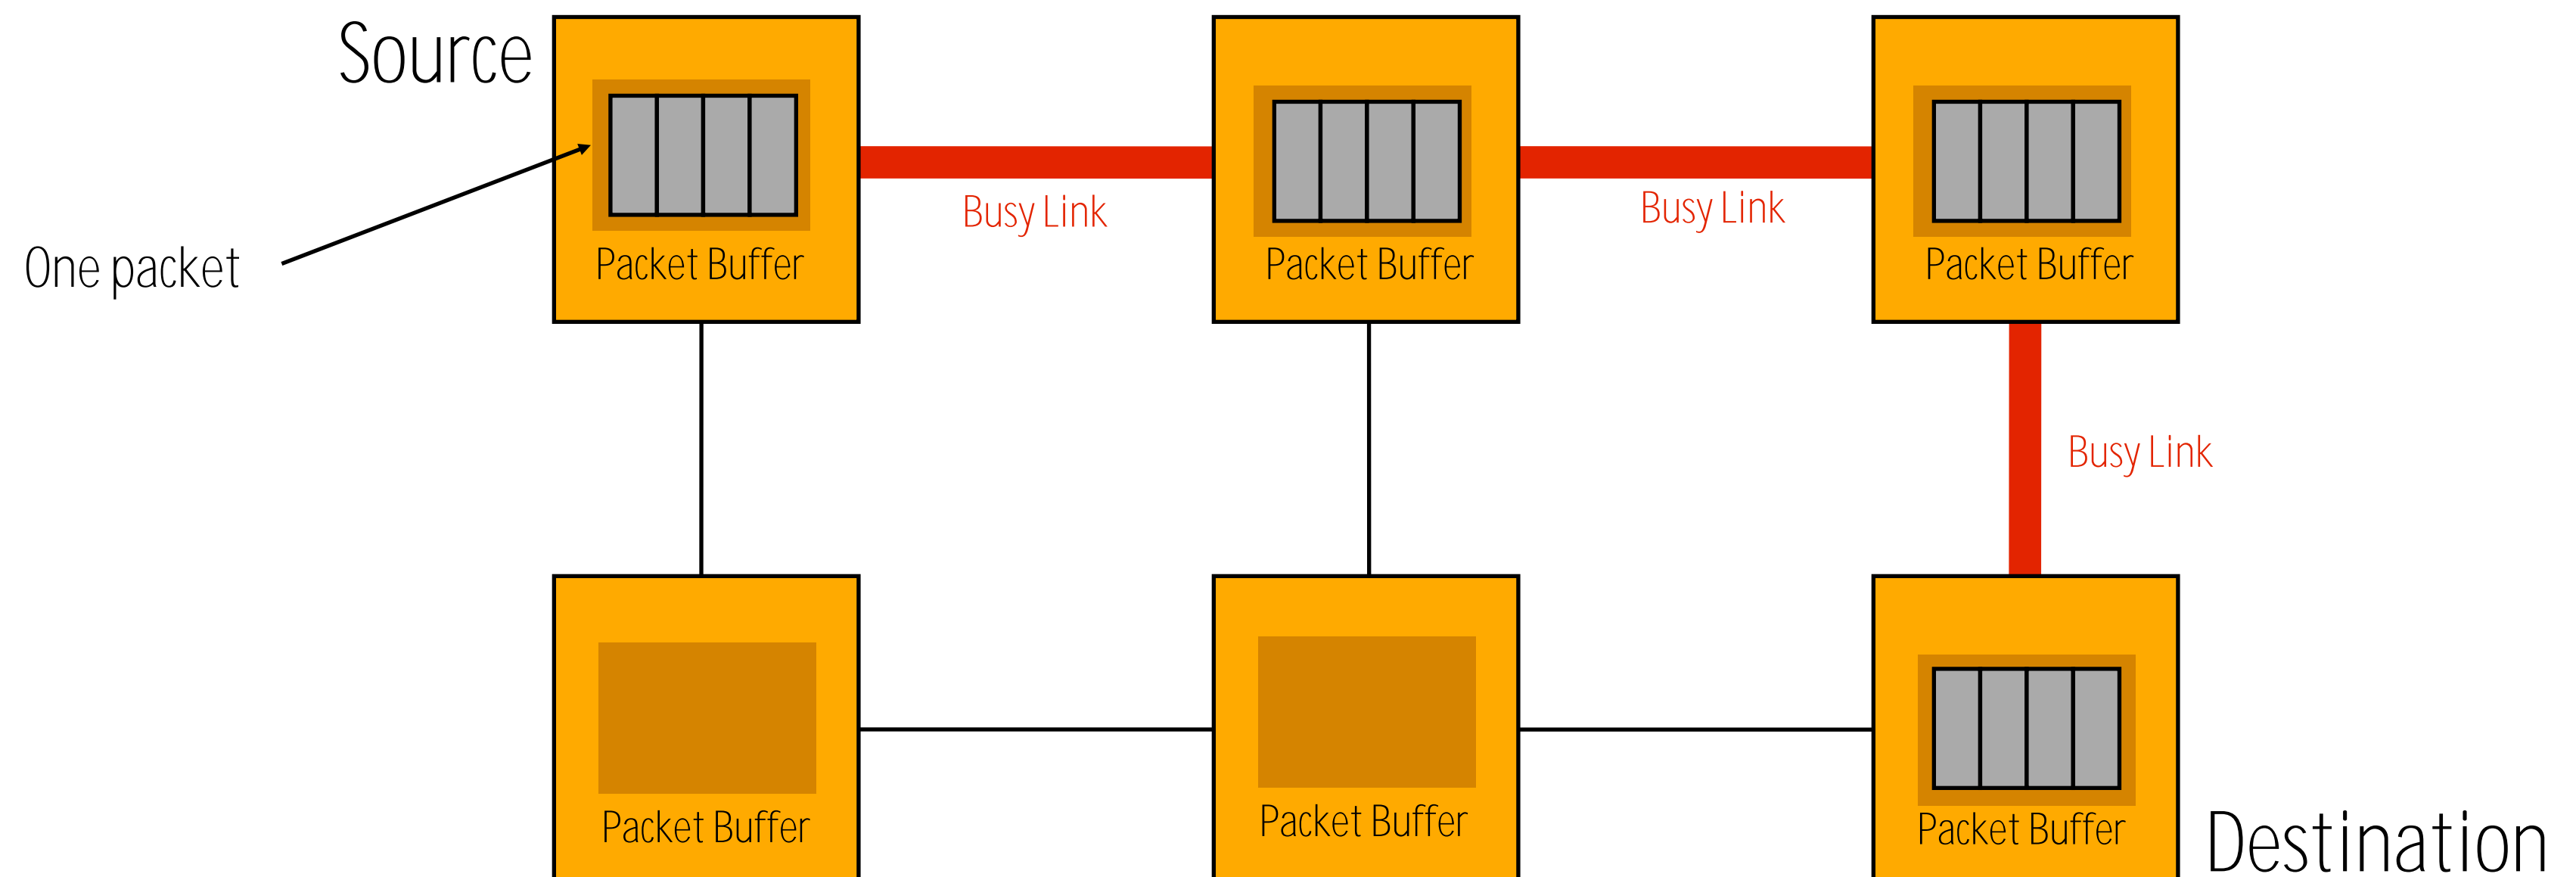
- High-granularity resource allocation
 - Main idea: pre-allocate all resources (links across multiple switches) **along entire network path for a message (“setup a flow”)**



- Costs
 - **Needs setup phase (“probe”) to set up the path (and to tear it down and release the resources when message complete)**
 - Lower link utilization. Transmission of two messages cannot share same link (even if some resources on a preallocated path are no longer utilized during a transmission)
- Benefits
 - No contention during transmission due to preallocation, so no need for buffering
 - Arbitrary message sizes (once path is set up, send data until done)

Store-and-forward (packet-based routing)

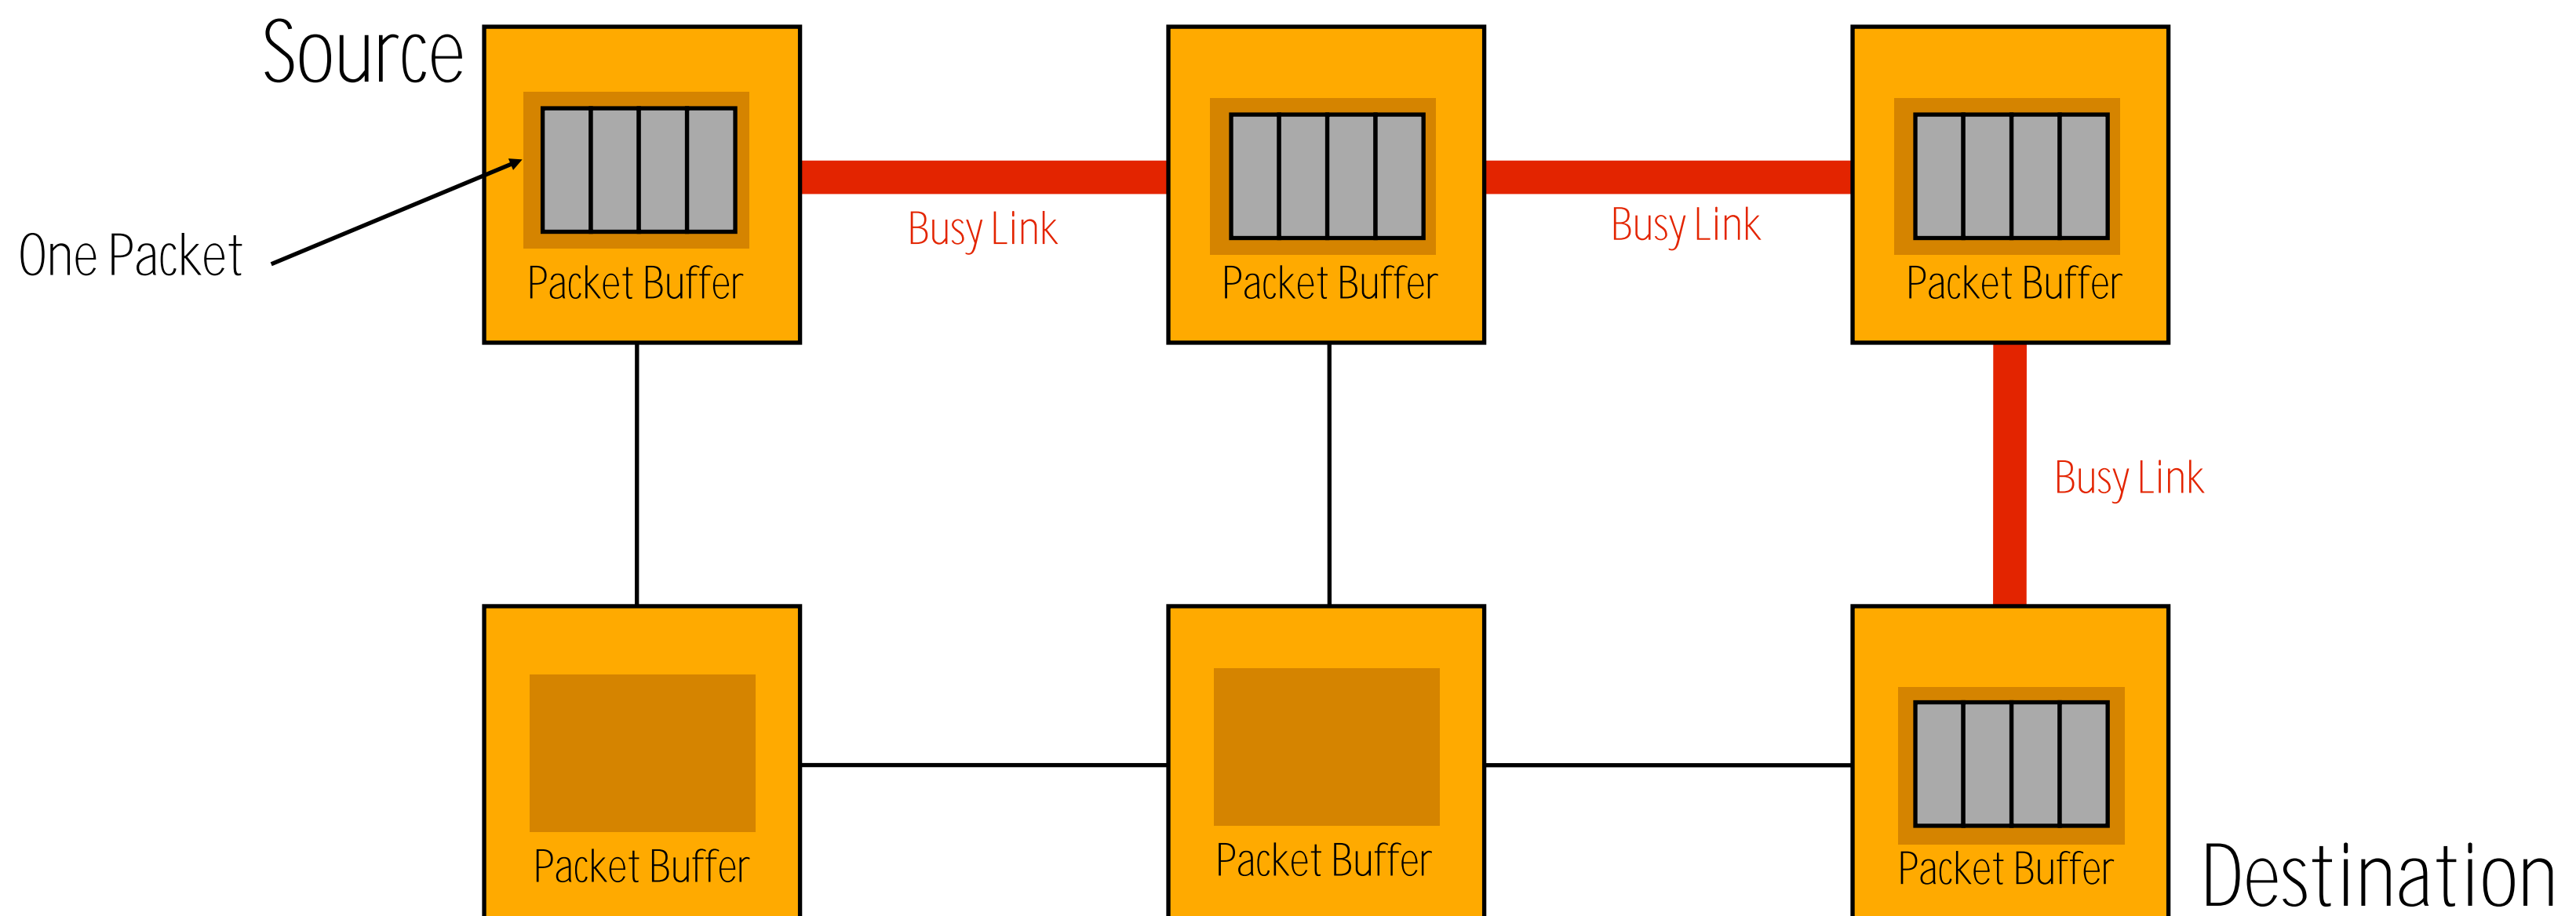
- Packet copied entirely into network switch before moving to next node
- Flow control unit is an entire packet
 - Different packets from the same message can take different routes, but all data in a packet is transmitted over the same route
- Requires buffering for entire packet in each router
- High per-packet latency (latency = packet transmission time on link x network distance)



Note to students: in lecture this slide was animated and the final build shown here is not illustrative of store-and-forward routing concept (please refer to lecture video)

Cut-through flow control (also packet-based)

- Switch starts forwarding data on next link as soon as packet header is received (header determines how much link bandwidth packet requires + where to route)
- Result: reduced transmission latency
 - Cut-through routing reduces to store-and-forward under high contention. Why?



Store and forward solution from previous slide: 3 hops x 4 units of time to transmit packet over a single link = 12 units of time

Cut-through solution: 3 steps of latency for head of packet to get to destination + 3 units of time for rest of packet = 6 units of time

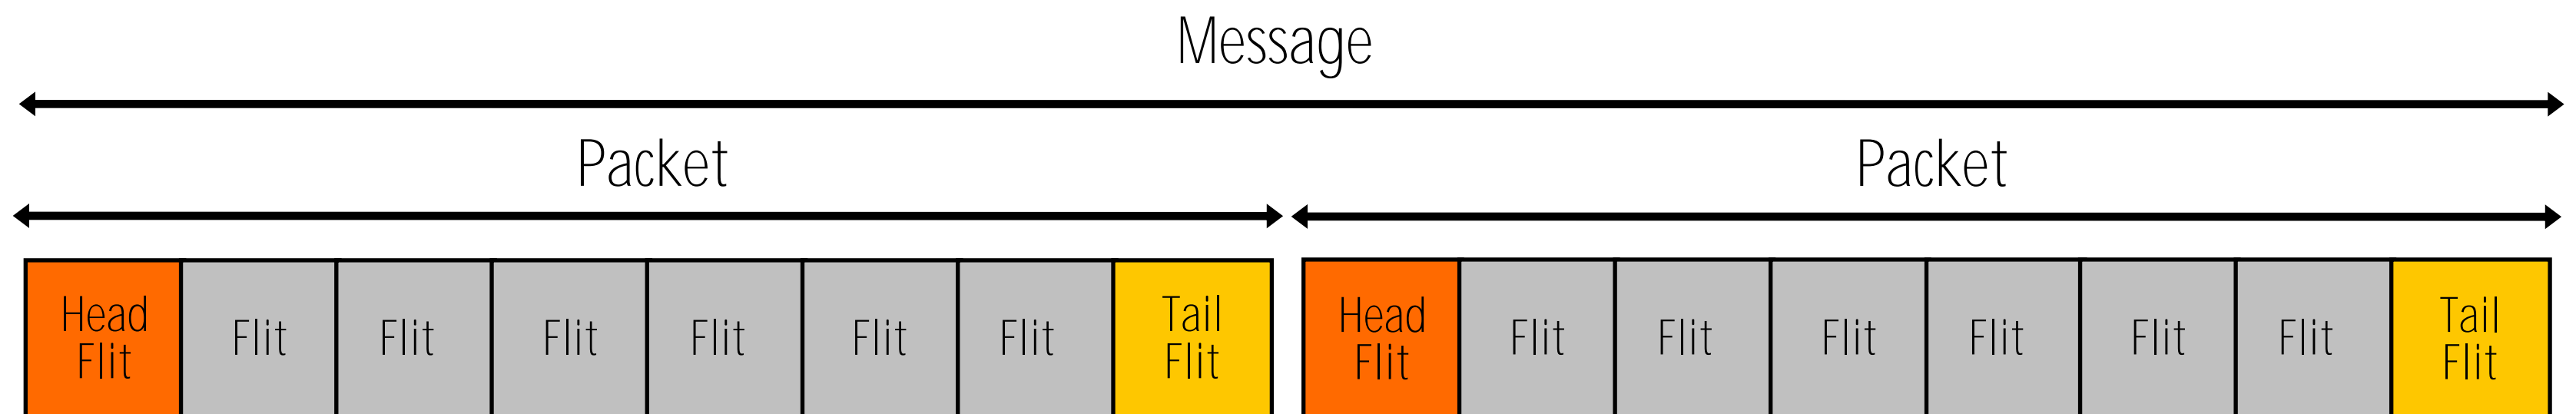
Note to students: in lecture this slide was animated and the final build shown here is not illustrative of the cut-through routing concept (please refer to lecture video)

Cut-through flow control

- If output link is blocked (cannot transmit head), transmission of tail can continue
 - Worst case: entire message is absorbed into a buffer in a switch (cut-through flow control degenerates to store-and-forward in this case)
 - Requires switches to have buffering for entire packet, just like store-and-forward

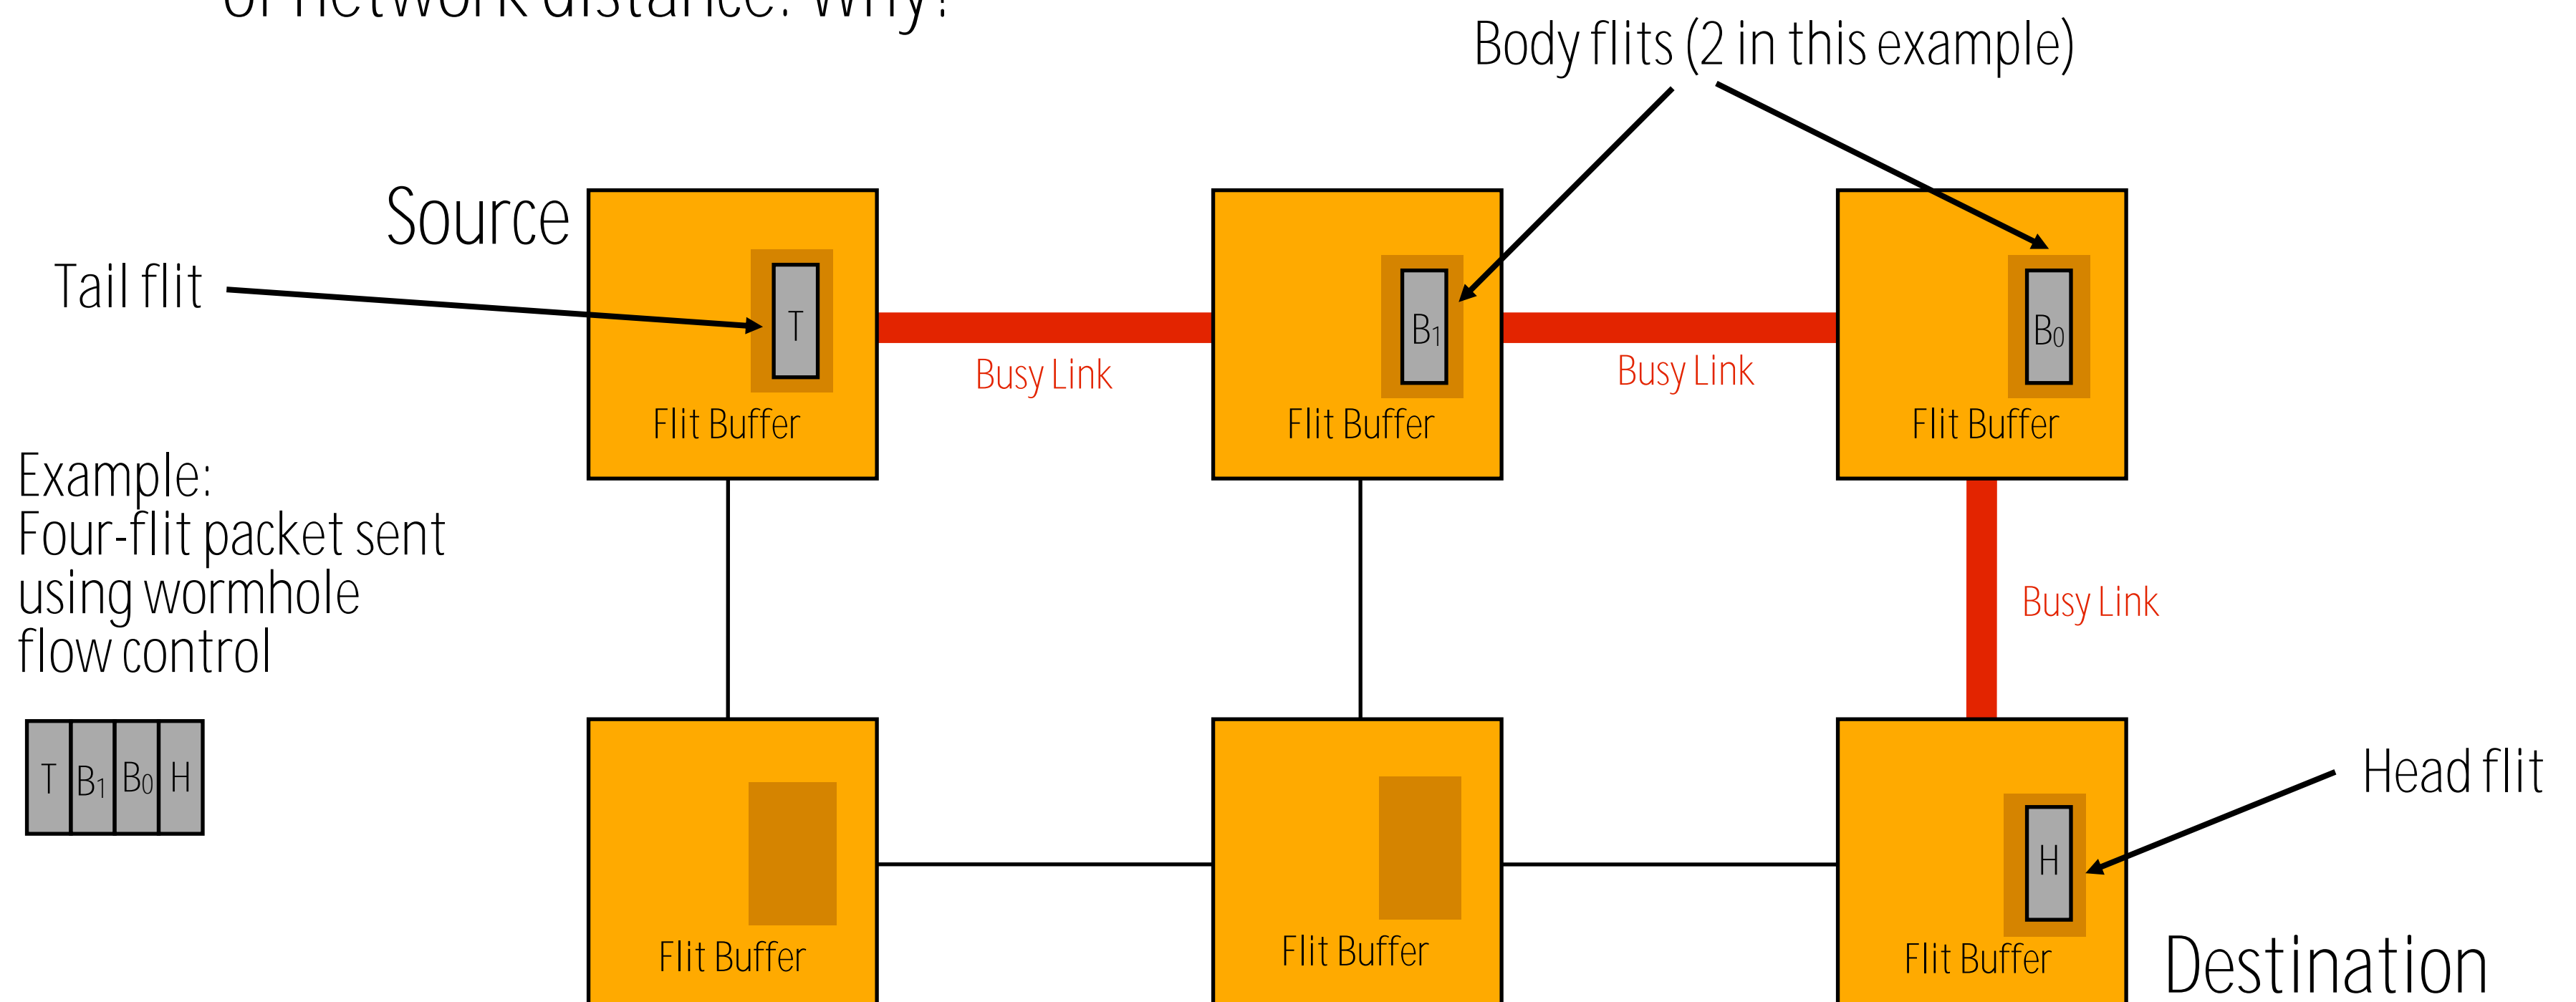
Wormhole flow control

- Flit (flow control digit)
 - **Packets broken into smaller units called “flits”**
 - **Flit: (“flow control digit”) a unit of flow control in the network**
 - Flits become minimum granularity of routing/buffering
 - Recall: up until now, packets were the granularity of transfer AND flow control and buffering (store-and-forward, cut-through routing)

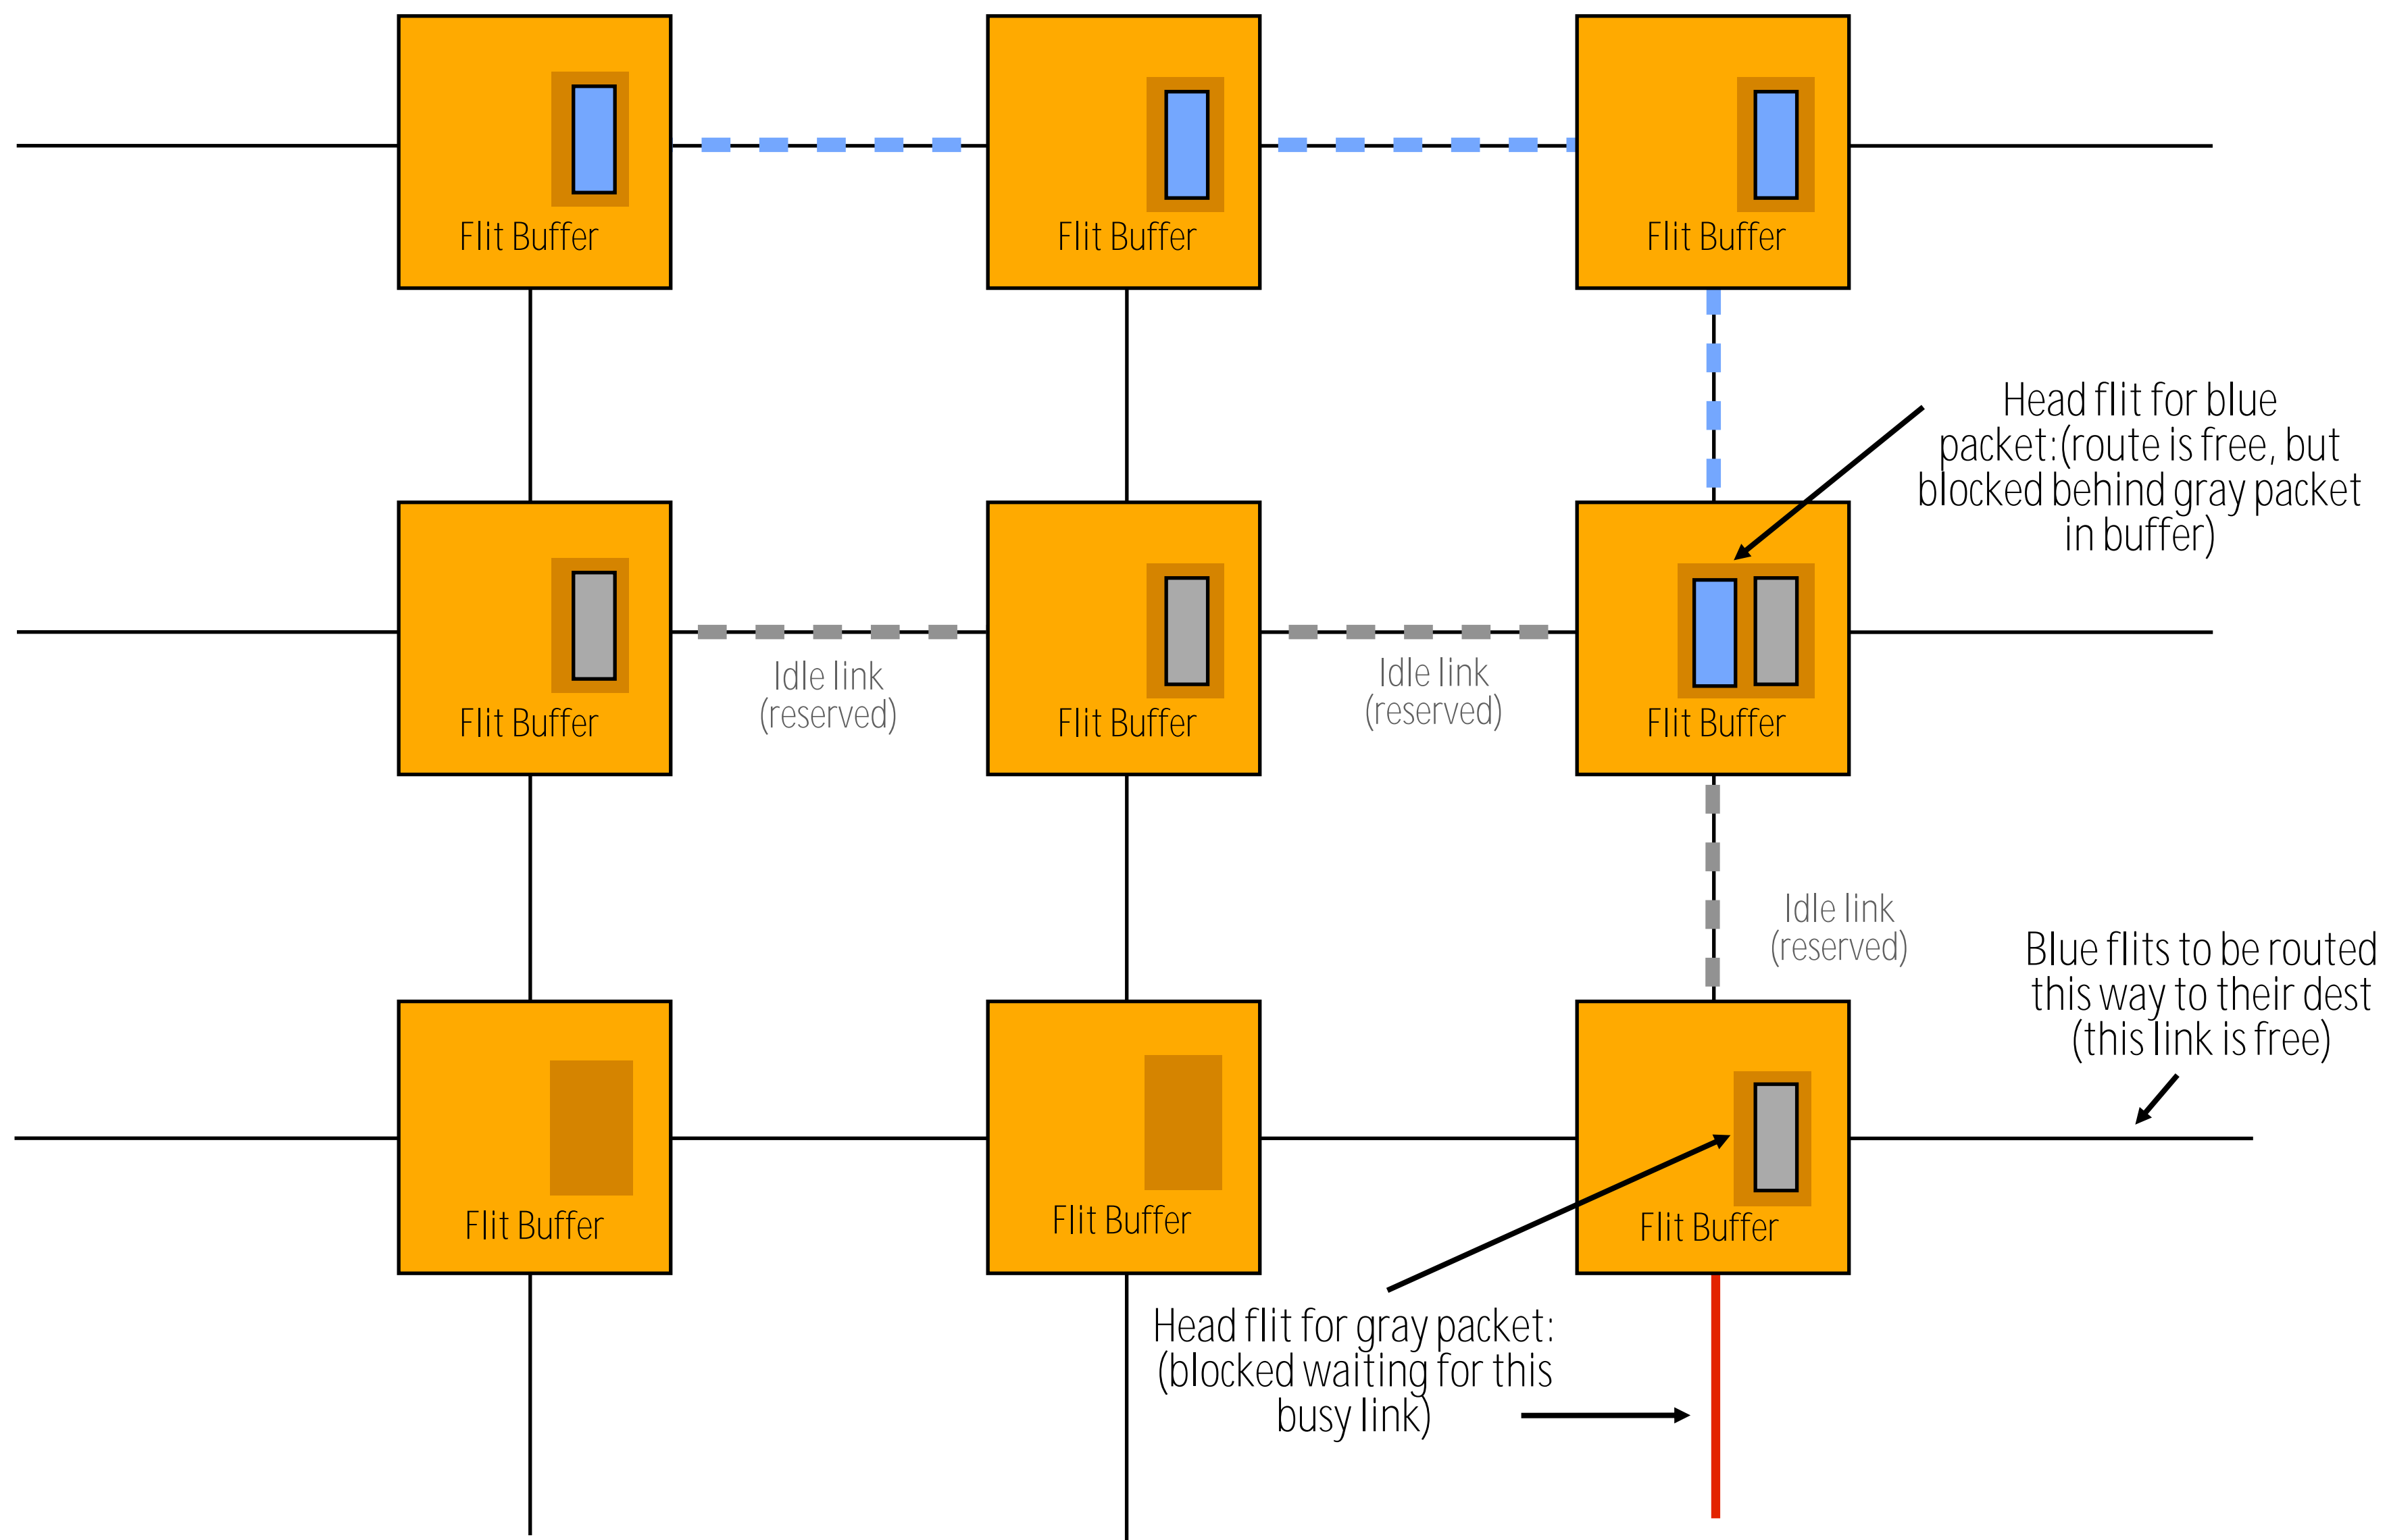


Wormhole flow control

- Routing information only in head flit
- Body flits follows head, tail flit flows body
- If head flit blocks, rest of packet stops
- Completely pipelined transmission
 - For long messages, latency is almost entirely independent of network distance. Why?

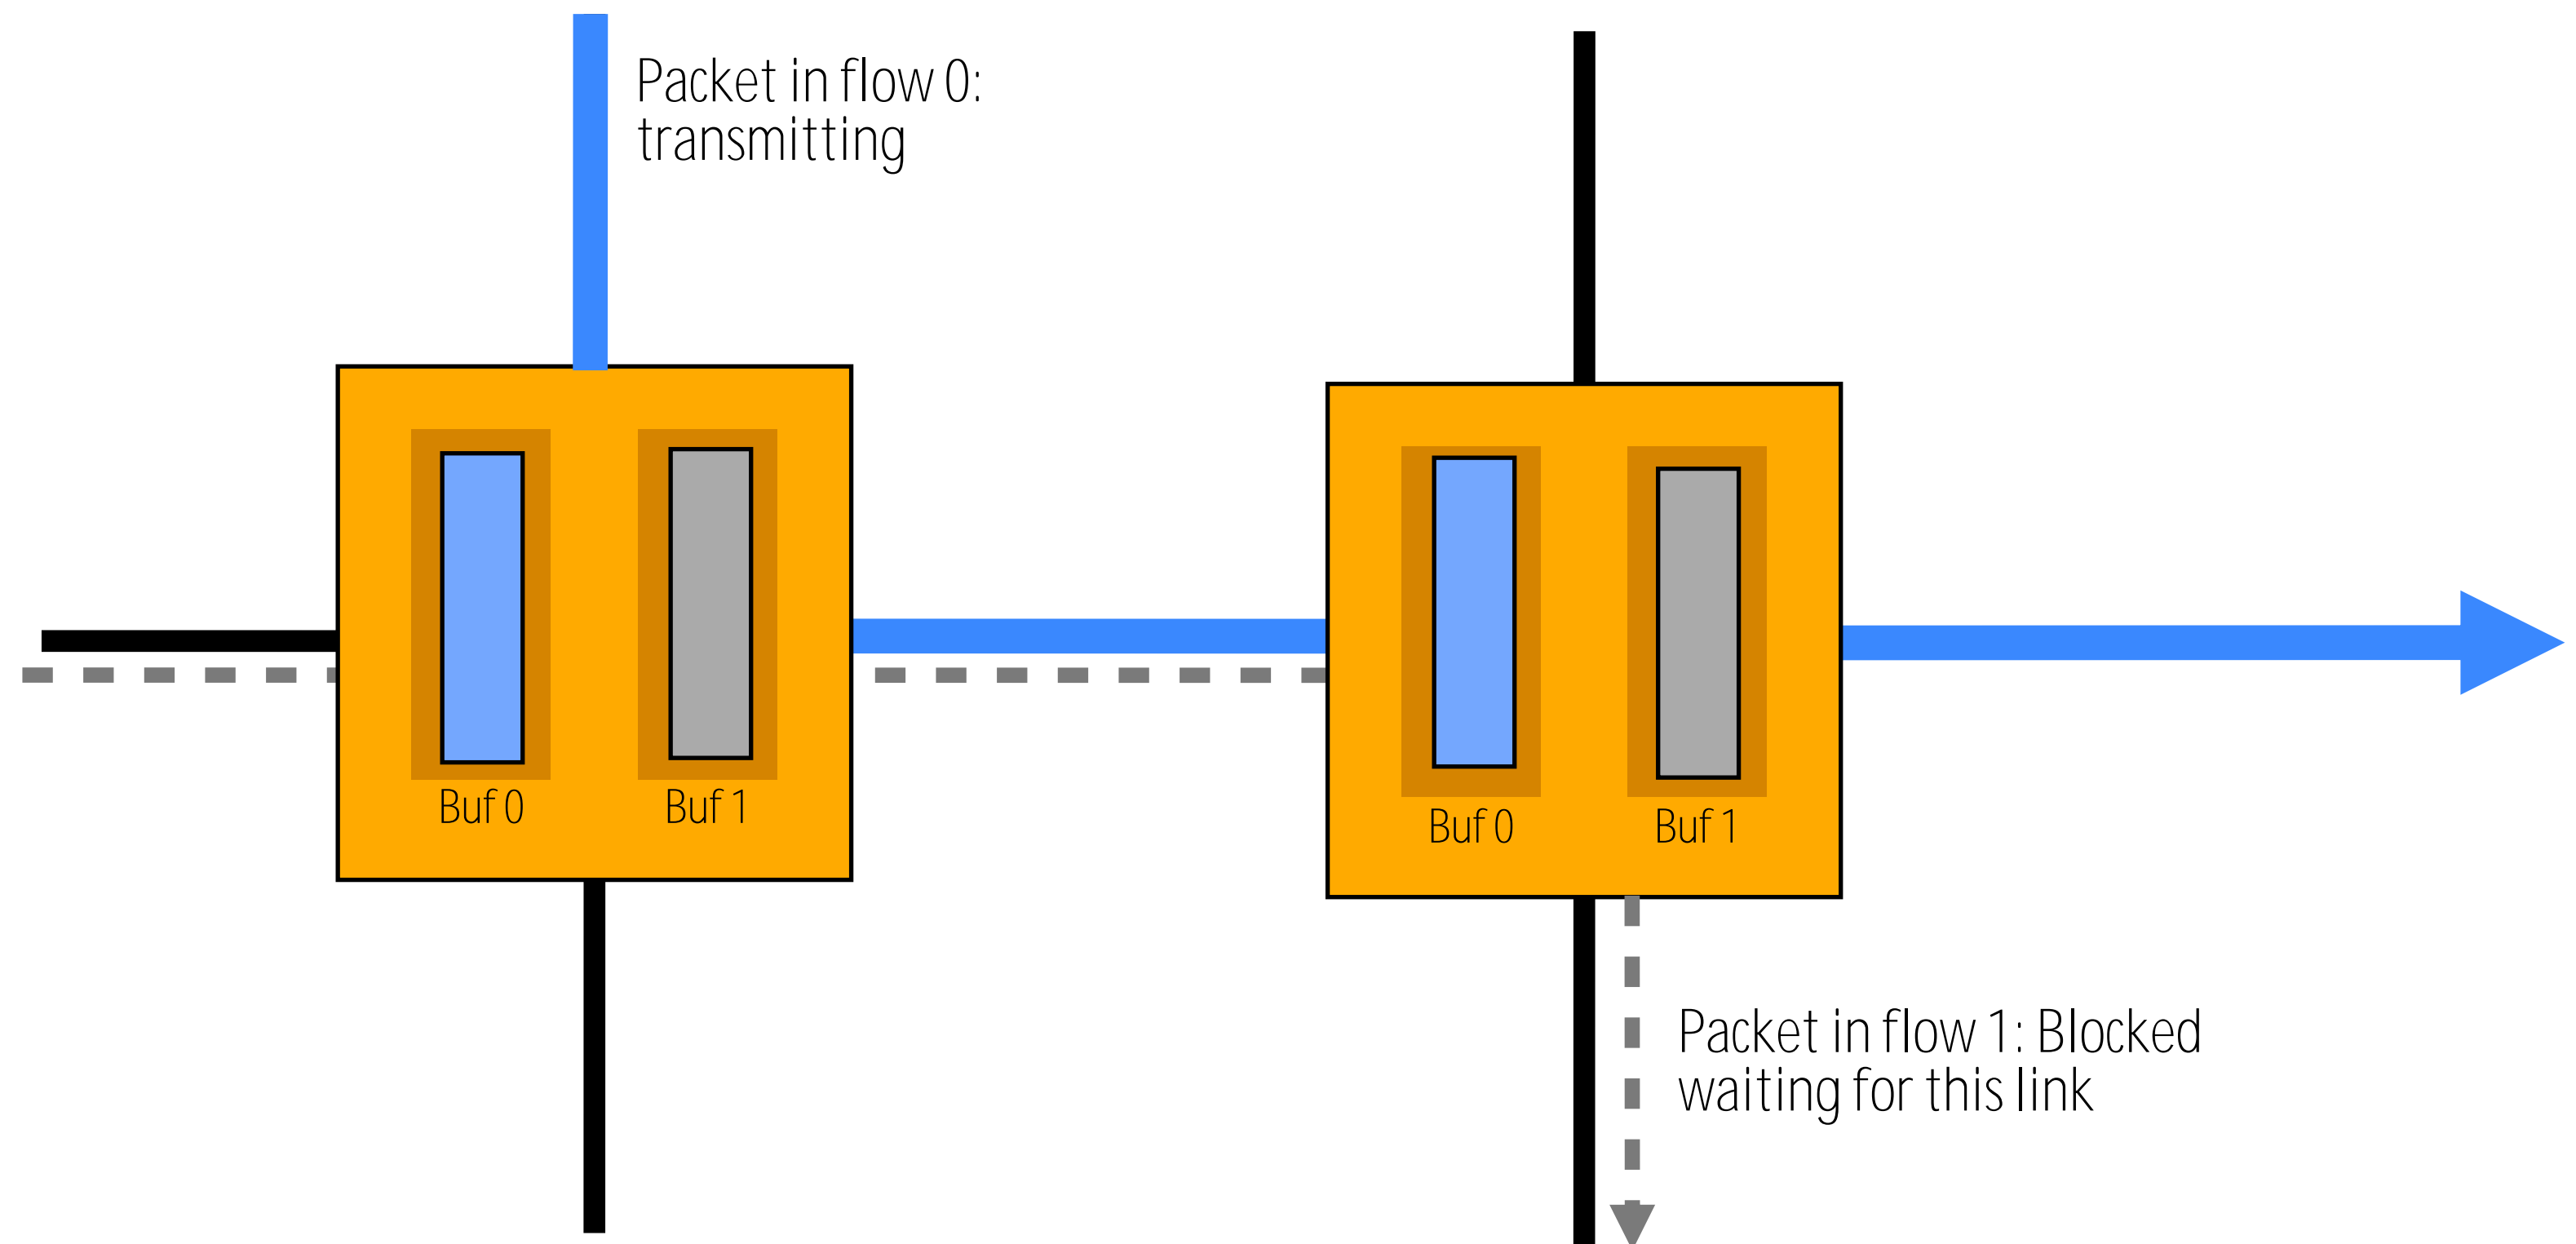


Problem: head-of-line blocking



Virtual channel flow control

- Multiplex multiple operations over single physical channel
- **Divide switch's input buffer into multiple buffers sharing a single physical channel**
- Reduces head-of-line blocking



Other uses of virtual channels

- Deadlock avoidance
 - Can be used to break cyclic dependency of resources
 - Prevent cycles by ensuring requests and responses use different virtual channels
 - **“Escape” VCs: retain at least one virtual channel that uses deadlock-free routing**
- Prioritization of traffic classes
 - Provide quality-of-service guarantees
 - Some virtual channels have higher priority than others

Current research topics

- Energy efficiency of interconnections
 - Interconnect can be energy intensive (~35% of total chip power in MIT RAW research processor)
 - Bufferless networks
 - Other techniques: turn on/off regions of network, use fast and slow networks
- Prioritization and quality-of-service guarantees
 - Prioritize packets to improve multi-processor performance (e.g., some applications may be more sensitive to network performance than others)
 - Throttle endpoints (e.g., cores) based on network feedback
- New/emerging technologies
 - Die stacking (3D chips)
 - Photonic networks-on-chip (use optical waveguides instead of wires)
 - Reconfigurable devices (FPGAs): create custom interconnects tailored to application (see CMU projects: CONNECT, CoRAM, Shrinkwrap)

Summary

- The performance of the interconnection network in a modern multi-processor is critical to overall system performance
 - Buses do not scale to many nodes
 - Historically interconnect was off-chip network connecting sockets, boards, racks
 - Today, all these issues apply to the design of on-chip networks
- Network topologies differ in performance, cost, complexity tradeoffs
 - e.g., crossbar, ring, mesh, torus, multi-stage network, fat tree, hypercube
- Challenge: efficiently routing data through network
 - Interconnect is a precious resource (communication is expensive!)
 - Flit-based flow control: fine-grained flow control to make good use of available link bandwidth
 - If interested, much more to learn about (not discussed in this class): ensuring quality-of-service, prioritization, reliability, deadlock, livelock, etc.